

Package ‘TempleCBE’

September 18, 2026

Type Package

Title Temple CBE Biostatistics, Clinical Data Science, and Modeling Utilities

Version 0.3.4672.2026.09.17.10.39

Description An open-source R package developed for Temple University's Center for Biostatistics and Epidemiology (CBE). Provides domain-agnostic biostatistical testing, data quality and missingness visualizations, normalization functions, custom Tidymodels recipe steps, and clinical survival metrics.

License GPL-3 | MIT + file LICENSE

Encoding UTF-8

URL <https://github.com/jkylearmstrong/TempleCBE>, <https://jkylearmstrong.github.io/TempleCBE/>

BugReports <https://github.com/jkylearmstrong/TempleCBE/issues>

Depends R (>= 4.2.0)

Imports broom,
corrplot,
dplyr (>= 1.1.0),
exact2x2,
fs,
generics,
ggplot2 (>= 3.5.0),
ggribes (>= 0.5.0),
ggsurvfit,
glue,
gtsummary,
hardhat (>= 1.3.0),
here,
grDevices,
grid,
igraph,
janitor,
labelled,
methods,
nortest,
openssl,
pdftools (>= 2.0),
purrr (>= 1.0.0),
r2rtf,

readxl,
 recipes ($\geq 1.0.0$),
 rlang ($\geq 1.1.0$),
 scales ($\geq 0.5.0$),
 stats,
 stringr,
 survival ($\geq 3.5.0$),
 tibble ($\geq 3.2.0$),
 tidygraph,
 tidyr ($\geq 1.3.0$),
 tools,
 utils,
 vctrs ($\geq 0.6.0$),
 writexl,
 xml2

Suggests arsenal,
 baguette ($\geq 1.1.0$),
 callr,
 caret,
 DALEX ($\geq 2.5.0$),
 DALEXtra ($\geq 2.3.0$),
 devtools,
 dials,
 DiagrammeR,
 embed ($\geq 1.1.0$),
 FactoMineR,
 furrr ($\geq 0.2.0$),
 future,
 ggpubr,
 ggraph,
 ggVennDiagram,
 glmnet ($\geq 4.1.0$),
 gt,
 haven,
 htmlwidgets,
 jsonlite,
 knitr,
 missForest,
 missRanger ($\geq 2.4.0$),
 openxlsx,
 parsnip,
 pkgload,
 probably ($\geq 1.2.0$),
 profvis,
 quarto (≥ 1.4),
 reticulate,
 rmarkdown,
 rsample ($\geq 1.1.0$),
 stacks ($\geq 1.1.1$),
 survex ($\geq 1.2.0$),
 testthat ($\geq 3.1.7$),
 tidymodels ($\geq 1.1.0$),

tune ($\geq$ 1.1.0),
 visNetwork,
 withr ($\geq$ 2.3.0),
 workflows,
 workflowsets ($\geq$ 1.1.0),
 yaml,
 yardstick ($\geq$ 1.3.0),
 zip ($\geq$ 2.3.0)

VignetteBuilder knitr

Config/testthat/edition 3

Config/roxygen2/version 8.1.0

Contents

add_graf_weights	7
add_missing	8
anonymize_pi	9
apply_database_metadata	11
as_igraph.cbe_database_relationships	12
as_pipeline_graph	12
as_tbl_graph.list	13
audit_report_deliverables	13
autoplot.cbe_compare_df	14
autoplot.cbe_database_relationships	14
autoplot.cbe_key_integrity	15
autoplot.cbe_loco_mp_coxnet	15
autoplot.cv_joint_model	16
build_manual_versioned	16
calculate_fences	17
cbe_balloon_plot	17
cbe_bar_plot	18
cbe_compare_df	19
cbe_contingency_plot	20
cbe_cox_check	22
cbe_cox_multi	23
cbe_cox_single	24
cbe_cox_table	25
cbe_database_relationships	26
cbe_database_venn	27
cbe_docx_review_extract	28
cbe_exact2x2	29
cbe_exact2x2_ci	31
cbe_explain	32
cbe_explain_survival	33
cbe_factor_reference	34
cbe_four_quadrant_report	35
cbe_heatmap_plot	36
cbe_km_single	37
cbe_labels	38
cbe_loco_mp_coxnet	39
cbe_mosaic_plot	41

cbe_pairwise_combos	41
cbe_palette	42
cbe_predict_parts_shap	42
cbe_sas_macro_dir	43
cbe_sas_macro_path	43
cbe_square_plot	44
cbe_survex_loss_ibs	45
cbe_test_categorical	46
cbe_theme_survival	48
cbe_variable_roles	48
censoring_km	50
check_docx_toolchain	51
clean_names	51
collapse_by_stage	52
compute_graph	52
controls_datasets	53
convert_pdfs_to_docx	53
convert_pdf_to_docx	54
copy_data_manifest	55
correlation_diff	55
correlation_diff_heatmap	56
correlation_plot	57
correlation_plot_split	58
corr_test_all	60
coxnet	61
create_qmd_renderer	62
create_report	63
create_toc_from_sas_pdf	64
cv_coxnet	65
cv_coxnet-methods	69
cv_joint_model	70
database_metadata_io	71
delete_nul_files	71
demo_cbe_mapping	72
detect_outliers	72
distribution_plot	73
distribution_test	73
export_interactive_pipeline	74
export_subgraph	74
extract_all_xlsx_tokens	75
extract_surv_components	76
extract_win_posix_paths	76
extract_yaml_formats	77
features_percent_miss	77
FileOutputs	78
FileOutputs-class	79
FilePath	79
FilePath-class	80
FileUses	81
FileUses-class	82
file_meta_fs	82
find_code	83

find_correlation	84
find_python	85
find_sas	86
find_section_file	86
find_soffice	87
fit_n_rank	87
flag_outliers	88
fmt_hr	89
fmt_num	89
fmt_p	90
fmt_pct	90
fmt_sig	91
generate_last_names	91
generate_pi_names	92
generate_pseudonym_token	93
get_dataset_info	94
get_model_parameters	95
get_render_plan	96
glmnet_IBS	96
graf_weights	98
infix_helpers	99
install.packages.no_lock	100
is.int	101
is_normal	101
is_poisson	102
joint_model	103
join_pipelines	104
keep_only	105
km_summary_to_prism	106
list_stage_data_artifacts	107
make_excel_names	107
manhattan_plot	108
min_max_norm	109
missforest_impute_by_mtry	109
missforest_oob_by_mtry	110
missforest_sweep_mtry	111
missmap	113
missranger_max_mtry	114
missranger_oob_by_mtry	115
missranger_sweep_mtry	117
multiple_t_test	119
my_summary_table	120
nested_cv_coxnet	121
nested_cv_joint_model	123
normalize_safely	124
one_vs_rest_t_test	125
package_deliverables	126
parse_here_call_vec	126
pca_biplot	127
pca_eqns	128
pca_feature_loading_heatmap	129
pca_loading_diff	129

pca_loading_diff_heatmap	130
pca_percent_var_explained	131
pca_plot	131
pca_scree_plot	132
pca_variables_plot	133
pdf_to_rtf	134
pformat	135
pipeline_config	135
pipeline_summary	136
pi_anonymizer	136
plot.features_percent_miss	137
plot_categorical_associations	137
plot_cox_forest	138
plot_cox_forest_multi	140
plot_cox_marginal	141
plot_cox_survival	141
plot_dynamic_trajectory	142
plot_group_comparison	143
plot_missingness	144
plot_pca_bi	144
plot_survival_km	145
predict.coxnet_model	145
predict.joint_model	146
print.cbe_cox	147
print.cbe_cox_check	148
print.cbe_cox_multi	148
print.cbe_km	148
print.joint_model	149
print_pipeline	149
proc_pca	150
profvis_summary	151
range_norm	152
read_database_from_excel	152
read_data_manifest	153
read_excel_multiple_headers	154
read_mapped_section_data	154
read_raw_table	155
read_search	155
read_workbook	156
render	157
rossi_data	159
rotation_matrix	160
run_sas_script	160
R_names	161
scale_color_cbe	162
scale_colour_temple	162
scale_fill_cbe	163
scan_data_io	163
sd.error	165
significance_stars	165
simulate_section_data	166
single_t_test	166

step_famd	167
step_lencode_coxnet	169
step_lencode_joint_model	171
stop_if_invalid_manifest	172
summarize_section_by_time	172
summarize_tune_results	173
summary.tbl_graph	174
SumNa	174
surv_subject_truth	175
table_two_by_two	176
temple_brand_path	177
temple_colors	177
temple_pal	178
theme_cbe	178
theme_cbe_deck	179
theme_temple	179
tidy.cbe_loco_mp_coxnet	180
tidy.coxnet_model	180
tidy.joint_model	181
tidy_tmerge_cox	181
tumor_long	182
tumor_wide	183
tune_over_alpha	184
use_temple_brand	185
validate_column_mapping	186
validate_data_manifest	186
visualize_pipeline	187
visualize_pipeline_interactive	188
volcano_plot	188
words	189
write_database_to_excel	190
write_search	191
zip_render	192
zip_reports	193
z_norm	194
Index	195

add_graf_weights

*Add Censoring Weights to Survival Predictions for yardstick***Description**

Fills the `weight_censored` column of each element of a `pred` list-column with `[graf_weights()]`, so the data can go straight to `yardstick::brier_survival()`, `brier_survival_integrated()`, or `roc_auc_survival()`. Use it to score any model's per-subject survival predictions, including models fit to start/stop data.

Usage

```
add_graf_weights(
  data,
  truth = ".truth",
  estimate = ".pred",
  censoring,
  trunc = 0.05
)
```

Arguments

data	A data frame with one row per subject.
truth	Column of 'data' holding right-censored 'Surv' truth.
estimate	List-column of data frames with 'eval_time' and 'pred_survival', one data frame per row, all with the same evaluation times.
censoring	A [censoring_km()] object, estimated on the training data.
trunc	Lower bound for G , capping weights at '1 / trunc' (default 0.05, as in tidymodels' 'parsnip').

Value

'data', with 'weight_censored' added to (or replaced in) every data frame in 'estimate'.

See Also

[surv_subject_truth()], [censoring_km()]

Examples

```
if (requireNamespace("survival", quietly = TRUE) &&
    requireNamespace("yardstick", quietly = TRUE)) {
  train <- survival::Surv(c(2, 4, 6, 8), c(0, 1, 0, 1))
  scored <- tibble::tibble(
    .truth = survival::Surv(c(5, 7), c(1, 0)),
    .pred = list(
      tibble::tibble(eval_time = c(3, 6), .pred_survival = c(0.9, 0.5)),
      tibble::tibble(eval_time = c(3, 6), .pred_survival = c(0.8, 0.6))
    )
  )
  scored <- add_graf_weights(scored, censoring = censoring_km(train))
  yardstick::brier_survival(scored, truth = .truth, .pred)
}
```

add_missing

Add Missing Values Completely at Random

Description

Sets a fixed share of cells in each selected column to NA, choosing the rows independently for each column (missing completely at random). Use it to test imputation: start from complete data, add missingness, impute, and compare the imputed cells with the originals.

Usage

```
add_missing(data, cols = dplyr::everything(), pct_na = 0.1)
```

Arguments

<code>data</code>	A data frame or tibble.
<code>cols</code>	Columns to add missing values to, as a tidyselect expression (default <code>dplyr::everything()</code>).
<code>pct_na</code>	Proportion of cells to set to NA in each selected column, between 0 and 1 (default 0.1). Can be a single number applied to all selected columns, or a numeric vector with length equal to the number of selected columns specifying column-specific amputation rates.

Details

Every selected column must start with no missing values, so afterwards its share of NAs equals `pct_na` (to the nearest whole row), matching `PctNa` from [features_percent_miss](#). The observed share is $1 - \text{pct_na}$, so only `pct_na` is given.

Value

data with missing values added. Its `"missing_cells"` attribute is a tibble with one row per masked cell, giving its row number and feature (column name). Most dplyr verbs drop this attribute, so read it before transforming the result.

Examples

```
set.seed(1)
amputed <- add_missing(mtcars, c(mpg, hp), pct_na = 0.25)
features_percent_miss(amputed)
head(attr(amputed, "missing_cells"))

# Column-specific missingness proportions with a vector:
p_vec <- c(mpg = 0.1, hp = 0.3)
amputed2 <- add_missing(mtcars, c(mpg, hp), pct_na = p_vec)
features_percent_miss(amputed2)
```

anonymize_pi

Anonymize Investigator Names (Vectorized)

Description

Anonymizes investigator names into pseudonyms. Supports both single scalar names and character vectors (e.g. data frame columns like `df$investigator`). Offers both stateless keyed-HMAC hashing and stateful mapping table persistence.

Usage

```
anonymize_pi(
  name,
  method = c("hmac_token", "hmac_surname", "token", "synthetic"),
  key = Sys.getenv("TEMPLECBE_SECRET_KEY", unset = ""),
  secrets_path = NULL,
  exclude = character(),
  prefix = "PI_",
  n_chars = 16,
  rename_fn = NULL,
  seed = NULL
)
```

Arguments

name	Character vector of investigator names or surnames to mask.
method	Character string specifying anonymization approach: "hmac_token", "hmac_surname", "token", or "synthetic".
key	Character string, secret key / pepper for HMAC methods. Defaults to 'Sys.getenv("TEMPLECBE_SECRET_KEY", unset = "")'.
secrets_path	Path to confidential mapping file when 'method = "token"'. Defaults to 'tools::R_user_dir("TempleCBE", which = "data")/pi_mapping.json'.
exclude	Character vector of names that must never be generated (passed to [generate_last_names()] when method is "hmac_surname" or "synthetic").
prefix	Character string, prefix for tokens (default "PI_").
n_chars	Integer, number of hex characters for tokens (default 8). Can be shortened for short names.
rename_fn	Optional function to rename or format generated tokens/pseudonyms.
seed	Optional integer, random seed when method is "synthetic".

Details

Methods:

- "hmac_token" (default): Stateless deterministic token (e.g., 'PI_5c1ebfd9') computed via 'openssl::sha256(name, key = key)'. No mapping file needed on disk.
- "hmac_surname": Stateless deterministic procedural surname generated via keyed HMAC seed and [generate_last_names()]. Consistent across runs without disk state.
- "token": Stateful cryptographic tokens persisted to a user-level JSON mapping file ('tools::R_user_dir("TempleCBE", which = "data")/pi_mapping.json'). Explicitly refuses to write into the Git repository tree to prevent accidental commits of PHI.
- "synthetic": Random procedural surnames via [generate_last_names()].

Value

Character vector of anonymized tokens or pseudonyms, matching the length of 'name'.

Examples

```
# Vectorized token mapping (default method = "hmac_token")
anonymize_pi(c("Franklin", "Taylor", "Patel"), key = "study_salt")

# Short tokens (4 characters)
anonymize_pi(c("Franklin", "Taylor"), n_chars = 4, key = "study_salt")

# Custom renaming function
anonymize_pi(c("Franklin", "Taylor"), rename_fn = tolower, key = "study_salt")

# Vectorized realistic pseudonyms
anonymize_pi(c("Franklin", "Taylor"), method = "hmac_surname", key = "study_salt")
```

apply_database_metadata

Apply Metadata and Variable Roles to an R Database

Description

Synchronizes updated metadata (such as variable labels, dataset descriptions, database names, and variable roles) back into a database list or data frame.

Usage

```
apply_database_metadata(db, metadata)
```

Arguments

db	A named list of data frames (an R database) or a single data frame.
metadata	A data frame/tibble of metadata containing columns columns and optionally dataset_name, labels, dataset_label, database_name, role, X_var, Y_var, ID_var, Time_var.

Value

The modified database or data frame with attributes and labels updated.

Examples

```
db <- list(demo = data.frame(id = 1:2, age = c(20, 30)))
meta <- get_dataset_info(db)
meta$labels[meta$columns == "age"] <- "Age at Enrollment (years)"
db <- apply_database_metadata(db, meta)
```

```
as_igraph.cbe_database_relationships
```

Convert to an enriched igraph graph with rich metadata

Description

Convert to an enriched igraph graph with rich metadata

Usage

```
## S3 method for class 'cbe_database_relationships'
as_igraph(x, ...)

as_igraph(x, ...)

## Default S3 method:
as_igraph(x, stage_colors = NULL, ...)
```

Arguments

x	list of FilePath, FileUses, and FileOutputs objects, or another object with an as_igraph method (e.g. cbe_database_relationships).
...	Additional arguments passed to methods.
stage_colors	Optional named list of stage colors. If NULL, uses pipeline_config() options.

Value

an igraph object containing complete node and edge attributes

```
as_pipeline_graph
```

Convert pipeline objects or igraph to a tidygraph tbl_graph

Description

Enables native dplyr verbs (activate, filter, mutate, group_by, summarise) on computational graphs.

Usage

```
as_pipeline_graph(x, ...)
```

Arguments

x	A list of FilePath/FileUses/FileOutputs objects, or an igraph object.
...	Additional arguments passed to tidygraph::as_tbl_graph.

Value

A tbl_graph object.

as_tbl_graph.list	<i>Convert pipeline list to tidygraph</i>
-------------------	---

Description

Convert pipeline list to tidygraph

Usage

```
## S3 method for class 'list'
as_tbl_graph(x, ...)
```

Arguments

x	A list of FilePath/FileUses/FileOutputs objects.
...	Additional arguments passed to as_pipeline_graph .

Value

A tbl_graph object.

audit_report_deliverables	<i>Audit Deliverable File Tokens in Analysis Scripts</i>
---------------------------	--

Description

Scans source Quarto and R scripts to extract read/write file references and verify whether declared data and report deliverables exist on disk.

Usage

```
audit_report_deliverables(
  analysis_path = ".",
  file_pattern = "\\.(xlsx|rds|pdf|docx|html)$"
)
```

Arguments

analysis_path	Directory containing Quarto (.qmd) and R (.R) scripts.
file_pattern	Regular expression matching deliverable extensions (default: "\\.(xlsx rds pdf docx html)\$").

Value

A tibble summarizing referenced deliverable files, source lines, and on-disk existence.

```
autoplot.cbe_compare_df
```

Autoplot Method for CBE Data Frame Comparison

Description

Generates ggplot2 visualizations of data frame comparisons: 2-set Venn diagrams of observation/key overlap or variable concordance (leveraging **ggVennDiagram** when installed), or discrepancy bar charts across variables.

Usage

```
## S3 method for class 'cbe_compare_df'
autoplot(object, type = c("observations", "variables", "discrepancies"), ...)
```

Arguments

object	A cbe_compare_df object.
type	Character string specifying the plot type: "observations" (default; Venn of rows/keys), "variables" (Venn of common/unique columns), or "discrepancies" (bar chart of value differences).
...	Additional arguments passed to methods or ggVennDiagram .

Value

A [ggplot](#) object.

```
autoplot.cbe_database_relationships
```

Autoplot Method for Database Relationships

Description

Autoplot Method for Database Relationships

Usage

```
## S3 method for class 'cbe_database_relationships'
autoplot(object, type = c("graph", "venn"), db = NULL, id_col = NULL, ...)
```

Arguments

object	A cbe_database_relationships object.
type	Plot type: "graph" for network ER graph, or "venn" for key overlap.
db	Optional parent database list required when type = "venn".
id_col	Identifier column for Venn overlap (default: first shared key).
...	Additional arguments passed to plotting backends.

Value

A ggplot2 or plot object.

`autoplot.cbe_key_integrity`*Autoplot Method for Key Integrity*

Description

Autoplot Method for Key Integrity

Usage

```
## S3 method for class 'cbe_key_integrity'  
autoplot(object, ...)
```

Arguments

<code>object</code>	A cbe_check_key_integrity object.
<code>...</code>	Additional arguments passed to methods.

Value

A ggplot2 object.

`autoplot.cbe_loco_mp_coxnet`*Autoplot Method for LOCO-MP Feature Importance*

Description

Generates a ggplot2 forest/lollipop chart displaying Leave-One-Covariate-Out importance scores, confidence intervals, and significance thresholds.

Usage

```
## S3 method for class 'cbe_loco_mp_coxnet'  
autoplot(object, ...)
```

Arguments

<code>object</code>	A <code>cbe_loco_mp_coxnet</code> object.
<code>...</code>	Additional arguments.

Value

A ggplot2 plot object.

autoplot.cv_joint_model

Autoplot Method for Cross-Validated Joint Models

Description

Visualizes comparative predictive performance across paradigms, plotting the distribution of Integrated Brier Scores and Concordance.

Usage

```
## S3 method for class 'cv_joint_model'
autoplot(object, ...)
```

Arguments

object	A cv_joint_model object.
...	Additional arguments.

Value

A ggplot2 object.

build_manual_versioned

Build a Versioned PDF Reference Manual

Description

Renders the package's Rd documentation to a PDF manual (via R CMD Rd2pdf) named with the current package version, so successive builds never collide or overwrite each other. Writes into pkgdown/assets/manual/ by default: pkgdown copies pkgdown/assets/* verbatim into docs/ on every pkgdown::build_site(), which is what makes a navbar link to the manual resolve on the published site. (inst/manual/ would not be copied there, and would also ship with every installed copy of the package.)

Usage

```
build_manual_versioned(pkg = ".", path = NULL, latest = TRUE)
```

Arguments

pkg	Path to the package, passed to as.package .
path	Directory to write the PDF into. Defaults to pkgdown/assets/manual/ under the package root.
latest	Logical (default TRUE); if TRUE, also writes a stable <package>_latest.pdf copy alongside the versioned file, so a permanent link (e.g. in a pkgdown navbar) doesn't need to change on every version bump.

Value

The path to the versioned PDF, invisibly.

Examples

```
## Not run:
build_manual_versioned()

## End(Not run)
```

calculate_fences	<i>Calculate Inner and Outer IQR Fences</i>
------------------	---

Description

Calculate Inner and Outer IQR Fences

Usage

```
calculate_fences(col)
```

Arguments

col A numeric vector or column.

Value

A one-row tibble with columns lower_inner_fence, upper_inner_fence (at 1.5 x IQR — conventional "mild" outlier boundary) and lower_outer_fence, upper_outer_fence (at 3 x IQR — "extreme" outlier boundary).

Examples

```
calculate_fences(c(1, 2, 3, 4, 5, 100))
```

cbe_balloon_plot	<i>Balloon Plot for 2-Way Contingency Tables</i>
------------------	--

Description

Dedicated wrapper for `cbe_contingency_plot` with type = "balloon".

Usage

```
cbe_balloon_plot(  
  data,  
  var1 = NULL,  
  var2 = NULL,  
  label1 = NULL,  
  label2 = NULL,  
  engine = c("ggplot", "corrplot", "ggpubr"),  
  size_range = c(6, 22),  
  fill_colors = c("#F4F4F4", "#9D2235"),  
  show_counts = TRUE,  
  title = NULL,  
  ...  
)
```

Arguments

data	A data frame containing the categorical variables, or a 2-way table/matrix.
var1	Character string of the row variable column name (ignored if data is a table).
var2	Character string of the column variable column name (ignored if data is a table).
label1	Optional display label for var1 (defaults to var1).
label2	Optional display label for var2 (defaults to var2).
engine	Engine for balloon/square rendering: "ggplot" (default), "corrplot", or "ggpubr".
size_range	Numeric vector of length 2 giving min and max glyph sizes for balloon and square plots (default c(6, 22)).
fill_colors	Character vector of colors for the palette or gradient. Defaults to Temple University brand palette (c("#F4F4F4", "#9D2235")).
show_counts	Logical (default TRUE); whether to print cell counts inside plots.
title	Plot title. If NULL, defaults to "{label1} \u00d7 {label2}".
...	Additional arguments passed to underlying plotting functions.

Value

A ggplot object or invisibly the contingency table.

cbe_bar_plot	<i>Bar Chart for 2-Way Contingency Tables</i>
--------------	---

Description

Dedicated wrapper for [cbe_contingency_plot](#) with type = "bar".

Usage

```
cbe_bar_plot(  
  data,  
  var1 = NULL,  
  var2 = NULL,  
  label1 = NULL,  
  label2 = NULL,  
  bar_position = c("fill", "dodge", "stack"),  
  show_counts = TRUE,  
  title = NULL,  
  ...  
)
```

Arguments

data	A data frame containing the categorical variables, or a 2-way table/matrix.
var1	Character string of the row variable column name (ignored if data is a table).
var2	Character string of the column variable column name (ignored if data is a table).
label1	Optional display label for var1 (defaults to var1).
label2	Optional display label for var2 (defaults to var2).
bar_position	Position adjustment when type = "bar": "fill" (100% proportional stacked bar, default), "dodge" (side-by-side grouped), or "stack" (count stacked).
show_counts	Logical (default TRUE); whether to print cell counts inside plots.
title	Plot title. If NULL, defaults to "{label1} \u00d7 {label2}".
...	Additional arguments passed to underlying plotting functions.

Value

A ggplot object.

cbe_compare_df	<i>Compare Two Data Frames (SAS PROC COMPARE Parity)</i>
----------------	--

Description

Compares two data frames (or tibbles) at the dataset, variable, and observation levels, providing an audit-ready, tidy alternative to SAS PROC COMPARE and legacy comparison utilities (such as `arsenal::comparedf`). Evaluates variable presence, data types, variable labels (from **labelled**), row counts, and value-level discrepancies within a user-specified numerical tolerance.

Usage

```
cbe_compare_df(  
  base,  
  compare,  
  by = NULL,  
  tolerance = 1e-07,  
  base_name = NULL,  
  compare_name = NULL,  
  max_diffs = 100  
)
```

Arguments

base	The base data frame or tibble (equivalent to SAS base=).
compare	The comparison data frame or tibble (equivalent to SAS compare=).
by	Optional character vector of column names specifying key variables used to align observations across datasets (equivalent to SAS id or by statement). If NULL (default), observations are compared by row order.
tolerance	Non-negative numeric threshold for numeric differences (default: 1e-7). Differences with absolute magnitude less than or equal to tolerance are considered matches.
base_name	Optional character string identifying the base dataset. If NULL, deparsed from the base argument.
compare_name	Optional character string identifying the comparison dataset. If NULL, deparsed from the compare argument.
max_diffs	Maximum number of discrepant rows to store per variable (default: 100).

Value

An S3 object of class "cbe_compare_df" containing:

meta	List containing metadata on dataset names, dimensions, tolerance, and keys.
variables	List with elements common, base_only, and compare_only.
summary	Tibble summarizing comparison status for each variable (types, match status, difference counts, max difference, RMSE).
observations	List summarizing observation matching (counts, keys matched, unmatched).
diffs	Tibble containing detailed cell-by-cell discrepancies for all variables exceeding tolerance.
is_concordant	Logical indicating whether the datasets match completely within tolerance on all common variables and observations.

Examples

```
df1 <- data.frame(id = 1:5, x = c(1, 2, 3, 4, 5), y = c("a", "b", "c", "d", "e"))
df2 <- data.frame(id = 1:5, x = c(1, 2, 3.0001, 4, 5), y = c("a", "b", "c", "d", "f"))
cmp <- cbe_compare_df(df1, df2, by = "id", tolerance = 1e-3)
print(cmp)
generics::tidy(cmp)
```

cbe_contingency_plot *Visualizations for 2-Way Contingency Tables and Categorical Tests*

Description

Generates publication-ready clinical visualizations for 2-way contingency tables with hypothesis test statistics (Central Fisher's Exact Test with mid-p adjustment when zero-cells occur, or Pearson Chi-squared test) automatically formatted in the caption.

Usage

```
cbe_contingency_plot(
  data,
  var1 = NULL,
  var2 = NULL,
  label1 = NULL,
  label2 = NULL,
  type = c("balloon", "bar", "mosaic", "heatmap", "square", "corrplot"),
  bar_position = c("fill", "dodge", "stack"),
  engine = c("ggplot", "corrplot", "ggpubr"),
  test = c("auto", "exact", "chisq", "fisher"),
  correct = FALSE,
  size_range = c(6, 22),
  fill_colors = c("#F4F4F4", "#9D2235"),
  show_counts = TRUE,
  title = NULL,
  caption = NULL,
  p_format = NULL,
  ...
)
```

Arguments

<code>data</code>	A data frame containing the categorical variables, or a 2-way table/matrix.
<code>var1</code>	Character string of the row variable column name (ignored if data is a table).
<code>var2</code>	Character string of the column variable column name (ignored if data is a table).
<code>label1</code>	Optional display label for var1 (defaults to var1).
<code>label2</code>	Optional display label for var2 (defaults to var2).
<code>type</code>	Visualization type: "balloon" (default), "bar", "mosaic", "heatmap", "square", or "corrplot".
<code>bar_position</code>	Position adjustment when <code>type = "bar"</code> : "fill" (100% proportional stacked bar, default), "dodge" (side-by-side grouped), or "stack" (count stacked).
<code>engine</code>	Engine for balloon/square rendering: "ggplot" (default), "corrplot", or "ggpubr".
<code>test</code>	Hypothesis test engine: "auto" (default CBE hierarchy), "exact" (force Central Fisher exact test for 2x2 or simulated Fisher for RxC), "chisq" (force Pearson's Chi-squared test via chisq.test), or "fisher" (force fisher.test).
<code>correct</code>	Logical; whether to apply continuity correction when <code>test = "chisq"</code> (default FALSE following CBE standard protocol).
<code>size_range</code>	Numeric vector of length 2 giving min and max glyph sizes for balloon and square plots (default <code>c(6, 22)</code>).
<code>fill_colors</code>	Character vector of colors for the palette or gradient. Defaults to Temple University brand palette (<code>c("#F4F4F4", "#9D2235")</code>).
<code>show_counts</code>	Logical (default TRUE); whether to print cell counts inside plots.
<code>title</code>	Plot title. If NULL, defaults to "{label1} \u00d7 {label2}".
<code>caption</code>	Optional plot caption. If NULL, automatically computed and formatted using pformat with test description.
<code>p_format</code>	Optional p-value formatting function (default pformat).
<code>...</code>	Additional arguments passed to underlying plotting functions.

Details

Supported visualization types:

- "balloon" (default): Proportional bubble/circle plot with counts displayed inside circles.
- "bar": Grouped, stacked, or 100% proportional fill bar chart.
- "mosaic": Two-dimensional mosaic plot with tile area proportional to joint frequency.
- "heatmap": Shaded tile heatmap with cell frequencies and overall percentages.
- "square": Standard 4-quadrant report for 2x2 tables (q1 | q2 // q3 | q4 // p = pformat), or proportional square tiles for RxC tables.
- "corrplot": Circular or square matrix plot rendered via **corrplot**.

Value

A ggplot object (for "balloon", "bar", "mosaic", "heatmap", "square"), or invisibly the contingency table (for "corrplot").

Examples

```
df <- data.frame(
  Treatment = factor(c(rep("Active", 20), rep("Control", 20))),
  Outcome   = factor(c(rep("Response", 14), rep("None", 6),
                        rep("Response", 5), rep("None", 15)))
)

# 1. Balloon Plot (Temple Brand Colors)
cbe_contingency_plot(df, "Outcome", "Treatment", type = "balloon")

# 2. Proportional Bar Plot (100% Fill)
cbe_contingency_plot(df, "Outcome", "Treatment", type = "bar", bar_position = "fill")

# 3. Mosaic Plot
cbe_contingency_plot(df, "Outcome", "Treatment", type = "mosaic")

# 4. Heatmap
cbe_contingency_plot(df, "Outcome", "Treatment", type = "heatmap")

# 5. Standard 4-Quadrant Square Plot (q1 | q2 // q3 | q4 // p = pformat)
cbe_contingency_plot(df, "Outcome", "Treatment", type = "square")

# 6. Force Pearson Chi-Square Test
cbe_contingency_plot(df, "Outcome", "Treatment", type = "square", test = "chisq")
```

cbe_cox_check

Tidy Proportional Hazards Diagnostics for Cox Models

Description

Standalone proportional hazards (PH) assumption diagnostics, usable on any fitted `survival::coxph` model, or on a `cbe_cox` or `cbe_cox_multi` object. Wraps `survival::cox.zph()` with a tidy per-term table, a per-term violation flag, and an automated text summary.

Usage

```
cbe_cox_check(fit)
```

Arguments

`fit` A fitted `survival::coxph` model, or a `cbe_cox` or `cbe_cox_multi` object.

Value

An object of class `cbe_cox_check` containing:

- `zph`: The `survival::cox.zph` object (or `NULL` if it could not be computed).
- `zph_table`: Data frame version of `zph$table` (one row per term, plus a `GLOBAL` row for multivariable fits).
- `zph_violated`: Named logical vector, one entry per term (excluding `GLOBAL`), `TRUE` where the assumption is violated ($p < 0.05$). For a univariable fit this is a length-1 named logical vector.
- `zph_text`: Automated summary sentence (a single string for a univariable fit with one term; a two-element character vector, per-term then global, for a multivariable fit).

See Also

`[cbe_cox_single()]`, `[cbe_cox_multi()]`

`cbe_cox_multi`

Multivariable Cox Proportional Hazards Modeling Engine

Description

Fits a multivariable Cox proportional hazards model, validates the proportional hazards assumption for every term (including the multivariate global test), formats coefficients with explicit reference levels grouped by variable, and gathers model fit and convergence metrics.

Usage

```
cbe_cox_multi(
  data,
  formula = NULL,
  outcome = "outcome",
  features = NULL,
  conf_level = 0.95,
  ...
)
```

Arguments

<code>data</code>	A data frame containing survival data and the candidate features.
<code>formula</code>	Optional. A model formula with a <code>survival::Surv()</code> outcome, e.g. <code>Surv(time, status) ~ age + sex</code> . When supplied, outcome and features are ignored.
<code>outcome</code>	Character string naming the <code>Surv</code> object or outcome column in data (e.g., <code>"outcome"</code> or <code>"Surv(survival_time, status)"</code>). Ignored if formula is supplied.

features	Character vector of candidate predictor columns in data. Required if formula is not supplied.
conf_level	Numeric confidence level (default: 0.95).
...	Additional arguments passed to <code>survival::coxph</code> (e.g., <code>id</code> , <code>ties</code> , <code>weights</code>).

Value

An object of class `cbe_cox_multi` containing:

- `model`: The fitted `survival::coxph` object.
- `table`: Tidy coefficient table (HR, log-HR, CI, p-value) with explicit reference rows per factor, grouped by `Variable`.
- `glance`: One-row data frame of model goodness-of-fit metrics.
- `zph`: A `cbe_cox_check` object (includes the multivariate global test).
- `converged`: Logical; whether the fit converged within the iteration limit.
- `n_iterations`: Number of Newton-Raphson iterations used by the fit.
- `features`: Character vector of feature names used.
- `var_labels`: Named character vector mapping feature names to labels.

See Also

[`cbe_cox_single()`], [`cbe_cox_check()`], [`cbe_cox_table()`], [`plot_cox_forest_multi()`]

`cbe_cox_single`

Univariable Cox Proportional Hazards Screening Engine

Description

Fits a univariable Cox proportional hazards model, validates the proportional hazards assumption, formats coefficients with explicit reference levels, generates plain-language clinical interpretations, and gathers model fit metrics.

Usage

```
cbe_cox_single(data, outcome = "outcome", feature, conf_level = 0.95, ...)
```

Arguments

data	A data frame containing survival data and the candidate feature.
outcome	Character string naming the <code>Surv</code> object or outcome column in data (e.g., <code>"outcome"</code> or <code>"Surv(survival_time, status)"</code>).
feature	Character string naming the candidate predictor column in data.
conf_level	Numeric confidence level (default: 0.95).
...	Additional arguments passed to <code>survival::coxph</code> (e.g., <code>weights</code> , <code>ties</code>).

Value

An object of class `cbe_cox` containing:

- `model`: The fitted `survival::coxph` object.
- `zph`: The `survival::cox.zph` proportional hazards test object.
- `zph_table`: Formatted assumption test table.
- `zph_violated`: Logical indicating whether the assumption was violated ($p < 0.05$).
- `zph_text`: Automated sentence summarizing the assumption test.
- `table`: Clean coefficient table with explicit reference level rows for factors.
- `interpretation`: Plain-language automated interpretation of the hazard ratio(s).
- `glance`: One-row data frame of model goodness-of-fit metrics.
- `feature`: Feature name string.
- `var_label`: Human-readable variable label.
- `is_numeric`: Logical flag indicating whether predictor is continuous/numeric.

See Also

[`cbe_cox_multi()`], [`cbe_cox_check()`], [`cbe_km_single()`], [`cbe_cox_table()`], [`plot_cox_forest()`], [`plot_cox_survival()`], [`plot_cox_marginal()`]

<code>cbe_cox_table</code>	<i>Presentation-Ready Cox Coefficient Table</i>
----------------------------	---

Description

Formats the table element of a [`cbe_cox_single()`] or [`cbe_cox_multi()`] result into a presentation table, adding a `log(HR)` column (if not already present) and optional sorting and significance annotation.

Usage

```
cbe_cox_table(x, sort = c("none", "magnitude", "pvalue"), significance = FALSE)
```

Arguments

<code>x</code>	A <code>cbe_cox</code> (from [<code>cbe_cox_single()</code>]) or <code>cbe_cox_multi</code> (from [<code>cbe_cox_multi()</code>]) object.
<code>sort</code>	One of "none" (default; original model order), "magnitude" (sorts by <code>abs(log(HR))</code> descending; reference rows, which are always exactly 0, sort last), or "pvalue" (sorts by p-value ascending; reference rows, which have no p-value, sort last).
<code>significance</code>	Logical; if TRUE, appends a <code>sig</code> column of significance stars (three asterisks for $p < 0.001$, two for $p < 0.01$, one for $p < 0.05$) and bolds the <code>p.value</code> text of significant rows.

Value

A tibble formatted for presentation, with columns `Variable`, `Level`, `Role`, `HR`, `log(HR)`, `95% CI`, `p.value`, and (if `significance = TRUE`) `sig`.

See Also

[cbe_cox_single()], [cbe_cox_multi()]

cbe_database_relationships

Analyze Relationships and Linkages Across Datasets in an R Database

Description

Evaluates common key variables (e.g., subject identifiers, visit times) across multiple datasets in an R database list. Computes cross-table overlap, record coverage, key cardinality (1:1, 1:many, many:many), and identifies orphan records.

Usage

```
cbe_database_relationships(db, id_cols = NULL)

## S3 method for class 'cbe_database_relationships'
as_tbl_graph(x, ...)

cbe_find_shared_keys(db)

cbe_check_key_integrity(
  db,
  id_col,
  master_dataset = NULL,
  compare = FALSE,
  tolerance = 1e-07,
  ...
)
```

Arguments

db	A named list of data frames (an R database).
id_cols	Optional character vector of key/ID column names to evaluate. If NULL, shared keys are automatically detected across datasets.
x	A cbe_database_relationships object.
...	Additional arguments passed to cbe_compare_df .
id_col	Character string specifying the identifier column to check.
master_dataset	Character string naming the primary cohort table. If provided, checks which IDs in other tables are absent from the master cohort table.
compare	Logical; if TRUE and master_dataset is supplied, runs cbe_compare_df on overlapping variables between the master dataset and child datasets, attaching the comparison results.
tolerance	Numeric comparison tolerance passed to cbe_compare_df .

Value

A tibble summarizing dataset pairs, shared key columns, record counts, overlap counts, coverage percentages, and cardinality.

Examples

```
db <- list(
  patients = data.frame(id = 1:5, age = c(20, 30, 40, 50, 60)),
  vitals = data.frame(id = c(1, 1, 2, 2, 3, 6), bp = c(120, 122, 130, 132, 110, 115))
)
cbe_database_relationships(db)
```

cbe_database_venn

Multi-Dataset Key Overlap Venn Diagram

Description

Visualizes shared subject identifiers or key overlap across all tables in an R database list using **ggVennDiagram** (or native ggplot2 fallback).

Usage

```
cbe_database_venn(db, id_col = "id", title = NULL, ...)
```

Arguments

db	A named list of data frames (an R database).
id_col	Character string specifying the identifier column to intersect across datasets.
title	Optional plot title.
...	Additional arguments passed to ggVennDiagram .

Value

A [ggplot](#) object.

Examples

```
## Not run:
db <- list(
  inputs = data.frame(id = 1:5),
  abg = data.frame(id = 1:3),
  survival = data.frame(id = 2:6)
)
cbe_database_venn(db, id_col = "id")

## End(Not run)
```

cbe_docx_review_extract

DOCX Review Extractor & Multi-Reviewer Collaborative Tracker

Description

Extracts comments, tracked changes, and reconstructed paragraph redlines directly from Microsoft Word (‘.docx’) files and maintains a collaborative review workflow in a formatted multi-sheet Excel workbook and UTF-8 CSV files.

Supports multi-reviewer workflows with per-reviewer fork files (‘review_tracker_<reviewer>.xlsx’), live Excel formulas, dropdown data validations, conditional formatting, non-destructive merging across document revisions, and analytical compute graph ordering.

Usage

```
cbe_docx_review_extract(
  input_dir = NULL,
  output_dir = NULL,
  tracker = "review_tracker.xlsx",
  min_revision_length = 2,
  fork_reviewers = c("jka", "darina", "zhao"),
  manifest_path = NULL,
  verbose = FALSE
)
```

Arguments

input_dir	Character path to a directory containing ‘.docx’ files, or to a single ‘.docx’ file. Defaults to finding an edits directory or the current working directory.
output_dir	Character path to the directory where review tracker workbooks and CSV files should be saved. If ‘NULL’ (default), uses ‘<input_dir>/review_extract’ (or ‘<input_dir>’ if it is already named ‘review_extract’).
tracker	Character filename for the master Excel workbook. Defaults to “review_tracker.xlsx”.
min_revision_length	Integer minimum character count required to record a tracked change. Defaults to ‘2’.
fork_reviewers	Character vector of reviewer identifiers who receive their own single-reviewer fork workbooks. Defaults to ‘c("jka", "darina", "zhao")’.
manifest_path	Optional character path to ‘reports_to_render.xlsx’. If ‘NULL’, the function attempts to locate it automatically in standard repository paths.
verbose	Logical indicating whether to print detailed progress messages. Defaults to ‘FALSE’.

Value

An object of class ‘cbe_review_extract’, which is a named list containing:

comments A tibble of merged comment records.

revisions A tibble of merged raw tracked changes.

redlines A tibble of reconstructed paragraph before/after redlines.

documents A tibble of document-level summary metrics and review statuses.

docxwalk A tibble crosswalking reviewed DOCX files to source ‘.qmd’ and rendered outputs.

summary A tibble summarizing recurring comments across documents.

errors A tibble of any processing errors encountered.

paths A named list of generated file paths.

Examples

```
## Not run:
res <- cbe_docx_review_extract("tasks/edits", verbose = TRUE)
print(res)

## End(Not run)
```

cbe_exact2x2

Exact Test for 2x2 Tables with Automatic Zero-Cell Mid-p Default

Description

Performs an exact test for 2x2 contingency tables using **exact2x2**. If any cell count in the 2x2 table is zero ($\min(\text{tab}) == 0$) and `midp` is unspecified, the function defaults to the mid-p version of Central Fisher’s exact test (`midp = TRUE`). This prevents the extreme conservatism and loss of power of standard conditional exact tests on boundary tables. When all cells are non-zero, it defaults to the standard Central Fisher’s exact test (`midp = FALSE`), which guarantees central confidence intervals that invert the two one-sided tests.

Usage

```
cbe_exact2x2(
  x,
  y = NULL,
  data = NULL,
  variable = NULL,
  by = NULL,
  midp = NULL,
  conf.level = 0.95,
  alternative = "two.sided",
  tsmethod = "central",
  ...
)
```

Arguments

<code>x</code>	A 2x2 numeric matrix, table, or a categorical vector. Can also be a data frame if <code>variable</code> and <code>by</code> are supplied.
<code>y</code>	Optional second categorical vector when <code>x</code> is a vector.
<code>data</code>	Optional data frame when used with <code>variable</code> and <code>by</code> .
<code>variable</code>	Character string of the column name to test when called by gtsummary .

<code>by</code>	Character string of the grouping/stratifying column when called by gtsummary .
<code>midp</code>	Logical. If NULL (the default), automatically set to TRUE when any table cell is 0, and FALSE when all cells are positive. Can be explicitly set to TRUE or FALSE to override.
<code>conf.level</code>	Confidence level for the returned confidence interval (default 0.95).
<code>alternative</code>	Alternative hypothesis direction: "two.sided" (default), "greater", or "less".
<code>tsmethod</code>	Two-sided method passed to exact2x2 : "two.sided" (Central Fisher's exact test, default), "minlike", or "blaker".
<code>...</code>	Additional arguments passed to exact2x2 .

Details

The function accepts a 2x2 table/matrix, two categorical vectors `x` and `y`, or (`data`, `variable`, `by`) arguments for direct compatibility as a custom test in **gtsummary**'s [add_p](#).

Value

A tibble with columns:

<code>estimate</code>	Estimated odds ratio (conditional MLE or median unbiased estimate).
<code>p.value</code>	Two-sided or one-sided p-value (mid-p adjusted when <code>midp</code> = TRUE).
<code>conf.low</code>	Lower bound of the confidence interval.
<code>conf.high</code>	Upper bound of the confidence interval.
<code>statistic</code>	Point estimate of the odds ratio (for gtsummary compatibility).
<code>method</code>	Descriptive name of the exact test performed.
<code>alternative</code>	Alternative hypothesis.
<code>midp</code>	Logical flag indicating whether mid-p adjustment was used.
<code>has_zero</code>	Logical flag indicating whether any cell in the table was zero.

Examples

```
# 2x2 table with a zero cell: automatically triggers mid-p adjustment
tab_zero <- matrix(c(0, 10, 5, 15), nrow = 2,
  dimnames = list(c("Treated", "Control"), c("Event", "No Event")))
cbe_exact2x2(tab_zero)

# 2x2 table without zero cells: standard Central Fisher's exact test
tab_nonzero <- matrix(c(12, 8, 5, 15), nrow = 2)
cbe_exact2x2(tab_nonzero)
```

cbe_exact2x2_ci

*Format Exact 2x2 Odds Ratio and Confidence Interval***Description**

Computes the odds ratio and confidence interval from `cbe_exact2x2` and returns a formatted character string (e.g. "0.8 (0.3, 2.1)"). Useful for inline text reporting or table presentation.

Usage

```
cbe_exact2x2_ci(
  data,
  variable = NULL,
  by = NULL,
  digits = 1,
  conf.level = 0.95,
  midp = NULL,
  ...
)
```

Arguments

<code>data</code>	A data frame or a 2x2 table/matrix.
<code>variable</code>	Character string column name (if data is a data frame).
<code>by</code>	Character string grouping column name (if data is a data frame).
<code>digits</code>	Integer number of decimal places for rounding (default 1).
<code>conf.level</code>	Confidence level (default 0.95).
<code>midp</code>	Logical or NULL (default). If NULL, automatically defaults to TRUE if any cell count is zero.
<code>...</code>	Additional arguments passed to <code>cbe_exact2x2</code> .

Value

A character string with the formatted odds ratio and confidence interval.

Examples

```
tab <- matrix(c(0, 10, 5, 15), nrow = 2)
cbe_exact2x2_ci(tab)
```

cbe_explain

*Explain Joint Model Components via DALEX and survex***Description**

Extracts and creates explainers for all components of a `joint_model`: the survival model (via `survex`), the binary status model (via `DALEX`), and the continuous duration time model (via `DALEX`).

Usage

```
cbe_explain(
  model,
  data = NULL,
  type = c("all", "survival", "status", "time"),
  times = NULL,
  label = NULL,
  verbose = FALSE,
  ...
)

cbe_explain_status(model, data = NULL, label = NULL, verbose = FALSE, ...)

cbe_explain_time(model, data = NULL, label = NULL, verbose = FALSE, ...)
```

Arguments

<code>model</code>	A fitted <code>joint_model</code> object.
<code>data</code>	Evaluation data frame. If <code>NULL</code> , training predictors are used.
<code>type</code>	Character indicating which component(s) to explain: "all", "survival", "status", or "time".
<code>times</code>	Numeric vector of evaluation time points for survival explanations.
<code>label</code>	Optional custom label prefix.
<code>verbose</code>	Logical; if <code>TRUE</code> , progress messages are shown. Default is <code>FALSE</code> .
<code>...</code>	Additional arguments passed to the underlying explainers.

Value

If `type = "all"`, a list of class `"cbe_joint_explainer"` containing `$survival`, `$status`, and `$time` explainers. Otherwise, the individual requested explainer.

cbe_explain_survival *Explain Survival Models via survex and DALEX*

Description

Adapts Tidymodels survival models (workflows, parsnip / censored, coxnet_model, and joint_model) to create a survex survival explainer. Standardizes survival predictions into $N \times K$ survival curves, bridging the gap identified in survex issue #98.

Usage

```
cbe_explain_survival(
  model,
  data = NULL,
  y = NULL,
  predict_survival_function = NULL,
  predict_risk_function = NULL,
  times = NULL,
  label = NULL,
  verbose = FALSE,
  ...
)
```

Arguments

model	A fitted model object: joint_model, coxnet_model, cv_coxnet, workflow, or model_fit.
data	A data frame containing predictor columns used to calculate explanations. If NULL and model is a joint_model, the training predictors are used.
y	A survival object or response variable (e.g. Surv(time, status)). If NULL and model is a joint_model, the original survival object is used.
predict_survival_function	Optional custom prediction function returning an $N \times K$ data frame or matrix of survival probabilities $S(t)$.
predict_risk_function	Optional custom prediction function returning a numeric vector of risk scores or cumulative hazard.
times	A numeric vector of evaluation time points. If NULL, defaults to unique event times in y.
label	Character string naming the model in plots and summaries.
verbose	Logical; if TRUE, progress messages are shown. Default is FALSE.
...	Additional arguments passed to survex::explain_survival().

Value

A survex_explainer object compatible with survex::model_parts(), survex::model_performance(), and survex::predict_parts().

Examples

```
## Not run:
if (requireNamespace("survex", quietly = TRUE) && requireNamespace("survival", quietly = TRUE)) {
  lung <- survival::lung
  df <- na.omit(lung[, c("time", "status", "age", "sex")])
  fit <- joint_model(df, survival::Surv(time, status) ~ age + sex)
  expl <- cbe_explain_survival(fit)
  mp <- survex::model_parts(expl)
}

## End(Not run)
```

cbe_factor_reference *Relevel a Factor's Reference Level*

Description

Standalone helper for choosing the reference (baseline) level of a factor before passing it to `[cbe_cox_single()]` or `[cbe_cox_multi()]`. Wraps `[stats::relevel()]`, defaulting to the first level (matching the reference-row convention already used by `cbe_cox_single()/cbe_cox_multi()`), and attaches a "cbe_reference_level" attribute recording the chosen level.

Usage

```
cbe_factor_reference(x, ref_level = NULL)
```

Arguments

<code>x</code>	A factor, or a vector coercible to one via <code>as.factor()</code> .
<code>ref_level</code>	Character string naming the level to use as reference. Defaults to the first level of <code>x</code> (its current or natural ordering).

Details

`cbe_cox_single()` and `cbe_cox_multi()` are not rewired to require this helper; it is an optional convenience for callers who want to choose the reference level explicitly before fitting.

Value

The releveled factor, with attribute "cbe_reference_level" set to the chosen reference level.

See Also

`[cbe_cox_single()]`, `[cbe_cox_multi()]`

cbe_four_quadrant_report

Standard 4-Quadrant Clinical Contingency Report

Description

Generates a standard 4-quadrant clinical report for a 2x2 contingency table:

```
q1 | q2
q3 | q4
p = pformat
```

where q1 is (Row 1, Col 1), q2 is (Row 1, Col 2), q3 is (Row 2, Col 1), and q4 is (Row 2, Col 2). If any cell count is zero, hypothesis testing automatically defaults to the mid-p version of Central Fisher's exact test via [cbe_exact2x2](#).

Usage

```
cbe_four_quadrant_report(
  data,
  var1 = NULL,
  var2 = NULL,
  label1 = NULL,
  label2 = NULL,
  test = c("auto", "exact", "chisq", "fisher"),
  correct = FALSE,
  p_format = NULL,
  ...
)
```

Arguments

data	A data frame or a 2x2 matrix/table.
var1	Character string of row variable column name (if data is a data frame).
var2	Character string of column variable column name (if data is a data frame).
label1	Display label for row variable.
label2	Display label for column variable.
test	Hypothesis test engine: "auto" (default CBE hierarchy: mid-p exact if zero-cell, otherwise Central Fisher's exact), "exact" (force exact test), "chisq" (force Pearson's Chi-squared test via chisq.test), or "fisher" (force fisher.test).
correct	Logical; whether to apply continuity correction for "chisq" (default FALSE).
p_format	Function for p-value formatting (default pformat).
...	Additional arguments passed to cbe_contingency_plot .

Value

A list containing:

quadrants	A tibble with counts and percentages for q1, q2, q3, and q4.
p_value	Numeric exact test p-value.
p_formatted	Formatted p-value string (e.g. "p = 0.024").
test_method	Name of the exact test used.
compact_report	Compact 3-line string formatted as q1 q2 // q3 q4 // p = pformat.
text_card	Console-ready formatted 4-quadrant text card.
plot	A ggplot object rendering the 4-quadrant square report.

Examples

```
tab <- matrix(c(14, 6, 4, 16), nrow = 2,
              dimnames = list(c("Yes", "No"), c("Active", "Placebo")))
rep <- cbe_four_quadrant_report(tab, label1 = "Response", label2 = "Treatment")
cat(rep$compact_report, "\n\n")
cat(rep$text_card)

# Force Pearson Chi-squared test:
rep_chi <- cbe_four_quadrant_report(tab, test = "chisq")
cat(rep_chi$compact_report, "\n")
```

cbe_heatmap_plot	<i>Heatmap Plot for 2-Way Contingency Tables</i>
------------------	--

Description

Dedicated wrapper for [cbe_contingency_plot](#) with type = "heatmap".

Usage

```
cbe_heatmap_plot(
  data,
  var1 = NULL,
  var2 = NULL,
  label1 = NULL,
  label2 = NULL,
  fill_colors = c("#F4F4F4", "#9D2235"),
  show_counts = TRUE,
  title = NULL,
  ...
)
```

Arguments

data	A data frame containing the categorical variables, or a 2-way table/matrix.
var1	Character string of the row variable column name (ignored if data is a table).
var2	Character string of the column variable column name (ignored if data is a table).

label1	Optional display label for var1 (defaults to var1).
label2	Optional display label for var2 (defaults to var2).
fill_colors	Character vector of colors for the palette or gradient. Defaults to Temple University brand palette (c("#F4F4F4", "#9D2235")).
show_counts	Logical (default TRUE); whether to print cell counts inside plots.
title	Plot title. If NULL, defaults to "{label1} \u00d7 {label2}".
...	Additional arguments passed to underlying plotting functions.

Value

A ggplot object.

cbe_km_single	<i>Univariable Kaplan-Meier and Cox Screening Engine</i>
---------------	--

Description

Pairs a univariable [cbe_cox_single()] Cox proportional hazards fit with the matching Kaplan-Meier stratified survival curve for the same predictor, so both are produced from a single call without refitting the Cox model twice.

Usage

```
cbe_km_single(data, outcome = "outcome", feature, conf_level = 0.95)
```

Arguments

data	A data frame containing survival data and the candidate feature.
outcome	Character string naming the Surv object or outcome column in data (e.g., "outcome" or "Surv(survival_time, status)").
feature	Character string naming the candidate predictor column in data. Factor (or coercible) predictors are used directly as Kaplan-Meier strata; numeric predictors are binned into quartiles for the Kaplan-Meier strata (matching [plot_cox_survival()]'s convention), while the Cox side treats them as continuous.
conf_level	Numeric confidence level (default: 0.95).

Value

An object of class cbe_km containing:

- `cox`: The cbe_cox object from [cbe_cox_single()].
- `km_fit`: The `survival::survfit` object for the Kaplan-Meier curve.
- `km_tidy`: `broom::tidy()` table of `km_fit`.
- `direction`: "increases" or "decreases", pulled from the Cox hazard ratio direction.
- `summary`: One-row-per-stratum tidy summary combining Cox hazard ratios and Kaplan-Meier median survival.
- `feature`: Feature name string.
- `var_label`: Human-readable variable label.

See Also

[cbe_cox_single()], [plot_cox_survival()]

cbe_labels	<i>Get or Set Dataset and Database Labels</i>
------------	---

Description

Extends **labelled** variable-level conventions to dataset-level and database-level metadata attributes.

Usage

```
cbe_dataset_label(x)

cbe_dataset_label(x) <- value

cbe_database_name(x)

cbe_database_name(x) <- value

cbe_database_label(x)

cbe_database_label(x) <- value

cbe_set_dataset_labels(x, labels)

cbe_get_dataset_labels(x)
```

Arguments

<code>x</code>	A data frame (for dataset functions) or a named list of data frames (for database functions).
<code>value</code>	A character string specifying the label or name.
<code>labels</code>	A named character vector or list mapping dataset names to dataset labels.

Value

For getters, the label or name string (or NULL if unset). For setters, the modified object `x` with the attribute attached.

Examples

```
df <- mtcars
cbe_dataset_label(df) <- "1974 Motor Trend Car Road Tests"
cbe_dataset_label(df)
```

cbe_loco_mp_coxnet	<i>Leave-One-Covariate-Out Inference with MiniPatch Ensembles (LOCO-MP) for Cox Models</i>
--------------------	--

Description

Implements the LOCO-MP statistical inference framework (Gan, Zheng, & Allen 2022) for penalized Cox proportional hazards regression. Combines random minipatch ensembling (subsampling both observations/subjects and covariates) with out-of-bag Leave-One-Covariate-Out evaluation to produce distribution-free feature importance estimates, standard errors, asymptotic z-tests, and confidence intervals for survival models.

Usage

```
cbe_loco_mp_coxnet(
  formula = NULL,
  data = NULL,
  x = NULL,
  y = NULL,
  B = 100,
  n_ratio = 0.7,
  m_ratio = 0.5,
  penalty = NULL,
  mixture = 1,
  eval_time = NULL,
  subject_id = NULL,
  alpha = 0.05,
  p_adjust = c("bonferroni", "BH", "fdr", "holm", "none"),
  trunc = 0.05,
  parallel = FALSE,
  seed = NULL,
  ...
)
```

Arguments

formula	A formula with a Surv() outcome, e.g. Surv(time, status) ~ x1 + x2.
data	A data frame containing the variables in the model.
x	An optional predictor matrix (used if formula and data are NULL).
y	An optional Surv outcome object (used if formula and data are NULL).
B	Number of minipatches to generate. Default is 100.
n_ratio	Subsampling fraction for subjects/observations without replacement. Default is 0.7.
m_ratio	Subsampling fraction for predictors without replacement. Default is 0.5.
penalty	Penalty parameter for coxnet. If NULL, defaults to 0 (unpenalized).
mixture	Elastic-net mixing parameter ($\alpha \in [0, 1]$). Default is 1 (lasso).
eval_time	Numeric vector of evaluation time points for IPCW Brier loss integration. If NULL, defaults to deciles of event times.

subject_id	Optional character column name identifying subjects for clustered / counting-process start/stop data.
alpha	Significance level for confidence intervals. Default is 0.05.
p_adjust	Multiple testing correction method for p-values. Options include "bonferroni", "BH", "fdr", "holm", or "none". Default is "bonferroni".
trunc	Lower truncation threshold for Kaplan-Meier censoring weights. Default is 0.05.
parallel	Logical; if TRUE and future is available, runs patches in parallel.
seed	Optional random seed for reproducible subsampling.
...	Additional arguments passed to <code>coxnet()</code> .

Value

An S3 object of class "cbe_loco_mp_coxnet" containing:

results	A tibble with columns term, importance, std_error, statistic, p_value, p_adjusted, conf_low, conf_high.
eval_time	Evaluation time grid used for IPCW loss integration.
B	Number of successfully fitted minipatches.
n_ratio, m_ratio	Subsampling ratios used.
alpha	Significance level.
p_adjust	Multiple testing adjustment method.

See Also

[`nested_cv_coxnet()`], [`cv_coxnet()`], [`coxnet()`]

Examples

```
if (requireNamespace("glmnet", quietly = TRUE) && requireNamespace("survival", quietly = TRUE)) {
  set.seed(42)
  n <- 60
  df <- data.frame(
    time = stats::rexp(n, 0.1) + 0.1,
    status = stats::rbinom(n, 1, 0.6),
    x1 = stats::rnorm(n),
    x2 = stats::rnorm(n),
    x3 = stats::rnorm(n)
  )
  fit <- cbe_loco_mp_coxnet(survival::Surv(time, status) ~ x1 + x2 + x3, data = df, B = 20)
  fit
  generics::tidy(fit)
}
```

cbe_mosaic_plot	<i>Mosaic Plot for 2-Way Contingency Tables</i>
-----------------	---

Description

Dedicated wrapper for `cbe_contingency_plot` with `type = "mosaic"`.

Usage

```
cbe_mosaic_plot(
  data,
  var1 = NULL,
  var2 = NULL,
  label1 = NULL,
  label2 = NULL,
  show_counts = TRUE,
  title = NULL,
  ...
)
```

Arguments

<code>data</code>	A data frame containing the categorical variables, or a 2-way table/matrix.
<code>var1</code>	Character string of the row variable column name (ignored if data is a table).
<code>var2</code>	Character string of the column variable column name (ignored if data is a table).
<code>label1</code>	Optional display label for var1 (defaults to var1).
<code>label2</code>	Optional display label for var2 (defaults to var2).
<code>show_counts</code>	Logical (default TRUE); whether to print cell counts inside plots.
<code>title</code>	Plot title. If NULL, defaults to "{label1} \u00d7 {label2}".
<code>...</code>	Additional arguments passed to underlying plotting functions.

Value

A ggplot object.

cbe_pairwise_combos	<i>Generate All Pairwise Categorical Combinations</i>
---------------------	---

Description

Helper for report generation to enumerate all pairwise combinations of categorical variables in a dataset with sequential integer IDs and labels.

Usage

```
cbe_pairwise_combos(data, cols = NULL, labels_df = NULL)
```

Arguments

data	A data frame or tibble.
cols	Optional character vector of column names to include. If NULL, all factor and character columns are included.
labels_df	Optional lookup data frame with columns variable and label.

Value

A tibble with columns id, var1, var2, label1, label2, and comparison_label.

Examples

```
df <- data.frame(  
  A = factor(c("X", "Y")),  
  B = factor(c("1", "2")),  
  C = factor(c("M", "N"))  
)  
cbe_pairwise_combos(df)
```

cbe_palette	<i>Temple CBE Color Palette</i>
-------------	---------------------------------

Description

Official and complementary colors for CBE publications and slide decks.

Usage

```
cbe_palette
```

cbe_predict_parts_shap	<i>SHAP Variable Attributions across Survival and Joint Models</i>
------------------------	--

Description

Computes individual observation SHAP attributions: using time-dependent SurvSHAP(t) via survex::predict_parts(type = "survshap") for survival explainers, and DALEX::predict_parts(type = "shap") for classification/regression explainers.

Usage

```
cbe_predict_parts_shap(  
  explainer,  
  new_observation,  
  type = c("auto", "shap", "survshap"),  
  ...  
)
```

Arguments

explainer	An explainer object produced by <code>cbe_explain_survival()</code> , <code>cbe_explain()</code> , or <code>DALEX::explain()</code> .
new_observation	A 1-row data frame containing the observation to explain.
type	Character specifying the attribution type: "auto" (the default), "shap", or "survshap".
...	Additional arguments passed to <code>survex::predict_parts()</code> or <code>DALEX::predict_parts()</code> .

Value

A `predict_parts` attribution object, printable and plottable with `plot()`.

cbe_sas_macro_dir	<i>Directory of TempleCBE SAS Macros</i>
-------------------	--

Description

Returns the file system path to the directory containing TempleCBE's bundled SAS macros ('cbe_brier_score.sas', 'cbe_cox_phreg.sas', 'cbe_counting_process.sas', and 'cbe_macros.sas').

Usage

```
cbe_sas_macro_dir()
```

Value

Absolute path to the SAS macro directory.

See Also

[`cbe_sas_macro_path()`], [`run_sas_script()`], [`find_sas()`]

Examples

```
try(cbe_sas_macro_dir())
```

cbe_sas_macro_path	<i>Path to a TempleCBE SAS Macro File</i>
--------------------	---

Description

Returns the absolute path to an installed TempleCBE SAS macro or benchmark script.

Usage

```
cbe_sas_macro_path(macro = "cbe_macros.sas")
```

Arguments

macro Filename of the SAS macro or script (default: "cbe_macros.sas"). Available files include "cbe_macros.sas", "cbe_brier_score.sas", "cbe_cox_phreg.sas", "cbe_counting_process.sas", "coxtvc.sas", "cpdata.sas", "benchmark_brier_lung.sas", and "example_85_7.sas".

Value

Absolute path to the requested '.sas' file.

See Also

[cbe_sas_macro_dir()], [run_sas_script()], [find_sas()]

Examples

```
try(cbe_sas_macro_path("cbe_macros.sas"))
```

cbe_square_plot	<i>Square Plot: Standard 4-Quadrant Contingency Report</i>
-----------------	--

Description

Dedicated wrapper for [cbe_contingency_plot](#) with type = "square". For 2x2 contingency tables, renders the standard 4-quadrant report:

```
q1 | q2
q3 | q4
p = pformat
```

If any cell count is zero, hypothesis testing automatically defaults to the mid-p version of Central Fisher’s exact test via [cbe_exact2x2](#).

Usage

```
cbe_square_plot(
  data,
  var1 = NULL,
  var2 = NULL,
  label1 = NULL,
  label2 = NULL,
  test = c("auto", "exact", "chisq", "fisher"),
  correct = FALSE,
  engine = c("ggplot", "corrplot"),
  size_range = c(6, 22),
  fill_colors = c("#F4F4F4", "#9D2235"),
  show_counts = TRUE,
  title = NULL,
  caption = NULL,
  p_format = NULL,
  ...
)
```

Arguments

<code>data</code>	A data frame containing the categorical variables, or a 2-way table/matrix.
<code>var1</code>	Character string of the row variable column name (ignored if data is a table).
<code>var2</code>	Character string of the column variable column name (ignored if data is a table).
<code>label1</code>	Optional display label for var1 (defaults to var1).
<code>label2</code>	Optional display label for var2 (defaults to var2).
<code>test</code>	Hypothesis test engine: "auto" (default CBE hierarchy), "exact" (force Central Fisher exact test for 2x2 or simulated Fisher for RxC), "chisq" (force Pearson's Chi-squared test via chisq.test), or "fisher" (force fisher.test).
<code>correct</code>	Logical; whether to apply continuity correction when test = "chisq" (default FALSE following CBE standard protocol).
<code>engine</code>	Engine for balloon/square rendering: "ggplot" (default), "corrplot", or "ggpubr".
<code>size_range</code>	Numeric vector of length 2 giving min and max glyph sizes for balloon and square plots (default c(6, 22)).
<code>fill_colors</code>	Character vector of colors for the palette or gradient. Defaults to Temple University brand palette (c("#F4F4F4", "#9D2235")).
<code>show_counts</code>	Logical (default TRUE); whether to print cell counts inside plots.
<code>title</code>	Plot title. If NULL, defaults to "{label1} \u00d7 {label2}".
<code>caption</code>	Optional plot caption. If NULL, automatically computed and formatted using pformat with test description.
<code>p_format</code>	Optional p-value formatting function (default pformat).
<code>...</code>	Additional arguments passed to underlying plotting functions.

Value

A ggplot object or invisibly the contingency table.

`cbe_survex_loss_ibs` *TempleCBE Custom Loss Functions for survex*

Description

Provides survival loss functions implementing TempleCBE's inverse-probability-of-censoring weighted (IPCW) Graf Integrated Brier Score and time-dependent Brier Score for use with `survex::model_parts()` and `survex::model_performance()`.

Usage

```
cbe_survex_loss_ibs(normalization = NULL, max_quantile = 1, trunc = 0.05)
```

```
cbe_survex_loss_brier(trunc = 0.05)
```

Arguments

<code>normalization</code>	Optional normalization parameter for survex integral calculation. Can be NULL, "t_max", or "survival".
<code>max_quantile</code>	Upper quantile cutoff for evaluation time window. Default is 1.
<code>trunc</code>	Lower truncation threshold for censoring survival probabilities to prevent inflation by rare tails. Default is 0.05.

Value

A loss function with attributes "loss_type" ("integrated" or "time-dependent") and "loss_name", directly passable to `survex`.

Examples

```
## Not run:
if (requireNamespace("survex", quietly = TRUE)) {
  loss_ibs <- cbe_survex_loss_ibs()
  # mp <- survex::model_parts(explainer, loss_function = loss_ibs)
}

## End(Not run)
```

cbe_test_categorical	<i>Institutional Categorical Hypothesis Test for gtsummary and Batch Testing</i>
----------------------	--

Description

Standard categorical hypothesis testing engine conforming to the CBE statistical protocol. Operates as a drop-in custom test for `add_p` or as a standalone batch testing engine with optional **furrr** parallelization.

Usage

```
cbe_test_categorical(
  data,
  variable = NULL,
  by = NULL,
  test = c("auto", "exact", "chisq", "fisher"),
  correct = FALSE,
  B = 2000L,
  simulate.p.value = NULL,
  parallel = FALSE,
  n_chunks = 4L,
  group = NULL,
  type = NULL,
  test.args = NULL,
  adj.vars = NULL,
  conf.level = NULL,
  tbl = NULL,
  continuous_variable = NULL,
  ...
)
```

Arguments

`data` Data frame supplied directly or by **gtsummary**.

variable	Character string column name for the feature being tested, a character vector of variable names for batch testing, or a two-sided formula of the form <code>response ~ grouping_var</code> or <code>var1 + var2 ~ grouping_var</code> . If <code>NULL</code> , all categorical/factor columns in data (excluding <code>by</code>) are tested.
by	Character string column name for the stratifying/grouping variable. Can be omitted when <code>variable</code> is a formula.
test	Hypothesis test engine: <code>"auto"</code> (default CBE hierarchy), <code>"exact"</code> (force Central Fisher exact test for 2x2 or simulated Fisher for RxC), <code>"chisq"</code> (force Pearson's Chi-squared test via chisq.test), or <code>"fisher"</code> (force fisher.test).
correct	Logical; whether to apply continuity correction when <code>test = "chisq"</code> (default <code>FALSE</code> following CBE standard protocol).
B	Integer; number of Monte Carlo replicates when Fisher's exact test with simulation is triggered (default 2000L).
simulate.p.value	Optional logical; if specified, explicitly enables or disables Monte Carlo simulation for Fisher's test. If <code>NULL</code> (default), auto-triggers simulation when table total exceeds 500 or any dimension exceeds 2.
parallel	Logical; whether to use furrr for parallel execution (for batch testing across multiple variables or chunking high-B simulations). Requires setting a plan beforehand.
n_chunks	Integer; number of chunks to partition B into when running parallel simulations with <code>B >= 10000</code> (default 4L).
group, type, test.args, adj.vars, conf.level, tbl, continuous_variable	Unused; accepted and ignored so this function satisfies gtsummary 's custom-test calling convention (see tests) without those arguments leaking into ... and reaching chisq.test / fisher.test , neither of which accepts extra named arguments.
...	Additional arguments passed to exact2x2 or fisher.test .

Details

The function follows a rigorous 4-rule hierarchy:

1. **2x2 table with zero cell:** Uses Central Fisher's exact test with mid-p adjustment ([exact2x2](#) with `midp = TRUE`) to avoid extreme conservatism.
2. **2x2 table without zero cell:** Uses Central Fisher's exact test ([exact2x2](#) with `midp = FALSE`).
3. **RxC table with sparse counts:** If any expected cell count is less than 5, runs Fisher's exact test with Monte Carlo simulation ([fisher.test](#) with `simulate.p.value = TRUE`).
4. **RxC table with adequate counts:** If all expected cell counts are at least 5, runs Pearson's Chi-squared test without continuity correction ([chisq.test](#) with `correct = FALSE`).

Value

When `variable` is a single string (such as when called by [gtsummary](#)), returns a one-row tibble compliant with [gtsummary](#)'s custom test specification (`p.value`, `statistic`, `parameter`, `method`). When `variable` is a vector or `NULL`, returns a tidy tibble of test results across all evaluated variables.

Examples

```
library(gtsummary)
trial |>
  tbl_summary(by = trt, include = c(response, death, grade)) |>
  add_p(test = all_categorical() ~ cbe_test_categorical) |>
  separate_p_footnotes()

# Standalone batch testing across multiple features:
cbe_test_categorical(trial, variable = c("response", "death", "grade"), by = "trt")

# Formula interface:
cbe_test_categorical(trial, response ~ trt)
```

cbe_theme_survival	<i>CBE ggplot2 Theme for Survival Analysis Visualizations</i>
--------------------	---

Description

A variant of `theme_cbe()` tuned for Cox model and Kaplan-Meier plots (forest plots, predicted/observed survival curves, marginal risk curves), giving a single consistent look across `plot_cox_forest()`, `plot_cox_forest_multi()`, `plot_cox_survival()`, and `plot_cox_marginal()`.

Usage

```
cbe_theme_survival(base_size = 12, base_family = "")
```

Arguments

- `base_size` Base font size (default: 12)
- `base_family` Base font family (default: "")

Value

A ggplot2 theme object

See Also

```
[theme_cbe()], [theme_cbe_deck()]
```

cbe_variable_roles	<i>Define and Manage Clinical Variable Roles</i>
--------------------	--

Description

Assigns functional roles (e.g. predictor, outcome, identifier, time variable, strata, weights) to variables across a single data frame or an entire R database list. Produces standardized schemas compatible with [validate_column_mapping](#), **recipes**, and modeling workflows.

Usage

```

cbe_variable_roles(
  data,
  id = NULL,
  outcome = NULL,
  time = NULL,
  predictors = NULL,
  strata = NULL,
  weight = NULL,
  ignore = NULL,
  roles = NULL
)

cbe_set_roles(data, ...)

cbe_get_roles(data)

cbe_get_predictors(roles)

cbe_get_outcomes(roles)

cbe_get_id_cols(roles)

cbe_get_time_cols(roles)

```

Arguments

<code>data</code>	A data frame or a named list of data frames (an R database).
<code>id</code>	Character vector of column names acting as subject/patient/cluster identifiers.
<code>outcome</code>	Character vector of column names acting as response/outcomes (<code>Y_var</code>).
<code>time</code>	Character vector of column names acting as longitudinal visit/time variables (<code>Time_var</code>).
<code>predictors</code>	Character vector of column names acting as features/predictors (<code>X_var</code>).
<code>strata</code>	Character vector of column names used for stratification or subgroup analysis.
<code>weight</code>	Character vector of column names used for case weights or survey weights.
<code>ignore</code>	Character vector of column names to ignore/exclude from modeling.
<code>roles</code>	Role table (e.g. from cbe_variable_roles) or data object.

Value

A tibble with columns `columns`, `role`, `X_var`, `Y_var`, `ID_var`, `Time_var`, and `dataset_name` (when multi-table).

Examples

```

df <- data.frame(id = 1:5, time = 1:5, sbp = c(120, 130, 115, 140, 125), death = c(0, 0, 1, 0, 1))
roles <- cbe_variable_roles(df, id = "id", time = "time", outcome = "death", predictors = "sbp")

```

censoring_km	<i>Kaplan-Meier Estimate of the Censoring Distribution</i>
--------------	--

Description

Fits the "reverse" Kaplan-Meier estimator $G(t) = P(C > t)$, treating censoring as the event, for inverse-probability-of-censoring weights. Estimate it on the data a model was trained on, and apply it to the data being scored.

Usage

```
censoring_km(truth)
```

```
## S3 method for class 'censoring_km'
predict(object, time, left = FALSE, trunc = 0, ...)
```

Arguments

truth	A right-censored [survival::Surv()] object with one element per subject, such as the <code>‘.truth’</code> column of [surv_subject_truth()].
object	A <code>‘censoring_km’</code> object.
time	Numeric vector of times.
left	If <code>‘TRUE’</code> , the left limit $G(t^-)$, i.e. censoring strictly before <code>‘time’</code> .
trunc	Lower bound applied to G , so weights stay finite.
...	Not used.

Value

An object of class `‘censoring_km’`, with `‘time’` and `‘surv’`; use `‘predict(object, time)’` for $G(t)$.

See Also

[graf_weights()]

Examples

```
if (requireNamespace("survival", quietly = TRUE)) {
  cens <- censoring_km(survival::Surv(c(2, 4, 6, 8), c(0, 1, 0, 1)))
  predict(cens, c(1, 2, 5, 7))
}
```

check_docx_toolchain	<i>Report Which PDF -> DOCX Backends Are Usable on This Machine</i>
----------------------	--

Description

Safe to call at any time; performs discovery only and converts nothing.

Usage

```
check_docx_toolchain(quiet = FALSE)
```

Arguments

quiet Suppress the printed diagnosis.

Value

Invisibly, a list with elements python_with_pdf2docx, python_any, soffice, word_com and backend.

See Also

[convert_pdf_to_docx()], [convert_pdfs_to_docx()]

Examples

```
## Not run:
check_docx_toolchain()

## End(Not run)
```

clean_names	<i>Clean and Standardize Variable Names</i>
-------------	---

Description

Converts variable/column names to lower snake_case and removes invalid special characters.

Usage

```
clean_names(data)
```

Arguments

data A data frame, tibble, or character vector of names.

Value

Cleaned data frame or character vector.

Examples

```
clean_names(c("First Name", "ZIP Code", "Total ($)"))
```

collapse_by_stage	<i>Collapse an igraph pipeline graph to one node per pipeline ‘stage’</i>
-------------------	---

Description

The full per-file compute graph (~90+ nodes) is the right level of detail for the interactive HTML, where a reader can zoom and pan – but it cannot be made legible as a single static image for the PDF/DOCX deliverable at any reasonable page size. This produces a bird’s-eye version instead: one node per ‘stage’, with an edge between two stages whenever any file in the first stage feeds any file in the second (parallel edges collapsed).

Usage

```
collapse_by_stage(g, stage_labels = NULL)
```

Arguments

- g An igraph object (see as_igraph()).
- stage_labels Optional named character vector mapping stage identifiers to display labels.

Value

An igraph object with one node per stage.

compute_graph	<i>Computational Pipeline Dependency Graphs</i>
---------------	---

Description

S4 classes and helper functions for tracking file-level dependencies in an analysis pipeline (raw data, rendering scripts, derived artifacts) and for visualizing, summarizing, and topologically sorting the resulting dependency graph. Language-, format-, and project-agnostic: a node is any file on disk (R, Python, SAS, Quarto, spreadsheets, ...), identified by path rather than tied to a specific tool.

Details

The graph engine (**igraph/tidygraph**) is a hard dependency. Three rendering backends are optional (Suggests) and only required by the functions that use them: static PNG export ([print_pipeline](#)) needs **ggraph**; the interactive Graphviz view ([visualize_pipeline](#)) needs **DiagrammeR**; and the interactive HTML export ([visualize_pipeline_interactive](#), [export_interactive_pipeline](#), [export_subgraph](#)) needs **visNetwork** and **htmlwidgets**.

controls_datasets	<i>Check if a FileOutputs object controls datasets (outputs an rds, xlsx, etc. file)</i>
-------------------	--

Description

Check if a FileOutputs object controls datasets (outputs an rds, xlsx, etc. file)

Usage

```
controls_datasets(obj)
```

Arguments

obj A FileOutputs object.

Value

Logical.

convert_pdfs_to_docx	<i>Convert PDFs to DOCX Using the Best Available Backend</i>
----------------------	--

Description

Backend preference is python (pdf2docx) -> LibreOffice -> Word COM (Windows only). When the python backend fails on an individual file and LibreOffice is present, that file is retried once with LibreOffice.

Usage

```
convert_pdfs_to_docx(
  conversions,
  backend = "auto",
  strict = TRUE,
  timeout = 600
)
```

Arguments

conversions	A data frame with character columns src and dest, and optionally temp_dest, a second location each converted file is copied to.
backend	One of "auto", "python", "libreoffice", "word_com".
strict	If TRUE (default), error with an actionable message when no backend is available. Set FALSE to warn and skip.
timeout	Seconds before the Word COM backend gives up on a file.

Value

conversions with an added logical converted column.

See Also

[convert_pdf_to_docx()], [check_docx_toolchain()], and [zip_reports()], whose docx_from_pdf argument accepts convert_pdf_to_docx.

convert_pdf_to_docx *Convert a Single PDF to DOCX Using the Best Available Backend*

Description

Convenience wrapper around [convert_pdfs_to_docx()] for one file. Its src, dest signature fits the docx_from_pdf callback of [zip_reports()].

Usage

```
convert_pdf_to_docx(
    src,
    dest = sub("\\.pdf$", ".docx", src, ignore.case = TRUE),
    backend = "auto",
    strict = TRUE,
    timeout = 600
)
```

Arguments

src	Path to the input PDF file.
dest	Path for the output DOCX file. Defaults to src with .pdf replaced by .docx.
backend	One of "auto", "python", "libreoffice", "word_com".
strict	If TRUE (default), error with an actionable message when no backend is available. Set FALSE to warn and skip.
timeout	Seconds before the Word COM backend gives up on a file.

Value

Invisibly, dest on success, or FALSE if conversion failed.

Examples

```
## Not run:
convert_pdf_to_docx("report.pdf")
zip_reports(reports, output_formats = c("pdf", "docx"), docx_from_pdf = convert_pdf_to_docx)

## End(Not run)
```

copy_data_manifest	<i>Copy Datasets to Snapshot Directory</i>
--------------------	--

Description

Copies files listed in a manifest from their source paths to a destination directory, preserving timestamps and validating contents immediately upon copy.

Usage

```
copy_data_manifest(  
  dir = ".",  
  manifest = NULL,  
  project_root = here::here(),  
  overwrite = TRUE  
)
```

Arguments

dir	Destination directory for frozen copies.
manifest	Data frame with columns file and source. If NULL, read from <code>file.path(dir, "manifest.csv")</code> .
project_root	Absolute or relative root directory for resolving source paths (default: <code>here::here()</code>).
overwrite	Logical; whether to overwrite existing destination files (default: TRUE).

Value

A validation tibble produced by [validate_data_manifest](#).

correlation_diff	<i>Difference in Correlation Matrices Between Two Datasets</i>
------------------	--

Description

Compares the correlation matrix of `comparison_data` against that of `baseline_data` (e.g. a later timepoint vs. baseline, or a treatment group vs. a reference group), matching numeric variables by column name.

Usage

```
correlation_diff(  
  baseline_data,  
  comparison_data,  
  cor.use = "everything",  
  cor.method = "pearson",  
  na_omit = TRUE  
)
```

Arguments

`baseline_data` A data frame or tibble treated as the reference.
`comparison_data` A data frame or tibble to compare against `baseline_data`.
`cor.use`, `cor.method`, `na.omit`
 See [correlation_plot](#).

Details

Unlike PCA loadings (see [pca_loading_diff](#)), correlation coefficients have no sign-ambiguity to correct for – a correlation matrix is uniquely determined by the data, so the difference is simply (comparison correlation) minus (baseline correlation), with no sign-alignment step needed.

Variables are matched by name (numeric column names shared by both datasets). If the two datasets' numeric columns differ, only the intersection is used; no error is raised. Because a correlation matrix is symmetric, only one triangle is returned (no duplicate `var1/var2` vs. `var2/var1` rows). The diagonal is dropped: a variable's correlation with itself is always 1 in both datasets, so its difference is always 0 and carries no information.

Value

A tibble with one row per shared variable pair, with columns `var1`, `var2`, and `diff` (comparison correlation minus baseline correlation).

Examples

```
set.seed(1)
baseline <- mtcars
comparison <- mtcars[sample(nrow(mtcars), replace = TRUE), ]
correlation_diff(baseline, comparison)
```

correlation_diff_heatmap

Heatmap of Correlation Differences Between Two Datasets

Description

Renders the output of [correlation_diff](#) as a heatmap (variable by variable), using the same visual language as [pca_loading_diff_heatmap](#): a diverging fill scale centered at zero, so variable pairs with little change are white and larger correlation changes in either direction stand out in Temple Night Owl or cherry.

Usage

```
correlation_diff_heatmap(
  baseline_data,
  comparison_data,
  cor.use = "everything",
  cor.method = "pearson",
  na.omit = TRUE
)
```

Arguments

`baseline_data` A data frame or tibble treated as the reference.
`comparison_data` A data frame or tibble to compare against `baseline_data`.
`cor.use`, `cor.method`, `na.omit`
 See [correlation_plot](#).

Value

A ggplot object.

Examples

```
set.seed(1)
baseline <- mtcars
comparison <- mtcars[sample(nrow(mtcars), replace = TRUE), ]
correlation_diff_heatmap(baseline, comparison)
```

correlation_plot	<i>Correlation Plot</i>
------------------	-------------------------

Description

Plots a correlation matrix across all numeric columns of data using **corrplot**.

Usage

```
correlation_plot(
  data,
  cor.use = "everything",
  cor.method = "pearson",
  method = "ellipse",
  type = "upper",
  order = "FPC",
  title = "Correlation Coefficient Plot",
  na.omit = TRUE,
  tl.cex = 0.5,
  number.cex = 0.75,
  tl.srt = 45,
  mar = c(0, 0, 2, 0),
  show_coef = TRUE,
  ...
)
```

Arguments

`data` A data frame or tibble.
`cor.use` Passed to [cor](#): "everything" (default), "all.obs", "complete.obs", "na.or.complete", or "pairwise.complete.obs".
`cor.method` Passed to [cor](#): "pearson" (default), "kendall", or "spearman".

<code>method</code>	Visualization method passed to corrplot (default "ellipse").
<code>type</code>	"upper" (default), "full", or "lower".
<code>order</code>	Ordering method for the correlation matrix (default "FPC", first principal component order).
<code>title</code>	Plot title.
<code>na_omit</code>	Logical (default TRUE); drop rows with any NA among the numeric columns before computing correlations.
<code>tl.cex, number.cex, tl.srt</code>	Label/number sizing and rotation, passed to corrplot .
<code>mar</code>	Plot margin, passed to corrplot in the standard <code>par("mar")</code> form <code>c(bottom, left, top, right)</code> (default <code>c(0, 0, 2, 0)</code>). <code>corrplot()</code> does not otherwise reserve any extra space above the matrix for title, so with the default (zero) margin the title text collides with the 45-degree diagonal column labels sitting just below it; the default here adds two lines of top margin so the title clears the labels. Increase further if using a long title or a larger <code>tl.cex</code> .
<code>show_coef</code>	Logical (default TRUE); whether to draw the correlation coefficient inside each cell. On a small matrix the numbers add useful precision on top of the visual encoding, but on a large matrix (many variables) they quickly overlap each other and the diagonal labels. Set to FALSE to omit them entirely – this is the clean way to declutter a large matrix; shrinking <code>tl.cex</code> and <code>number.cex</code> down to near-zero to visually "hide" the numbers (as earlier callers did) also erases the variable name labels and should be avoided. See correlation_plot_split for an alternative that keeps coefficients readable by splitting a large matrix into several smaller plots instead of hiding them.
<code>...</code>	Additional arguments passed to corrplot .

Value

Invisibly, the correlation matrix (called for its plot side effect).

Examples

```
correlation_plot(mtcars, tl.cex = .7)
correlation_plot(mtcars, tl.cex = .7, show_coef = FALSE)
```

`correlation_plot_split`

Correlation Plot, Split Into Legible Sub-Plots

Description

Draws the same style of correlation matrix as [correlation_plot](#), but for matrices with too many variables to stay legible in a single plot (for example, ~40 clinical parameters), it first groups variables via hierarchical clustering on their correlation structure and then draws one [correlation_plot](#)-style plot per group, showing only the within-group correlations.

Usage

```
correlation_plot_split(
  data,
  cor.use = "everything",
  cor.method = "pearson",
  na.omit = TRUE,
  group_size = 12,
  title = "Correlation Coefficient Plot",
  ...
)
```

Arguments

<code>data</code>	A data frame or tibble.
<code>cor.use</code> , <code>cor.method</code> , <code>na.omit</code>	See correlation_plot .
<code>group_size</code>	Target number of variables per sub-plot (default 12). The number of groups is <code>ceiling(n_vars / group_size)</code> ; groups may end up somewhat smaller or larger than this target since cutree produces clusters of whatever sizes the dendrogram structure dictates, not exactly equal-sized groups.
<code>title</code>	Base plot title; each sub-plot's title has " (Group i of n)" appended so the sub-plots can be told apart.
<code>...</code>	Additional arguments passed on to correlation_plot for each sub-plot (e.g. <code>tl.cex</code> , <code>show_coef</code> , <code>method</code>).

Details

Variables are clustered on `as.dist(1 - abs(cor_mat))` – the same correlation-based distance `corrplot`'s own `order = "hclust"` uses internally – so that variables which move together end up in the same sub-plot instead of being split arbitrarily (e.g. alphabetically). The resulting dendrogram is cut into `ceiling(n_vars / group_size)` contiguous groups via [cutree](#).

The default `group_size` of 12 is chosen to match [correlation_plot](#)'s own defaults: at the default `tl.cex/number.cex`, a matrix of roughly a dozen variables is about as many as can fit one legible diagonal label and coefficient per cell without crowding – the same order of magnitude as the small examples (like `mtcars`'s 11 numeric columns) `correlation_plot()` was originally tuned against. Larger or smaller groups can be requested via `group_size` depending on label length and plot size.

Value

Invisibly, a named list of per-group correlation matrices (one matrix per sub-plot, named "Group 1", "Group 2", ...).

Examples

```
# A data frame with more numeric columns than fit legibly in one plot.
wide_data <- cbind(mtcars, iris[seq_len(nrow(mtcars)), sapply(iris, is.numeric)])
correlation_plot_split(wide_data, group_size = 6)
```

corr_test_all

*Pairwise Correlation Tests Across All Numeric Columns***Description**

Runs `cor.test` on every pair of numeric columns in a data frame and returns one row per pair.

Usage

```
corr_test_all(
  data,
  method = "pearson",
  use = c("pairwise.complete.obs", "complete.obs"),
  columns = c("compact", "tidy"),
  sort = c("p_value", "estimate", "abs_estimate", "none"),
  ...
)
```

Arguments

data	A data frame or tibble. Non-numeric columns are ignored.
method	Correlation method passed to <code>cor.test</code> : "pearson" (default), "kendall", or "spearman".
use	How missing values are handled. "pairwise.complete.obs" (default) tests each pair on the rows where both columns are observed, which is what <code>cor.test()</code> does natively. "complete.obs" first drops every row with a missing value in any numeric column, so all pairs are tested on the same rows.
columns	Output shape. "compact" (default) returns var1, var2, r, p_value. "tidy" returns var1, var2, then every column broom::tidy() reports for the test, with estimate renamed cor: cor, statistic, p.value, parameter, conf.low, conf.high, method, alternative. parameter and the confidence limits are only reported by methods that compute them (Pearson).
sort	Row order. "p_value" (default) is ascending p-value; "estimate" is descending correlation; "abs_estimate" is descending absolute correlation; "none" keeps pair order. Ties keep pair order, and NAs sort last.
...	Further arguments passed to <code>cor.test</code> , such as alternative, conf.level, or exact.

Details

Pairs are enumerated in column order (via `combn`), so var1 is always the column that appears first in data. A pair whose test cannot be computed (e.g. fewer than three complete observations) is kept with NA results rather than aborting the whole run.

Value

A tibble with one row per pair of numeric columns.

Examples

```
corr_test_all(iris[, 1:4])

# Every cor.test() statistic, strongest positive correlation first
corr_test_all(mtcars[, c("mpg", "hp", "wt", "qsec")], columns = "tidy", sort = "estimate")
```

coxnet	<i>Penalized Cox Regression for Right-Censored or Start/Stop Survival Data</i>
--------	--

Description

Fits an elastic-net penalized Cox proportional hazards model with `[glmnet::glmnet()]`, through the tidymodels **hardhat** interface: a formula, a recipe, or predictors and outcome given separately. The outcome is a `[survival::Surv()]` object, either right-censored, `'Surv(time, event)'`, or counting-process, `'Surv(start, stop, event)'`, for time-varying covariates.

Usage

```
coxnet(x, ...)

## Default S3 method:
coxnet(x, ...)

## S3 method for class 'data.frame'
coxnet(x, y, penalty = NULL, mixture = 1, path = NULL, ...)

## S3 method for class 'matrix'
coxnet(x, y, penalty = NULL, mixture = 1, path = NULL, ...)

## S3 method for class 'formula'
coxnet(formula, data, penalty = NULL, mixture = 1, path = NULL, ...)

## S3 method for class 'recipe'
coxnet(x, data, penalty = NULL, mixture = 1, path = NULL, ...)
```

Arguments

x	A data frame or matrix of predictors, or a <code>[recipes::recipe()]</code> whose outcome is a single <code>'Surv'</code> column.
...	Further arguments passed to <code>[glmnet::glmnet()]</code> , such as <code>'nlambda'</code> or <code>'cox.ties'</code> . Not <code>'x'</code> , <code>'y'</code> , <code>'family'</code> , <code>'alpha'</code> , or <code>'lambda'</code> . <code>'cox.ties'</code> defaults to <code>"breslow"</code> , matching the Breslow baseline hazard used for survival predictions, so results don't depend on glmnet's own default (Breslow in glmnet 5.0, Efron from 5.1).
y	A <code>[survival::Surv()]</code> object with one element per row of <code>'x'</code> .
penalty	Default penalty (glmnet's <code>'lambda'</code>) used by <code>'predict()'</code> and <code>'tidy()'</code> . A single non-negative number, or <code>'NULL'</code> to require it there.
mixture	Elastic-net mixing parameter (glmnet's <code>'alpha'</code>): 1 is the lasso, 0 is ridge regression.
path	Optional decreasing sequence of penalties to fit, instead of glmnet's default path.

formula	A formula with a 'Surv()' outcome, such as 'survival::Surv(tstart, tstop, status) ~ age + sex'.
data	A data frame containing the variables in 'formula' or used by the recipe.

Details

The whole regularization path is fit, as glmnet recommends; 'penalty' only sets the default penalty for [predict()][predict.coxnet_model()] and [tidy()][tidy.coxnet_model()]. To choose 'penalty' and 'mixture' by cross-validation, use [cv_coxnet()].

Predictors must be numeric. The formula interface expands factors into one indicator column per level, which suits penalized models; with a recipe, add 'recipes::step_dummy()'. In a formula, write 'survival::Surv()' unless **survival** is attached.

Value

A 'coxnet_model' object: the glmnet fit, 'penalty', 'mixture', the training predictor matrix 'x' and outcome 'y' (kept so survival can be predicted at any penalty), and the hardhat 'blueprint'.

See Also

[predict.coxnet_model()], [tidy.coxnet_model()], [cv_coxnet()], [nested_cv_coxnet()]

Examples

```
if (requireNamespace("glmnet", quietly = TRUE) &&
    requireNamespace("survival", quietly = TRUE)) {
  lung <- stats::na.omit(survival::lung[, c("time", "status", "age", "sex", "ph.ecog", "wt.loss")])
  fit <- coxnet(survival::Surv(time, status) ~ ., data = lung, penalty = 0.05, mixture = 0.5)
  generics::tidy(fit)
  predict(fit, lung[1:3, ], type = "survival", eval_time = c(180, 365))$.pred[[1]]
}
```

create_qmd_renderer	Create a FileOutputs object for a rendering QMD or R Markdown document
---------------------	--

Description

Create a FileOutputs object for a rendering QMD or R Markdown document

Usage

```
create_qmd_renderer(
  name,
  path,
  deps = list(),
  file_stage = NA_character_,
  output_format = getOption("pipeline.default_formats", "pdf"),
  description = NA_character_
)
```

Arguments

name	The nickname for the file.
path	The full path to the .qmd or .Rmd file.
deps	A list of dependency objects (e.g., list(og_DATA)).
file_stage	The 'stage' for this file.
output_format	Output format(s): e.g. "pdf", "html", c("pdf", "docx"), or "yaml" to automatically extract the output formats declared in the document's YAML frontmatter. Defaults to getOption("pipeline.default_formats", "pdf").
description	A brief summary of the document.

Value

A FileOutputs object.

create_report	<i>Scaffold a New Report From a Template</i>
---------------	--

Description

Copies a bundled report template (and its supporting bibliography/title files) into location.

Usage

```
create_report(
  location = getwd(),
  template_name = "t_test_example",
  child = TRUE,
  type = ".qmd",
  include_bib = TRUE,
  include_tex = TRUE,
  install_brand = FALSE,
  filename = NULL
)
```

Arguments

location	Directory to create the report in (default getwd()).
template_name	One of "t_test_example" (default), "example", or "template".
child	Logical (default TRUE); also copy the child-document template.
type	One of ".qmd" (default) or ".Rmd".
include_bib	Logical (default TRUE); also copy the .bib file.
include_tex	Logical (default TRUE); also copy the title .tex file.
install_brand	Logical (default FALSE); for the "template" template, also run use_template_brand (location), which downloads the extension into location itself. Leave this FALSE when the extension is already installed at the project root (see use_template_brand); it doesn't need reinstalling per report.

filename Base name (no extension) for the report file, default NULL uses `template_name`. The report is always written to location, so calling `create_report()` twice for the same location with the same `template_name` (or the same `filename`) overwrites the first report; pass a distinct `filename` for each report that shares a location, e.g. `create_report("analysis", filename = "analysis1")` and `create_report("analysis", filename = "analysis2")`.

Details

The "template" template renders with the `quarto_template_brand` extension's `template-html`, `template-pdf`, and `template-typst` formats, so it needs that extension installed somewhere Quarto can find it from location (see [use_template_brand](#)): either at an ancestor project root shared by every report, or, with `install_brand = TRUE`, in location itself for a one-off report. It uses no `title.tex` or child document, and is Quarto-only.

Value

A list indicating whether each file was created.

Examples

```
## Not run:
create_report(here::here("analysis"))

# A project with several "template" reports: install the extension once at
# the project root, then scaffold each report without install_brand.
use_template_brand(here::here())
create_report(here::here("analysis"), template_name = "template")
create_report(here::here("reports"), template_name = "template", filename = "q3")

# A single one-off "template" report instead installs beside itself.
create_report(here::here("analysis"), template_name = "template", install_brand = TRUE)

# Two reports sharing one location need distinct `filename`s, or the
# second call overwrites the first report file:
create_report(here::here("analysis"), template_name = "template", filename = "analysis1")
create_report(here::here("analysis"), template_name = "template", filename = "analysis2")

## End(Not run)
```

create_toc_from_sas_pdf

Build a Table of Contents from a SAS-Generated PDF

Description

Scans a PDF (as produced by SAS reporting output) for table titles and writes a Table-of-Contents .rtf file alongside it, prefixed `TOC_`.

Usage

```
create_toc_from_sas_pdf(input, collapse = FALSE, top_margin_height = 60)
```

Arguments

`input` Path to the input PDF.

`collapse` If FALSE (default), every occurrence of a title is listed; if TRUE, only the first in a numbered sequence is kept.

`top_margin_height` How far down the page (in PDF text-position units) to look for titles.

Value

Invisibly, a message string naming the output file path.

Examples

```
## Not run:
create_toc_from_sas_pdf("path/to/report.pdf", collapse = TRUE)

## End(Not run)
```

cv_coxnet

*Cross-Validate a Penalized Cox Model With yardstick Survival Metrics***Description**

A tidymodels counterpart to [glmnet::cv.glmnet()] for right-censored and start/stop survival data. Every combination of ‘mixture’ and ‘penalty’ is fit on each analysis set and scored on the assessment set with a **yardstick** survival metric set, and the best settings are chosen by one metric: the integrated Brier score by default.

Usage

```
cv_coxnet(x, ...)

## Default S3 method:
cv_coxnet(x, ...)

## S3 method for class 'data.frame'
cv_coxnet(
  x,
  y,
  subject_id = NULL,
  group = NULL,
  mixture = 1,
  penalty = NULL,
  resamples = NULL,
  v = 10,
  metrics = NULL,
  eval_time = NULL,
  metric = NULL,
  covariates = c("path", "baseline"),
  trunc = 0.05,
  parallel = FALSE,
```

```

    ...
)

## S3 method for class 'matrix'
cv_coxnet(
  x,
  y,
  subject_id = NULL,
  group = NULL,
  mixture = 1,
  penalty = NULL,
  resamples = NULL,
  v = 10,
  metrics = NULL,
  eval_time = NULL,
  metric = NULL,
  covariates = c("path", "baseline"),
  trunc = 0.05,
  parallel = FALSE,
  ...
)

## S3 method for class 'formula'
cv_coxnet(
  formula,
  data,
  subject_id = NULL,
  group = NULL,
  mixture = 1,
  penalty = NULL,
  resamples = NULL,
  v = 10,
  metrics = NULL,
  eval_time = NULL,
  metric = NULL,
  covariates = c("path", "baseline"),
  trunc = 0.05,
  parallel = FALSE,
  ...
)

## S3 method for class 'recipe'
cv_coxnet(
  x,
  data,
  subject_id = NULL,
  group = NULL,
  mixture = 1,
  penalty = NULL,
  resamples = NULL,
  v = 10,
  metrics = NULL,

```

```

    eval_time = NULL,
    metric = NULL,
    covariates = c("path", "baseline"),
    trunc = 0.05,
    parallel = FALSE,
    ...
  )

```

Arguments

x	A data frame or matrix of predictors, or a [recipes::recipe()] whose outcome is a single 'Surv' column.
...	Further arguments passed to [glmnet::glmnet()], such as 'cox.ties' or 'nlambda'.
y	For the data frame and matrix methods, a [survival::Surv()] outcome with one element per row of 'x'.
subject_id	Subject identifier: a column name for the formula and recipe methods (for a recipe, a single column with role "id" is used when 'NULL'), or a vector for the data frame and matrix methods. Required for start/stop data; for right-censored data, each row is a subject when 'NULL'.
group	Optional grouping for folds, in the same form as 'subject_id'. Defaults to 'subject_id'.
mixture	Numeric vector of elastic-net mixing values to try.
penalty	Optional decreasing penalty path; by default glmnet's path for each 'mixture'.
resamples	An 'rset' of 'data', such as [rsample::group_vfold_cv()]. Overrides 'v' and 'group'.
v	Number of folds when 'resamples' is 'NULL'.
metrics	A [yardstick::metric_set()] of survival metrics. Defaults to the integrated Brier score, concordance, time-specific Brier score, and time-dependent ROC AUC.
eval_time	Evaluation times. Defaults to deciles (10th to 90th percentile) of the observed event times.
metric	Name of the metric used to choose settings; the first metric in 'metrics' by default.
covariates	"path" or "baseline"; see Details.
trunc	Lower bound for the censoring probability in the Graf weights.
parallel	If 'TRUE', fit folds with [furrr::future_map()]; set a [future::plan()] first.
formula	A formula with a 'Surv()' outcome, as in [coxnet()].
data	A data frame with the outcome, predictors, 'subject_id', and 'group' columns.

Details

****Folds.**** Unless 'resamples' is given, folds come from [rsample::group_vfold_cv()] grouped by 'group' (default: 'subject_id'), so a subject's start/stop rows are never split between analysis and assessment sets, as they would be by 'cv.glmnet()'s row-level folds. Group by a coarser unit, such as study site, for leave-sites-out cross-validation; each subject must then belong to one group.

****Fitting.**** For each 'mixture', one penalty path is computed on all of 'data' (as 'cv.glmnet()' does) and every fold is fit along it. With a recipe, the recipe is prepped on each analysis set, so preprocessing is learned inside the fold.

****Predictions.**** Survival for each assessment-set subject is computed from the Breslow baseline hazard of the analysis set. With `covariates = "path"`, the cumulative hazard is integrated over the subject's start/stop covariate path, carrying each interval's values forward to the next interval and past the last one; with `"baseline"`, only the covariates of the first interval are used, which uses no information from after time 0. The static concordance metric uses predicted survival at the last `'eval_time'`.

****Scoring.**** Start/stop rows are collapsed to one row per subject (`[surv_subject_truth()]`), weighted for censoring with Graf weights from the analysis set (`[graf_weights()]`), and passed to `'metrics'`.

****Selection.**** `'lambda.min'` is the penalty, across all `'mixture'` values, with the best mean `'metric'` over folds (for time-specific metrics, at the first `'eval_time'`). `'lambda.1se'` is the largest penalty, at the same mixture, whose mean is within one standard error of that best mean.

Value

A `'cv_coxnet'` object, with: * `'metrics'`: mean, `'n'`, and `'std_err'` of each metric by `'mixture'`, `'penalty'`, and `'eval_time'`; * `'fold_metrics'`: every fold's metrics; * `'best_by_mixture'`: `'lambda_min'` and `'lambda_1se'` for each `'mixture'`; * `'mixture'`, `'lambda_min'`, `'lambda_1se'`, `'metric'`, `'direction'`; * `'fit'`: a `[coxnet()]` model fit to all of `'data'` at the chosen `'mixture'`, with `'lambda_min'` as its default penalty; * `'eval_time'`, `'resamples'`, and `'covariates'`.

Use `'predict()'`, `'tidy()'`, `'autoplot()'`, and `'tune::collect_metrics()'` on it.

See Also

`[coxnet()]`, `[nested_cv_coxnet()]`, `[surv_subject_truth()]`

Examples

```
if (requireNamespace("glmnet", quietly = TRUE) &&
    requireNamespace("survival", quietly = TRUE) &&
    requireNamespace("rsample", quietly = TRUE) &&
    requireNamespace("yardstick", quietly = TRUE)) {
  set.seed(1)
  # Start/stop data: up to 3 visits per subject, events only at the last one
  long <- do.call(rbind, lapply(1:80, function(i) {
    k <- sample(1:3, 1)
    stops <- cumsum(stats::runif(k, 2, 6))
    risk <- stats::rnorm(1)
    data.frame(subject_id = i, tstart = c(0, utils::head(stops, -1)), tstop = stops,
               status = c(rep(0, k - 1), stats::rbinom(1, 1, stats::plogis(risk))),
               x1 = risk + stats::rnorm(k, sd = 0.2), x2 = stats::rnorm(k), x3 = stats::rnorm(k))
  )))
  cv <- cv_coxnet(
    survival::Surv(tstart, tstop, status) ~ x1 + x2 + x3,
    data = long, subject_id = "subject_id", mixture = c(0.5, 1),
    v = 5, nlambdas = 20, cox.ties = "breslow"
  )
  cv
  generics::tidy(cv, penalty = "lambda.1se")
}
```

cv_coxnet-methods *Use a 'cv_coxnet' Result*

Description

Predict from, tidy, plot, or collect the metrics of a [cv_coxnet()] result. Predictions and coefficients come from the model refit to all the data at the chosen 'mixture'.

Usage

```
## S3 method for class 'cv_coxnet'
predict(
  object,
  new_data,
  type = c("linear_pred", "survival"),
  penalty = "lambda.min",
  eval_time = NULL,
  increasing = TRUE,
  ...
)

## S3 method for class 'cv_coxnet'
tidy(x, penalty = "lambda.min", ...)

## S3 method for class 'cv_coxnet'
collect_metrics(x, ..., summarize = TRUE)

## S3 method for class 'cv_coxnet'
autoplot(object, ...)
```

Arguments

object, x	A [cv_coxnet()] result.
new_data	A data frame of new predictors.
type, eval_time, increasing	As for [predict.coxnet_model()].
penalty	"lambda.min" (default), "lambda.1se", or a number.
...	Not used.
summarize	For 'collect_metrics()': 'TRUE' for means over resamples, 'FALSE' for each resample's metrics.

Value

'predict()': a tibble as from [predict.coxnet_model()]. 'tidy()': the coefficients at 'penalty'. 'collect_metrics()': a tibble of metrics. 'autoplot()': a ggplot of the selection metric against the penalty, with 'lambda.min' (solid) and 'lambda.1se' (dashed) marked.

cv_joint_model

*Cross-Validation for Joint Models***Description**

Evaluates the [joint_model](#) trio across cross-validation splits (e.g. via [vfold_cv](#) or [group_vfold_cv](#)), benchmarking survival, classification, and regression paradigms using the IPCW Integrated Brier Score (`brier_survival_integrated`), Concordance, and Calibration.

Usage

```
cv_joint_model(
  data,
  outcome = survival::Surv(time, status) ~ .,
  v = 5,
  resamples = NULL,
  subject_id = NULL,
  engine = c("glmnet", "baguette", "stacks"),
  calibration = TRUE,
  parallel = FALSE,
  ...
)
```

Arguments

<code>data</code>	A data frame.
<code>outcome</code>	Survival formula with a <code>Surv()</code> outcome.
<code>v</code>	Number of cross-validation folds (default 5).
<code>resamples</code>	Optional pre-constructed rsample object.
<code>subject_id</code>	Optional subject identifier column for grouped resampling.
<code>engine</code>	Modeling engine: "glmnet", "baguette", or "stacks".
<code>calibration</code>	Logical; whether to calibrate status predictions (default TRUE).
<code>parallel</code>	Logical; whether to run folds in parallel via furrr .
<code>...</code>	Additional arguments passed to joint_model .

Value

An S3 object of class `c("cv_joint_model", "tbl_df")` summarizing comparative metrics across folds.

Note

Counting-process `Surv(start, stop, event)` outcomes are not yet supported here (only in [joint_model](#) itself): scoring needs one row per subject, but predictions are one row per interval. Use `Surv(time, status)` outcomes for cross-validation.

database_metadata_io *Export or Import Dataset Metadata*

Description

Exports and imports data dictionaries/metadata to and from .csv or .xlsx formats for in-place review and editing.

Usage

```
write_database_metadata(metadata, path, ...)

read_database_metadata(path, sheet = NULL, ...)
```

Arguments

metadata	A tibble/data.frame of dataset metadata (e.g. from get_dataset_info).
path	File path ending in .csv or .xlsx.
...	Additional arguments passed to the underlying writer/reader.
sheet	Optional sheet name when reading from Excel (defaults to 1 or "METADATA").

Value

For writers, path invisibly. For readers, a tibble of metadata.

delete_nul_files *Delete Stray 'nul' Files*

Description

Windows-only. knitr occasionally leaves behind a file literally named nul as a side effect of redirecting output to the Windows NUL device. This deletes files whose *basename* is exactly "nul" (case-insensitive) under path.

Usage

```
delete_nul_files(
  path = here::here(),
  .dontask = FALSE,
  .verify_command = FALSE
)
```

Arguments

path	Directory to search, defaults to <code>here::here()</code> .
.dontask	Logical (default FALSE); skip the confirmation prompt.
.verify_command	Logical (default FALSE); if TRUE, return the shell command(s) instead of running them.

Value

Invisibly, the deleted file paths (or the commands, if `.verify_command = TRUE`).

demo_cbe_mapping	<i>Create Sample Plug-and-Play Mapping Dictionary</i>
------------------	---

Description

Generates a demonstration column mapping table for documentation and testing.

Usage

```
demo_cbe_mapping()
```

Value

A tibble with sample mapping entries.

detect_outliers	<i>Detect Outliers Across a Data Frame's Numeric Columns</i>
-----------------	--

Description

Runs `flag_outliers` on every numeric column of data.

Usage

```
detect_outliers(data, outliers_only = TRUE)
```

Arguments

data	A matrix, data frame, or tibble.
outliers_only	Logical (default TRUE); if TRUE, only rows actually flagged as outliers are returned.

Value

A tibble with a column identifying which feature each row came from, plus value, `.outlier`, and `.outlier_type`.

Examples

```
df <- data.frame(a = c(1, 2, 3, 4, 100), b = c(10, 12, 11, 9, 8))
detect_outliers(df)
detect_outliers(df, outliers_only = FALSE)
```

distribution_plot	<i>Distribution Plot</i>
-------------------	--------------------------

Description

Ridgeline density plot of every numeric column in data, optionally normalized/standardized first so columns on different scales stay comparable.

Usage

```
distribution_plot(data, method = "range")
```

Arguments

data	A data frame or tibble.
method	Normalization applied before plotting: "range" (default, via range_norm), "min_max" (via min_max_norm), "z" (via z_norm), or "none"/"raw" to plot the values as-is.

Value

A ggplot object.

Examples

```
df <- data.frame(normal_dist = rnorm(1000, -10, .5), poisson_dist = rpois(1000, 5))
distribution_plot(df)
distribution_plot(df, method = "z")
distribution_plot(df, method = "none")
```

distribution_test	<i>Check a Vector or Data Frame's Distribution</i>
-------------------	--

Description

Runs both [is_normal](#) and [is_poisson](#) against a numeric vector, or against every numeric column of a data frame.

Usage

```
distribution_test(x)
```

Arguments

x	A numeric vector, matrix, or data frame/tibble.
---	---

Value

A tibble of test results (with a feature column when x has multiple columns).

Examples

```
distribution_test(rpois(n = 1000, lambda = 2))
distribution_test(mtcars)
```

export_interactive_pipeline

Export the computational graph to a self-contained interactive HTML file for stakeholders

Description

Export the computational graph to a self-contained interactive HTML file for stakeholders

Usage

```
export_interactive_pipeline(
  all_objects,
  file = here::here("R", "workflow_viz", "full_compute_graph.html"),
  direction = "LR",
  hierarchical = TRUE,
  title = getOption("pipeline.study_name", "Computational Pipeline"),
  subtitle = "Computational Pipeline"
)
```

Arguments

all_objects	A list of FilePath objects or an igraph object.
file	Path to the output HTML file. Default is here::here("R", "workflow_viz", "full_compute_graph.html").
direction	"LR" (default) or "UD".
hierarchical	Logical. Default TRUE.
title	Chart title. Defaults to the 'pipeline.study_name' option, or "Computational Pipeline" if unset.
subtitle	Chart subtitle.

Value

The saved file path (invisibly).

export_subgraph

Export a focused interactive subgraph centered around a focal report or stage

Description

Export a focused interactive subgraph centered around a focal report or stage

Usage

```
export_subgraph(
  all_objects,
  focal_node = NULL,
  stage = NULL,
  order = 1,
  file = NULL
)
```

Arguments

<code>all_objects</code>	List of <code>FilePath</code> objects or an <code>igraph</code> object.
<code>focal_node</code>	Character name of the node to focus on (e.g. "EDA_Report").
<code>stage</code>	Character name of the stage to filter by (e.g. "03_Imputation").
<code>order</code>	Degree of neighborhood around <code>focal_node</code> (default 1).
<code>file</code>	Destination HTML path.

Value

File path invisibly.

`extract_all_xlsx_tokens`

Extract Every '.xlsx' File Name From Text

Description

Extract Every '.xlsx' File Name From Text

Usage

```
extract_all_xlsx_tokens(x, xlsx_token_re = xlsx_name_re)
```

Arguments

<code>x</code>	Character vector, e.g. code lines.
<code>xlsx_token_re</code>	Regular expression for the file-name part. The default matches a name ending in '.xlsx': right after a quote or path separator it may contain spaces and punctuation other than quotes and separators; elsewhere only letters, digits, '_', '.', and '-'. Names followed by '(' (function calls such as 'read.xlsx()') are skipped.

Value

A list the length of 'x', each element a character vector of the '.xlsx' file names found in that string (empty when none).

See Also

[`extract_win_posix_paths()`]

Examples

```
extract_all_xlsx_tokens(c("read_excel('out/a.xlsx'); write_xlsx(df, 'b c.xlsx')", "none"))
```

extract_surv_components

Extract Survival Outcome Components and Predictors

Description

Helper utility that parses a survival formula into clean components supporting both 2-parameter `Surv(time, status)` and 3-parameter start/stop `Surv(tstart, tstop, status)` counting process structures.

Usage

```
extract_surv_components(data, outcome, subject_id = NULL)
```

Arguments

<code>data</code>	A data frame.
<code>outcome</code>	A survival formula or <code>Surv</code> expression.
<code>subject_id</code>	Optional subject identifier column name.

Value

A list with elements `surv_obj`, `time`, `start`, `status`, `type`, `predictors`, `pred_names`, and `formula`.

extract_win_posix_paths

Extract Full '.xlsx' Paths From Text

Description

Finds the first full Windows or POSIX path ending in an `'xlsx'` file name in each string, and splits it into folder and file name. Separators may repeat (`'C:/data/out'`, or `'C:\\data'` as written in R code). Relative paths such as `'output/t1.xlsx'` are not full paths and give `'NA'`.

Usage

```
extract_win_posix_paths(x, xlsx_token_re = xlsx_name_re)
```

Arguments

<code>x</code>	Character vector, e.g. code lines.
<code>xlsx_token_re</code>	Regular expression for the file-name part. The default matches a name ending in <code>'xlsx'</code> : right after a quote or path separator it may contain spaces and punctuation other than quotes and separators; elsewhere only letters, digits, <code>'_'</code> , <code>'.'</code> , and <code>'-'</code> . Names followed by <code>'('</code> (function calls such as <code>'read.xlsx()'</code>) are skipped.

Value

A tibble with one row per element of `'x'`: `'dir'`, `'file'`, and `'full_path'` (`'NA'` where no path was found).

See Also

[extract_all_xlsx_tokens()]

Examples

```
extract_win_posix_paths(c("C:\\data\\table_1.xlsx", "/home/me/out/table-2.xlsx", "none"))
```

extract_yaml_formats *Extract Output Formats From Document YAML Front Matter*

Description

Inspects the YAML header of a Quarto (.qmd) or R Markdown (.Rmd) file and returns all declared output formats (e.g. "html", "pdf", "docx", "gfm").

Usage

```
extract_yaml_formats(file_path)
```

Arguments

file_path Path to the .qmd or .Rmd document.

Value

Character vector of lowercase format names, or empty character vector if none found.

features_percent_miss *Calculate Percentage of Missing Data Per Feature*

Description

Computes the count and percentage of missing and complete values for each column in a dataset.

Usage

```
features_percent_miss(data, na_list = NULL)
```

Arguments

data A data frame or tibble.

na_list Optional vector of additional values to treat as missing, passed to [SumNa](#) (beyond actual NA).

Value

A tibble with columns 'feature', 'SumNa', 'SumComp', 'PctNa', 'PctComp' sorted descending by 'PctNa'.

Examples

```
df <- data.frame(a = c(1, 2, NA, 4), b = c(NA, NA, 3, 4))
features_percent_miss(df)

df2 <- data.frame(a = c(1, "NA", 3), b = c("", 2, 3))
features_percent_miss(df2, na_list = c("NA", ""))
```

FileOutputs

*Construct a FileOutputs object***Description**

User-friendly constructor function for instantiating a [FileOutputs](#) object.

Usage

```
FileOutputs(
  name,
  path,
  dependencies = list(),
  output = list(),
  renders = TRUE,
  stage = NA_character_,
  artifact_role = "report_source",
  description = NA_character_
)
```

Arguments

name	Character. User-supplied human-readable name for the file.
path	Character. Absolute or relative file system path.
dependencies	List of FilePath objects that this file depends on.
output	List of FilePath objects produced by this file.
renders	Logical. Indicates whether the file is a rendering script. Defaults to TRUE.
stage	Character. Pipeline lifecycle stage.
artifact_role	Character. Pipeline role of the artifact. Defaults to "report_source".
description	Character. Brief description of the file.

Value

A new [FileOutputs](#) object.

FileOutputs-class	<i>FileOutputs S4 Class</i>
-------------------	-----------------------------

Description

Represents an executable or rendering pipeline file that produces downstream output artifacts.

Usage

```
## S4 method for signature 'FileOutputs'
show(object)
```

Arguments

object A [FileOutputs](#) object.

Details

Inherits from [FileUses](#) (inheriting upstream dependencies), while adding a list of downstream output [FilePath](#) objects produced when this file executes or renders.

Slots

output List of [FilePath](#) objects produced by this file.

See Also

[FilePath](#), [FileUses](#), [FileOutputs](#), [create_qmd_renderer](#)

FilePath	<i>Construct a FilePath object</i>
----------	------------------------------------

Description

User-friendly constructor function for instantiating a [FilePath](#) object.

Usage

```
FilePath(
  name,
  path,
  renders = FALSE,
  stage = NA_character_,
  artifact_role = NA_character_,
  description = NA_character_
)
```

Arguments

name	Character. User-supplied human-readable name for the file.
path	Character. Absolute or relative file system path.
renders	Logical. Indicates whether the file is a rendered file or rendering script. Defaults to FALSE.
stage	Character. Pipeline lifecycle stage.
artifact_role	Character. Pipeline role of the artifact.
description	Character. Brief description of the file.

Value

A new [FilePath](#) object.

FilePath-class	<i>FilePath S4 Class</i>
----------------	--------------------------

Description

File paths are unique identifiers for files and are used to track changes in the file system. A file path is a general R, Python, SAS, Quarto, shell, etc. agnostic way to refer to a file. It is not tied to any specific software or programming language, but rather to the file system itself.

Usage

```
## S4 method for signature 'FilePath'
show(object)
```

Arguments

object A [FilePath](#) object.

Details

Represents an individual file node in a computational pipeline dependency graph. A [FilePath](#) object encapsulates file-level metadata and state: tracking modification timestamps (`mtime`), file extensions (`file_ext`), workflow lifecycle milestones (`stage`), semantic pipeline roles (`artifact_role`), and rendering status (`renders`).

Upon instantiation, the object automatically inspects the target file on disk: if the file exists at `path`, `mtime` is populated from `file.info()$mtime` and `file_ext` is extracted via `tools::file_ext()`.

Slots

`name` Character. A user-supplied, human-readable name for the file, used to refer to the file across other parts of the pipeline and dependency graph.

`path` Character. The absolute or project-relative path to the file on the file system.

`mtime` POSIXct. The modification time of the file on the file system. Automatically populated from disk during object initialization if the file exists.

`file_ext` Character. The file extension of the file (e.g., `"xlsx"`, `"rds"`, `"qmd"`). Automatically populated from the file path during object initialization.

renders Logical. A logical indicating whether the file is a rendered file (i.e. a file that is produced by a rendering process) or a rendering source script. Defaults to FALSE.

stage Character. The workflow lifecycle stage of the file in the pipeline (e.g., "01_EDA", "02_Analysis_Original", "03_Imputation").

artifact_role Character. The semantic role of the file in the pipeline. Standard roles include:

- "raw_data": Original immutable input data (e.g. source spreadsheets).
- "derived_data": Primary processed baseline dataset generated from raw inputs.
- "enhanced_data": Imputed or augmented analytical dataset.
- "intermediate_data": Stage-internal calculation artifact or model.
- "report_source": Executable source document (e.g., Quarto .qmd).
- "deliverable_report": Rendered deliverable (e.g., .pdf, .html).
- "helper_script": Child or modular helper script invoked by parent reports.
- "reference": Lookup tables, crosswalks, or project configuration files.

description Character. A brief human-readable description of the file's contents or analytical purpose.

See Also

[FileUses](#), [FileOutputs](#), [FilePath](#), [create_qmd_renderer](#)

FileUses

Construct a FileUses object

Description

User-friendly constructor function for instantiating a [FileUses](#) object.

Usage

```
FileUses(
  name,
  path,
  dependencies = list(),
  renders = FALSE,
  stage = NA_character_,
  artifact_role = NA_character_,
  description = NA_character_
)
```

Arguments

name	Character. User-supplied human-readable name for the file.
path	Character. Absolute or relative file system path.
dependencies	List of FilePath objects that this file depends on.
renders	Logical. Indicates whether the file is a rendered file or rendering script. Defaults to FALSE.
stage	Character. Pipeline lifecycle stage.
artifact_role	Character. Pipeline role of the artifact.
description	Character. Brief description of the file.

Value

A new [FileUses](#) object.

FileUses-class	<i>FileUses S4 Class</i>
----------------	--------------------------

Description

Represents a pipeline file that depends on other upstream files. Extends [FilePath](#) by adding upstream input dependency tracking.

Usage

```
## S4 method for signature 'FileUses'
show(object)
```

Arguments

object A [FileUses](#) object.

Details

A FileUses node represents a computational artifact whose execution or validity relies upon one or more input [FilePath](#) objects listed in dependencies.

Slots

dependencies List of [FilePath](#) objects that this file depends on.

See Also

[FilePath](#), [FileOutputs](#), [FileUses](#)

file_meta_fs	<i>File Metadata as a Tibble</i>
--------------	----------------------------------

Description

Vectorized file information from `[fs::file_info()]`, with names and parent folders split out for joining.

Usage

```
file_meta_fs(paths)
```

Arguments

paths Character vector of file paths.

Value

A tibble with one row per path: 'path', 'file_name', 'dir_name' (parent folder name), 'dir_path', 'm_time' (modified), 'c_time' (created, falling back to changed, then modified), 'size' in bytes, and 'uname' (owner, where the platform reports one).

See Also

[scan_data_io()]

Examples

```
f <- tempfile(fileext = ".csv")
writelines("a,b", f)
file_meta_fs(f)
```

find_code

Search for Code Patterns Across a Directory Tree

Description

Recursively searches R-related files (.R, .Rmd, .qmd, etc.) under directory for a string or regex pattern.

Usage

```
find_code(
  directory,
  pattern,
  lines_before = 0,
  lines_after = 0,
  regex = FALSE,
  ignore_case = FALSE,
  include_comments = TRUE,
  exts = c("R", "r", "Rmd", "rmd", "qmd"),
  exclude_dirs = c(".git", "renv", "_freeze", "_book", "_site", ".quarto", "_extensions",
    "cache", ".Rproj.user"),
  return_all_matches = FALSE,
  match_extractor = NULL
)
```

Arguments

directory	Character. Root directory to search.
pattern	Character. String or regex to search for.
lines_before, lines_after	Integers. Context lines around each hit.
regex	Logical. If TRUE, treat pattern as a Perl regex.
ignore_case	Logical. Case-insensitive when TRUE (works correctly even in fixed-string mode, where base grep 's own ignore.case argument is silently ignored).

`include_comments` Logical. If FALSE, skip comment-only lines.
`exts` Character vector of file extensions to include.
`exclude_dirs` Character vector of directory names to exclude.
`return_all_matches` Logical. If TRUE, returns one row per extracted match token (via `match_extractor`) instead of one row per matching line.
`match_extractor` Optional function(`content_subset`) returning a list of tokens per line; only used when `return_all_matches = TRUE`.

Value

A tibble: file, path, line_number, line, and (when `return_all_matches = TRUE`) match.

Examples

```
find_code(system.file("R", package = "TempleCBE"), "roxygen2", ignore_case = TRUE)
```

find_correlation	<i>Find Highly Correlated Columns</i>
------------------	---------------------------------------

Description

Identifies numeric columns to drop for multicollinearity, via [findCorrelation](#) on the pairwise correlation matrix of data's numeric columns. Named `find_correlation()` (not `findCorrelation()`) so it doesn't shadow **caret**'s function of the same name for anyone with both packages loaded.

Usage

```
find_correlation(
  data,
  use = "pairwise.complete.obs",
  method = "pearson",
  cutoff = 0.9,
  verbose = FALSE,
  names = TRUE,
  exact = ncol(data) < 100
)
```

Arguments

`data` A data frame or tibble.
`use` Passed to [cor](#) (default "pairwise.complete.obs").
`method` Passed to [cor](#) (default "pearson").
`cutoff` Absolute correlation above which a column is flagged (default 0.9).
`verbose` Logical; passed to [findCorrelation](#).
`names` Logical; if TRUE (default) return column names instead of indices.
`exact` Passed to [findCorrelation](#); defaults to `ncol(data) < 100` (the original had a bug here referencing an undefined `x` instead of `data` — fixed).

Value

Character vector (or integer indices) of columns to remove.

Examples

```
if (requireNamespace("caret", quietly = TRUE)) {  
  find_correlation(mtcars, cutoff = 0.8)  
}
```

find_python

Locate a Python Interpreter That Can Import pdf2docx

Description

Search order (first verified hit wins):

1. `getOption("templecbe.python")`
2. `Sys.getenv("TEMPLECBE_PYTHON")`
3. `Sys.getenv("RETICULATE_PYTHON")` and `reticulate::py_exe()`
4. `Sys.which("python3")`, `Sys.which("python")`
5. Platform-specific well-known locations

Usage

```
find_python(verify = TRUE)
```

Arguments

`verify` If TRUE, require import pdf2docx to succeed.

Details

Each candidate is verified by actually running it, so Windows App Execution Alias stubs (which resolve on PATH but do nothing) are rejected. When `verify = TRUE` the candidate must additionally be able to import pdf2docx. Install the pinned Python requirements with `pip install -r on_system.file("python", "requirements.txt", package = "TempleCBE")`.

Value

Path to a usable interpreter, or NULL.

See Also

`[check_docx_toolchain()]`

find_sas

Locate a SAS Executable

Description

Checks, in order: 'getOption("templecbe.sas")', the 'SAS_EXE' environment variable, the 'SAS-ROOT' environment variable ('sas.exe'), 'sas' on the 'PATH', and default SAS install locations ('SASHome/SASFoundation/<version>' under 'C:/Program Files', 'C:/Program Files (x86)', 'C:', 'D:/Program Files', or 'D:' on Windows; under '/usr/local/SASHome', '/opt/sas/SASHome', or '/opt/SASHome' elsewhere), preferring the highest installed version.

Usage

```
find_sas()
```

Value

The path to the SAS executable, or 'NULL' if none is found.

See Also

[run_sas_script()], [cbe_sas_macro_path()], [cbe_sas_macro_dir()]

Examples

```
find_sas()
```

find_section_file

Find Section Data File in a Directory

Description

Find Section Data File in a Directory

Usage

```
find_section_file(data_dir, index)
```

Arguments

data_dir	Directory to search.
index	Section name/token to look for.

Value

File path.

find_soffice	<i>Locate a LibreOffice Headless Binary</i>
--------------	---

Description

Honours `getOption("templecbe soffice")` and the `TEMPLECBE_SOFFICE` environment variable before searching the `PATH` and default install locations.

Usage

```
find_soffice()
```

Value

Path to `soffice/libreoffice`, or `NULL`.

See Also

`[check_docx_toolchain()]`

fit_n_rank	<i>Fit the Configuration Ranked ‘rank’ in a Workflow Set</i>
------------	--

Description

Ranks the configurations of tuned workflow set results by ‘rank_metric’, picks the one at position ‘rank’, and fits it on the full training set with `[tune::fit_best()]`.

Usage

```
fit_n_rank(
  grid_results,
  rank_metric,
  .rank = 1,
  group_wflow = TRUE,
  select_best = FALSE,
  ...
)
```

Arguments

<code>grid_results</code>	Tuned workflow set results, the output of <code>[workflowsets::workflow_map()]</code> run with ‘control = <code>tune::control_grid(save_workflow = TRUE)</code> ’, which <code>[tune::fit_best()]</code> needs.
<code>rank_metric</code>	Name of the metric to rank by.
<code>.rank</code>	Position to fit; 1 (default) is the best.
<code>group_wflow</code>	If ‘TRUE’ (default), rank workflows, each by its best configuration, so ‘rank = 2’ is the second-best <i>*workflow*</i> . If ‘FALSE’, rank configurations directly, so ‘rank = 2’ can be a second configuration of the best workflow.

select_best	Passed to [workflowsets::rank_results()]: keep only each workflow's best configuration before ranking.
...	Passed to [tune::fit_best()].

Value

A list with 'model', the fitted workflow, and 'tuned_parameters', a one-row tibble of the workflow's 'wflow_id', 'preproc', 'model', tuning parameters, and its [workflowsets::rank_results()] row ('.config', '.metric', 'mean', 'std_err', 'n', 'rank', and 'overall_rank' when 'group_wflow = TRUE').

See Also

[get_model_parameters()]

Examples

```
## Not run:
results <- workflowsets::workflow_map(
  wset, "tune_grid", resamples = folds, grid = 5,
  control = tune::control_grid(save_workflow = TRUE)
)
best <- fit_n_rank(results, rank_metric = "rmse")
best$tuned_parameters
predict(best$model, new_data)

## End(Not run)
```

flag_outliers

Flag and Classify Outliers

Description

Flag and Classify Outliers

Usage

```
flag_outliers(col)
```

Arguments

col	A numeric vector or column.
-----	-----------------------------

Value

A one-column-input-turned-tibble with value, .outlier (factor: FALSE, TRUE) and .outlier_type (factor: "NONE", "MILD", or "EXTREME").

Examples

```
flag_outliers(c(1, 2, 3, 4, 5, 100))
```

fmt_hr*Format Hazard Ratio and Confidence Interval*

Description

Format Hazard Ratio and Confidence Interval

Usage

```
fmt_hr(hr, conf_low, conf_high, digits = 2)
```

Arguments

hr	Point estimate for Hazard Ratio
conf_low	Lower 95% confidence limit
conf_high	Upper 95% confidence limit
digits	Number of decimal places (default: 2)

Value

Formatted string (e.g., "HR 1.25 (95% CI 1.05–1.48)")

fmt_num*Format Numbers to Fixed Decimals*

Description

Format Numbers to Fixed Decimals

Usage

```
fmt_num(x, digits = 2)
```

Arguments

x	Numeric value
digits	Number of decimal places (default: 2)

Value

Formatted character string

`fmt_p`*Format P-Values*

Description

Formats p-values with institutional threshold (< 0.001).

Usage

```
fmt_p(p)
```

Arguments

<code>p</code>	Numeric p-value
----------------	-----------------

Value

Formatted string (e.g., "p = 0.034" or "p < 0.001")

`fmt_pct`*Format a Percentage*

Description

Format a Percentage

Usage

```
fmt_pct(x, digits = 1)
```

Arguments

<code>x</code>	Numeric proportion between 0 and 1
<code>digits</code>	Number of decimal places (default: 1)

Value

Formatted character string with `'%'`

fmt_sig	<i>Format Significant Figures</i>
---------	-----------------------------------

Description

Format Significant Figures

Usage

```
fmt_sig(x, digits = 3)
```

Arguments

x	Numeric value
digits	Number of significant figures (default: 3)

Value

Formatted character string

generate_last_names	<i>Procedural Synthetic Last Name Generator</i>
---------------------	---

Description

Generates realistic, pronounceable synthetic surnames using phonetic onsets, vowels, and common surname endings (e.g., '-berg', '-ton', '-man', '-wood', '-ford'). Supports strict exclusion lists (e.g. to guarantee real investigator surnames are never generated) and reproducible seed-based sampling without mutating the global RNG state.

Usage

```
generate_last_names(  
  n = 10,  
  seed = NULL,  
  exclude = character(),  
  max_tries = 100L  
)
```

Arguments

n	Integer, number of unique names to generate (default 10).
seed	Optional integer, random seed for reproducibility.
exclude	Character vector of surnames (case-insensitive) that must not be generated.
max_tries	Integer, maximum attempts multiplier before throwing an error if distinct names cannot be generated.

Value

A character vector of length ‘n’ containing synthetic surnames.

Examples

```
# Generate 5 random last names
generate_last_names(5)

# Reproducible generation with exclusions
generate_last_names(3, seed = 1, exclude = c("Madison", "Davison"))
```

generate_pi_names	<i>Generate Random PI Names or Tokens</i>
-------------------	---

Description

Generates random synthetic principal investigator (PI) names or cryptographic pseudonym tokens ('PI...').

Usage

```
generate_pi_names(
  n = 1,
  format = c("synthetic", "token"),
  prefix = "PI_",
  n_chars = 16,
  rename_fn = NULL,
  seed = NULL
)
```

Arguments

n	Integer, number of names to generate (default 1).
format	Character, either "synthetic" (realistic surnames) or "token" (cryptographic hash tokens).
prefix	Character, prefix when format is "token" (default "PI_").
n_chars	Integer, number of hex characters from the hash digest when 'format = "token"' (default 16).
rename_fn	Optional function to rename or format generated tokens/names.
seed	Optional integer, random seed for reproducibility.

Value

A character vector of length ‘n’ (or a single character string if ‘n = 1’).

Examples

```
generate_pi_names(1)
generate_pi_names(3, format = "token")
generate_pi_names(3, format = "token", n_chars = 8)
```

generate_pseudonym_token

Generate Pseudonym Tokens (Cryptographic Hash)

Description

Generates secure cryptographic hash pseudonym tokens. When ‘name’ is supplied, computes deterministic keyed HMAC-SHA256 tokens. When ‘name’ is ‘NULL’, generates random cryptographic hash tokens (or reproducible hash tokens when ‘seed’ is provided). Supports custom prefix, configurable character length (e.g. for short tokens), and custom renaming/formatting functions.

Usage

```
generate_pseudonym_token(
  name = NULL,
  n = 1,
  prefix = "PI_",
  key = Sys.getenv("TEMPLECBE_SECRET_KEY", unset = ""),
  n_chars = 16,
  rename_fn = NULL,
  seed = NULL
)
```

Arguments

name	Optional character vector of names to tokenize. If provided, computes a deterministic HMAC-SHA256 token for each non-NA name.
n	Integer, number of tokens to generate when ‘name’ is ‘NULL’ (default 1).
prefix	Character string, prefix prepended to each token (default ‘PI_’).
key	Character string, secret key / pepper for HMAC tokenization. Defaults to the environment variable ‘Sys.getenv(“TEMPLECBE_SECRET_KEY”, unset = “)’.
n_chars	Integer, number of hex characters from the hash digest to include (default 16). Can be shortened (e.g. 4 or 8) to produce compact/short tokens.
rename_fn	Optional function to transform or rename tokens (e.g., ‘tolower’, ‘toupper’, or a custom function such as ‘function(tok) paste0(“INV-”, substr(tok, 4, 8))’).
seed	Optional integer, random seed for reproducible random token generation when ‘name’ is ‘NULL’.

Value

Character vector of pseudonym tokens.

Examples

```
# Random cryptographic hash tokens
generate_pseudonym_token(n = 3)

# Short hash tokens (4 hex chars)
generate_pseudonym_token(n = 3, n_chars = 4)
```

```
# Custom renaming function
generate_pseudonym_token(n = 2, rename_fn = tolower)

# Stateless deterministic token with key
generate_pseudonym_token("Franklin", key = "study_salt")

# Vectorized tokenization
generate_pseudonym_token(c("Smith", "Jones"), key = "study_salt")
```

get_dataset_info

Summarize a Data Frame or Joint Model's Columns and Components

Description

Per-column metadata: class, variable label (if set via **labelled**), mean/sd for numeric columns, most-frequent value, distinct-value count, and missingness. `survival::Surv` columns are summarized from their underlying time/status or start/stop counting-process matrix rather than unrolled as plain numerics. When provided a `joint_model` object, summarizes the fitted training data and attaches model metadata. When provided a list of data frames (an R database), returns a consolidated data dictionary across all tables.

Usage

```
get_dataset_info(x, ...)

## S3 method for class 'data.frame'
get_dataset_info(
  x,
  subject_id = NULL,
  dataset_name = NULL,
  database_name = NULL,
  dataset_label = NULL,
  ...
)

## S3 method for class 'list'
get_dataset_info(x, subject_id = NULL, database_name = NULL, ...)

## S3 method for class 'joint_model'
get_dataset_info(x, ...)

proc_contents(x, ...)
```

Arguments

<code>x</code>	A data frame, tibble, named list of data frames, or a fitted <code>joint_model</code> object.
<code>...</code>	Additional arguments passed to methods.
<code>subject_id</code>	Optional character string specifying the subject identifier column for repeated-measures longitudinal datasets to audit time-varying vs. baseline features.
<code>dataset_name</code>	Optional character string overriding the displayed dataset name.
<code>database_name</code>	Optional character string specifying the parent database name.
<code>dataset_label</code>	Optional character string overriding the dataset-level description.

Value

A tibble with one row per column: dataset_name, (optional database_name and dataset_label), labels, columns, class, mean, sd, most_freq, n_distinct, SumNa, PctNa, and optionally variable_type when subject_id is specified.

Examples

```
get_dataset_info(mtcars)
```

get_model_parameters	<i>Tuning Parameters of a Ranked Workflow in a Workflow Set</i>
----------------------	---

Description

Ranks the workflows of tuned workflow set results by 'rank_metric', each workflow represented by its best configuration, and returns the preprocessor, model, and best tuning parameters of the workflow at position '.rank'.

Usage

```
get_model_parameters(workflowset_results, rank_metric, .rank = 1)
```

Arguments

workflowset_results	Tuned workflow set results, the output of [workflowsets::workflow_map()].
rank_metric	Name of the metric to rank by, e.g. "rmse" or "brier_survival_integrated".
.rank	Position to return; 1 (default) is the best workflow.

Value

A one-row tibble with 'wflow_id', 'preproc', 'model', and one column per tuning parameter.

See Also

[fit_n_rank()] to also fit the chosen configuration.

Examples

```
## Not run:
results <- workflowsets::workflow_map(wset, "tune_grid", resamples = folds, grid = 5)
get_model_parameters(results, "rmse")
get_model_parameters(results, "rmse", .rank = 2)

## End(Not run)
```

get_render_plan	<i>Get the topologically sorted list of QMDs that need to be re-rendered</i>
-----------------	--

Description

Get the topologically sorted list of QMDs that need to be re-rendered

Usage

```
get_render_plan(all_objects)
```

Arguments

`all_objects` The full list of FilePath, FileUses, and FileOutputs objects.

Value

A list of FileOutputs objects (the QMDs) in the correct render order.

glmnet_IBS	<i>Integrated Brier Score of a Penalized Cox Model on Start/Stop Survival Data</i>
------------	--

Description

Tunes an elastic-net Cox model on the analysis set of one resampling split and scores it on the assessment set by the integrated Brier score (IBS). Lower is better. A convenience wrapper around [cv_coxnet()] for data in counting-process (start/stop) layout, kept for existing analysis code; new code can call [cv_coxnet()] and [nested_cv_coxnet()] directly.

Usage

```
glmnet_IBS(
  object,
  alpha = 1,
  recipe,
  feature_names,
  time_data = NULL,
  formula = NULL,
  internal_folds = 5,
  id_col = "id",
  start_col = "tstart",
  stop_col = "tstop",
  status_col = "status",
  censoring_weights = c("ipcw", "none"),
  prep_on = c("baseline", "all"),
  eval_time = NULL,
  metric = "brier_survival_integrated",
  rule = c("min", "1se"),
  covariates = c("path", "baseline"),
```

```

    failure_ibs = NA_real_,
    ...
  )

```

Arguments

<code>object</code>	An 'rsplit' (e.g. one element of 'rsample' resamples).
<code>alpha</code>	Elastic net mixing parameter: 1 is lasso, 0 is ridge.
<code>recipe</code>	An unprepped [recipes::recipe()].
<code>feature_names</code>	Character vector of baked predictor column names, or a function that takes the baked analysis set and returns them – for recipes whose output columns vary by fold, such as 'step_pca(threshold =)'.
<code>time_data</code>	Optional data frame with columns 'start_col' and 'stop_col'; its positive stop times are used as 'eval_time' when that is 'NULL'.
<code>formula</code>	Optional character string of '+'-separated feature names (or a one-sided formula) to restrict the model to a subset of 'feature_names'. Defaults to all of them.
<code>internal_folds</code>	Number of subject-grouped folds used to choose the penalty.
<code>id_col, start_col, stop_col, status_col</code>	Column names of the subject identifier, interval start, interval stop, and event indicator (1 = event).
<code>censoring_weights</code>	Only "ipcw" is supported. "none", which scored interval rows without adjusting for censoring, was removed in TempleCBE 0.3.0; install TempleCBE 0.2.0 to reproduce results that used it.
<code>prep_on</code>	"baseline" (default) preps 'recipe' on each subject's first interval in the analysis set; "all" preps on every analysis row.
<code>eval_time</code>	Evaluation times for the IBS. Defaults to the stop times of 'time_data', or else deciles of the analysis-set event times.
<code>metric</code>	Name of the yardstick survival metric that chooses the penalty, e.g. "brier_survival_integrated" or "concordance_survival".
<code>rule</code>	"min" or "lse": report the penalty at 'lambda.min' or 'lambda.1se'.
<code>covariates</code>	"path" or "baseline"; see [cv_coxnet()].
<code>failure_ibs</code>	IBS reported when the model cannot be fit (default 'NA'); a warning gives the reason.
<code>...</code>	Further arguments passed to [glmnet::glmnet()], such as 'cox.ties' ("breslow" or "efron"). glmnet 5.0 defaults to Breslow and 5.1 to Efron, so pass 'cox.ties' to keep results stable across glmnet versions.

Details

1. **Preprocessing.** 'recipe' is prepped on the analysis set – only each subject's first interval when 'prep_on = "baseline"' – and baked onto both sets. The recipe must keep 'id_col', 'start_col', 'stop_col', and 'status_col' in its output (e.g. with an "id variable" role). 2. **Tuning.** [cv_coxnet()] chooses the penalty on the baked analysis set with 'internal_folds' folds grouped by subject, by 'metric' (the IBS by default). Earlier versions used [glmnet::cv.glmnet()], which chose the penalty by concordance on folds of rows, splitting subjects' intervals between folds. 3. **Scoring.** The model, refit on the whole analysis set, predicts each assessment-set subject's survival from the analysis-set Breslow baseline hazard along their covariate path (see [cv_coxnet()]), and is scored by [yardstick::brier_survival_integrated()] with inverse-probability-of-censoring (Graf) weights estimated on the analysis set.

Value

A tibble with columns 'IBS', 'lambda', 'term', 'estimate', and 'alpha': one row per feature (including those the penalty set to 0). When the fit fails, a single row with 'IBS = failure_ibs', 'lambda = NA', and 'alpha'.

References

Graf E, Schmoor C, Sauerbrei W, Schumacher M (1999). Assessment and comparison of prognostic classification schemes for survival data. **Statistics in Medicine**, 18(17-18), 2529-2545.

See Also

[cv_coxnet()], [nested_cv_coxnet()], [tune_over_alpha()], [summarize_tune_results()]

Examples

```
if (requireNamespace("glmnet", quietly = TRUE) &&
    requireNamespace("yardstick", quietly = TRUE) &&
    requireNamespace("rsample", quietly = TRUE)) {
  set.seed(1)
  # Synthetic start/stop data: 60 subjects, up to 3 intervals each
  long <- do.call(rbind, lapply(1:60, function(i) {
    k <- sample(1:3, 1)
    stops <- c(5, 10, 20)[1:k]
    risk <- rnorm(1)
    data.frame(subject_id = i, tstart = c(0, head(stops, -1)), tstop = stops,
               status = c(rep(0, k - 1), rbinom(1, 1, plogis(risk))),
               x1 = risk + rnorm(k, sd = 0.2), x2 = rnorm(k))
  )))
  rec <- recipes::recipe(~ ., data = long) |>
    recipes::update_role(subject_id, tstart, tstop, status, new_role = "id variable") |>
    recipes::step_range(recipes::all_numeric_predictors())
  split <- rsample::group_initial_split(long, group = subject_id)
  glmnet_IBS(split, alpha = 0.5, recipe = rec, feature_names = c("x1", "x2"),
             time_data = unique(long[, c("tstart", "tstop")]), internal_folds = 3,
             id_col = "subject_id", cox.ties = "breslow")
}
```

graf_weights

Inverse Probability of Censoring Weights (Graf et al.)

Description

Computes the weights time-dependent survival metrics need so that subjects lost to follow-up don't bias the score: at evaluation time t , a subject still under observation ($T > t$) gets $1/G(t)$; a subject with an event by t gets $1/G(T^-)$; a subject censored by t gets 0, as their status at t is unknown.

Usage

```
graf_weights(truth, eval_time, censoring, trunc = 0.05)
```

Arguments

truth	A right-censored [survival::Surv()] object, one element per subject (see [surv_subject_truth()]).
eval_time	Numeric vector of evaluation times.
censoring	A [censoring_km()] object, estimated on the training data.
trunc	Lower bound for G , capping weights at '1 / trunc' (default 0.05, as in tidymodels' 'parsnip').

Details

The weights assume follow-up starts at time 0 for every subject: they do not correct for delayed entry (left truncation).

Value

A numeric matrix with one row per element of 'truth' and one column per 'eval_time'.

References

Graf E, Schmoor C, Sauerbrei W, Schumacher M (1999). Assessment and comparison of prognostic classification schemes for survival data. **Statistics in Medicine**, 18(17-18), 2529-2545.

See Also

[add_graf_weights()], [censoring_km()]

Examples

```
if (requireNamespace("survival", quietly = TRUE)) {
  train <- survival::Surv(c(2, 4, 6, 8), c(0, 1, 0, 1))
  test <- survival::Surv(c(5, 7), c(1, 0))
  graf_weights(test, eval_time = c(3, 6), censoring = censoring_km(train))
}
```

Description

A small family of `grep1()/%in%` wrappers as base functions and matching infix operators.

Usage

```
like(vector, pattern, ignore.case = FALSE, fixed = FALSE, perl = FALSE)
```

```
vector %like% pattern
```

```
ilike(vector, pattern)
```

```
vector %ilike% pattern
```

```
flike(vector, pattern)
```

```

vector %flike% pattern

plike(vector, pattern)

vector %plike% pattern

notin(x, table)

x %!in% table

x %notin% table

```

Arguments

pattern	Pattern to match (for the like family).
ignore.case, fixed, perl	Logical flags controlling how pattern is matched, as in grepl .
x, vector	A vector to test.
table	Vector of values to test against (for notin()/%notin%/!in%).

Value

A logical vector.

```
install.packages.no_lock
```

Install a Package, Bypassing an Existing Lock

Description

Install a Package, Bypassing an Existing Lock

Usage

```
install.packages.no_lock(packages)
```

Arguments

packages	Character vector of package name(s) to install.
----------	---

Value

Invisibly, NULL (as per [install.packages](#)).

Examples

```

## Not run:
install.packages.no_lock("dplyr")

## End(Not run)

```

is.int

*Is a Vector Composed of Integer-Valued Numbers***Description**

Is a Vector Composed of Integer-Valued Numbers

Usage

```
is.int(col)
```

Arguments

col A numeric vector.

Value

A single logical.

Examples

```
is.int(sample(-100:100, size = 500, replace = TRUE))
is.int(runif(500))
```

is_normal

*Test Whether a Vector Looks Normally Distributed***Description**

Runs two normality tests on the finite values of 'col':

Usage

```
is_normal(col, alpha = 0.1)
```

Arguments

col A numeric vector. 'NA', 'NaN', and infinite values are dropped.
alpha Significance level for 'distribution.test'.

Details

****Shapiro-Wilk**** ([stats::shapiro.test()]), for 3 to 5000 values. It is left out above 5000 rather than run on a random subsample, so results never depend on the random number generator. ****Lilliefors**** ([nortest::lillie.test()]), for 5 or more values: the Kolmogorov-Smirnov test corrected for estimating the mean and standard deviation from the same data. A plain KS test against 'pnorm(mean(x), sd(x))' gives p-values that are far too large.

Both tests are deterministic. In large samples they reject normality for departures too small to matter, so read them alongside a plot such as [distribution_plot()] or a normal Q-Q plot.

Value

A tibble with one row per test run: ‘statistic’, ‘p.value’, ‘method’, ‘distribution.test’ (‘TRUE’ when ‘p.value >= alpha’, i.e. normality is not rejected), ‘p_value_sig’, and ‘distribution’. It has no rows when neither test applies (fewer than 3 values, or all values equal).

References

Lilliefors HW (1967). On the Kolmogorov-Smirnov test for normality with mean and variance unknown. **Journal of the American Statistical Association**, 62(318), 399-402.

Examples

```
set.seed(1)
is_normal(rnorm(1000, mean = 5, sd = 3))
is_normal(runif(1000, min = 2, max = 4))
```

is_poisson	<i>Test Whether a Vector Looks Poisson-Distributed</i>
------------	--

Description

Runs a chi-squared goodness-of-fit test: observed counts (binned at roughly equal-probability quantiles of the fitted Poisson distribution, so expected counts stay adequate regardless of where the data sits) vs. Poisson-expected counts, with one degree of freedom subtracted for the estimated rate. The test is deterministic — no simulated comparison sample is drawn.

Usage

```
is_poisson(col)
```

Arguments

col A numeric vector.

Details

A Kolmogorov-Smirnov test is deliberately **not** used here: the KS statistic’s null distribution assumes a continuous CDF, and the Poisson distribution is discrete with real point masses, which inflates the KS statistic (and deflates its p-value) regardless of true fit — the chi-squared test is the standard, correctly-calibrated tool for discrete/count goodness-of-fit. [is_normal](#) uses the Lilliefors variant of the Kolmogorov-Smirnov test, which suits the continuous normal distribution.

Since the Poisson distribution’s support is the non-negative integers, this returns an empty tibble for vectors containing negative values or fewer than 2 observations.

Value

A one-row tibble of test results (empty if col isn’t valid count data, e.g. it has negative values).

Examples

```
is_poisson(rpois(n = 1000, lambda = 2))
is_poisson(runif(1000, min = 2, max = 4))
```

joint_model

*Joint Survival-Status-Time Model***Description**

Fits a coordinated trio of predictive models on clinical survival data:

1. **Survival Model** (coxnet): Penalized Cox proportional hazards model on `Surv(time, status)` or start/stop counting-process data `Surv(tstart, tstop, status)` via `coxnet` and `cv_coxnet`, correctly handling right-censoring and time-varying intervals.
2. **Status Model** (status): Binary event classification on `status ~ x` via penalized logistic regression (`glmnet`) or bagged trees (`baguette`), with optional probability calibration via **probably**.
3. **Time Model** (time): Continuous duration regression on `time ~ x` via penalized linear regression (`glmnet`) or bagged trees (`baguette`).
4. **Stack Ensemble** (stack): Optional regularized meta-learner blending candidate predictions via **stacks**.

Usage

```
joint_model(
  data,
  outcome = survival::Surv(time, status) ~ .,
  subject_id = NULL,
  engine = c("glmnet", "baguette", "stacks"),
  calibration = TRUE,
  mixture = 1,
  penalty = NULL,
  eval_time = NULL,
  ...
)
```

Arguments

data	A data frame containing the survival outcome and predictors.
outcome	A formula containing a Surv outcome, such as <code>Surv(time, status) ~ .</code> or <code>Surv(tstart, tstop, status) ~ ..</code>
subject_id	Optional character string specifying the subject identifier column for counting-process (start/stop) data.
engine	Character string specifying the modeling engine: "glmnet" (default), "baguette" (bagged decision trees), or "stacks" (stacked ensemble).
calibration	Logical; whether to fit a probability calibration model on the status predictions using probably (default TRUE).
mixture	Elastic net mixing parameter for glmnet models (default 1 for lasso).
penalty	Penalty value for glmnet . If NULL (default), tuned automatically via internal cross-validation using the Integrated Brier Score.
eval_time	Optional vector of evaluation times for dynamic survival probabilities. Defaults to deciles of uncensored event times.
...	Additional arguments passed to <code>cv_coxnet</code> or <code>glmnet</code> .

Details

This joint framework enables clinical researchers to investigate and contrast true survival modeling against naive classification and regression proxies as discussed in clinical literature (e.g., Rizopoulos 2015, PMC4503792).

Value

An S3 object of class `c("joint_model", "list")` with elements:

<code>coxnet_model</code>	Fitted penalized Cox proportional hazards model.
<code>status_model</code>	Fitted binary event classification model.
<code>time_model</code>	Fitted continuous follow-up duration model.
<code>stack_model</code>	Fitted stacks ensemble (if <code>engine = "stacks"</code>).
<code>calibration_model</code>	Probability calibration model from probably (if <code>calibration = TRUE</code>).
<code>components</code>	List of extracted outcome variables, formulas, and baseline hazard.
<code>engine</code>	Selected modeling engine.
<code>eval_time</code>	Evaluation horizons used for survival scoring.

See Also

[cv_joint_model](#), [nested_cv_joint_model](#), [cv_coxnet](#), [glmnet_IBS](#)

Examples

```
if (requireNamespace("glmnet", quietly = TRUE) &&
    requireNamespace("survival", quietly = TRUE)) {
  set.seed(42)
  df <- data.frame(
    time = stats::rexp(60, rate = 0.05),
    status = stats::rbinom(60, 1, 0.6),
    x1 = stats::rnorm(60),
    x2 = stats::rnorm(60)
  )
  fit <- joint_model(df, survival::Surv(time, status) ~ x1 + x2)
  print(fit)
  preds <- predict(fit, new_data = df[1:5, ])
}
```

join_pipelines

Join two computational pipeline graphs

Description

Relational union of two pipeline subgraphs using `tidygraph::graph_join(by = "name")`. Nodes with identical names are coalesced, and their directed dependency and output edges are combined without duplicating nodes.

Usage

```
join_pipelines(graph_a, graph_b, ...)
```

Arguments

graph_a	A list of FilePath objects, igraph, or tbl_graph.
graph_b	A list of FilePath objects, igraph, or tbl_graph.
...	Additional arguments passed to tidygraph::graph_join.

Value

A merged tbl_graph object.

keep_only

Keep Only Specified Objects in an Environment

Description

Removes every object in the caller's environment except the ones named in vector. Prompts for confirmation in interactive sessions unless .dontask = TRUE; proceeds without prompting in non-interactive sessions (scripts, R CMD check, knitr rendering), since readline() would otherwise hang there.

Usage

```
keep_only(vector, .dontask = FALSE)
```

Arguments

vector	Character vector of object names to keep.
.dontask	Logical (default FALSE); skip the confirmation prompt.

Value

Invisibly, NULL.

Examples

```
e <- new.env()
local({a <- 1; b <- 2; keep_only("a", .dontask = TRUE)}), envir = e)
ls(e)
```

km_summary_to_prism	<i>Convert a Kaplan-Meier Summary Table to a GraphPad Prism Survival Table</i>
---------------------	--

Description

Expands a summary-by-time table of event and censoring counts (as from `summary(survival::survfit(...))`) into Prism's survival data layout: one row per subject, an 'X' column of times, and one column per group holding 1 for an event, 0 for censored, and 'NA' for the other groups.

Usage

```
km_summary_to_prism(
  x,
  sheet = NULL,
  time_col = "time",
  strata_col = "strata",
  n_event_col = "n.event",
  n_censor_col = "n.censor",
  strata_levels = NULL,
  drop_time0 = TRUE,
  validate_totals = TRUE,
  out_xlsx = NULL,
  out_xlsx_paste = NULL
)
```

Arguments

<code>x</code>	A data frame, or a path to an <code>'xlsx'</code> , <code>'xls'</code> , <code>'csv'</code> , or <code>'txt'</code> file, with one row per time per group.
<code>sheet</code>	Sheet name or index when <code>'x'</code> is an Excel file.
<code>time_col</code> , <code>strata_col</code> , <code>n_event_col</code> , <code>n_censor_col</code>	Column names of the time, group, number of events, and number censored.
<code>strata_levels</code>	Optional group order for the output columns. Every group in the data must be listed.
<code>drop_time0</code>	Drop rows at time 0 (default <code>'TRUE'</code>).
<code>validate_totals</code>	If <code>'TRUE'</code> (default) and <code>'x'</code> has an <code>'n.risk'</code> column, warn when a group's <code>'n.risk'</code> at its earliest time differs from its total events plus censored.
<code>out_xlsx</code>	Optional path to write the table as Excel (requires <code>'writexl'</code>).
<code>out_xlsx_paste</code>	Optional path to write a copy with blanks instead of <code>'NA'</code> , convenient for pasting into Prism.

Value

A data frame with column `'X'` and one integer column per group, with a `"notes"` attribute describing the layout.

Examples

```
if (requireNamespace("survival", quietly = TRUE)) {
  fit <- survival::survfit(survival::Surv(time, status) ~ sex, data = survival::lung)
  s <- summary(fit, censored = TRUE)
  km <- data.frame(time = s$time, strata = s$strata, n.risk = s$n.risk,
    n.event = s$n.event, n.censor = s$n.censor)
  head(km_summary_to_prism(km))
}
```

list_stage_data_artifacts	<i>Inspect and categorize all on-disk data artifacts for a pipeline stage</i>
---------------------------	---

Description

Scans a directory for data files (.rds, .xlsx, .csv) and distinguishes primary pipeline handoff datasets from intermediate analysis tables.

Usage

```
list_stage_data_artifacts(
  stage_dir,
  stage = NA_character_,
  primary_patterns = c("^DATA\\.rds$", "^INDEX\\.rds$", "Index_Scores",
    "imputed_DATA")
)
```

Arguments

- stage_dir Directory path to scan (e.g. analysis/2024_09/EDA/data).
- stage Optional stage identifier to tag the artifacts with.
- primary_patterns Regex pattern identifying primary handoff datasets.

Value

A tibble of data artifacts with name, path, stage, role, size_kb, and mtime.

make_excel_names	<i>Generate Excel-Compatible Column Names</i>
------------------	---

Description

Formats column names into sanitized, human-readable strings suitable for Excel output headers.

Usage

```
make_excel_names(names)
```

Arguments

names A character vector of column names.

Value

Sanitized character vector.

Examples

```
make_excel_names(c("patient_id", "body_mass_index"))
```

manhattan_plot	<i>Manhattan Plot</i>
----------------	-----------------------

Description

Manhattan Plot

Usage

```
manhattan_plot(.data, var, log_p, alpha = 0.05, highlight_significant = TRUE)
```

Arguments

- .data A data frame, e.g. the output of [multiple_t_test](#).
- var Unquoted column of variable labels.
- log_p Unquoted column of $-\log_{10}(p.value)$.
- alpha Significance threshold for the reference line/shading (default 0.05).
- highlight_significant Logical (default TRUE); shade the region above the $p < \alpha$ line.

Value

A ggplot object.

Examples

```
mtcars |>
  dplyr::mutate(am = factor(am)) |>
  multiple_t_test(class = "am") |>
  manhattan_plot(var = var, log_p = log_p)
```

min_max_norm

*Min-Max Data Normalization***Description**

Normalizes numerical values into $[0, 1]$. For a data frame, matrix, or tibble, each numeric column is normalized independently against its own min/max, and **non-numeric columns are dropped** — matching the original implementation this was migrated from.

Usage

```
min_max_norm(obj)
```

Arguments

obj A numeric vector, matrix, data frame, or tibble.

Value

A normalized object: a numeric vector stays a vector; a matrix stays a matrix; a data frame or tibble is reduced to just its (now normalized) numeric columns.

Examples

```
min_max_norm(c(10, 20, 30, 40, 50))
min_max_norm(data.frame(x = c(1, 2, 3), y = c(-10, 0, 10), id = c("a", "b", "c")))
```

missforest_impute_by_mtry

*Assemble Imputed Columns From Their Best-‘mtry’ Runs***Description**

Given a sweep produced by `[missforest_sweep_mtry()]` and a table naming the winning ‘mtry’ for each column, pulls each column out of the run that won it. Different columns may come from different runs.

Usage

```
missforest_impute_by_mtry(sweep, best)
```

Arguments

sweep A named list of `[missforest_oob_by_mtry()]` results, named by ‘mtry’ as a character string (as `[missforest_sweep_mtry()]` builds it).

best A data frame with one row per column, containing at least the columns ‘column’ and ‘mtry’ – typically the ‘best’ element returned by `[missforest_sweep_mtry()]`.

Details

Runs are looked up by ****name**** (`'sweep[[as.character(mtry)]]'`), not by position. Positional lookup happens to work only when the swept `'mtry'` values are exactly `'1:n'`; it silently returns the wrong run for any other sweep, such as one starting above 1 or skipping values.

Value

A tibble of imputed columns, one per row of `'best'`.

See Also

`[missforest_sweep_mtry()]`

`missforest_oob_by_mtry`

Run 'missForest' at a Single 'mtry' and Report Variablewise OOB Error

Description

Fits `[missForest::missForest()]` once at a given `'mtry'`, returning both the completed data and a tidy, per-column out-of-bag error table. This is the unit of work swept over by `[missforest_sweep_mtry()]`.

Usage

```
missforest_oob_by_mtry(data, mtry, ntree = 100, maxiter = 10)
```

Arguments

<code>data</code>	A data frame or tibble containing the columns to impute. Any identifier or time columns should already be removed; see <code>[missforest_sweep_mtry()]</code> 's <code>'exclude'</code> argument. Columns that are entirely <code>'NA'</code> are refused – see Details.
<code>mtry</code>	Number of variables randomly sampled as candidates at each split. A single positive integer.
<code>ntree</code>	Number of trees per forest. Passed to <code>[missForest::missForest()]</code> . Default <code>'100'</code> .
<code>maxiter</code>	Maximum number of imputation iterations. Passed to <code>[missForest::missForest()]</code> . Default <code>'10'</code> .

Details

`'missForest'` requires every column to be numeric or a factor, so character columns are converted to factors here. The conversion is idempotent – a caller that has already converted its own characters sees no change – which is what lets this single function replace call sites that previously converted at different points.

Columns that are entirely `'NA'` are refused with an error. `[missForest::missForest()]` silently drops them, which both breaks the per-column error table this function returns and would quietly shrink the predictor set that swept `'mtry'` values index into. Such columns arise easily – a mapping-driven reader that fills absent source columns with a bare `'NA'` produces full-length logical columns that pass name-based schema checks. Drop or fill them before imputing.

Value

A list of two elements:

‘ximp’ A tibble of the completed (imputed) data.

‘oob_error’ A tibble with one row per column of ‘data’, and columns ‘column’, ‘error_type’ (“MSE” for numeric columns, “PFC” for factors), ‘error’, and ‘mtry’.

Reproducibility

‘missForest’ is stochastic. This function does not set a seed – seeding is the caller’s responsibility, and [missforest_sweep_mtry()] handles it in a parallel-safe way. See that function’s Reproducibility section.

See Also

[missforest_sweep_mtry()], [missforest_impute_by_mtry()]

Examples

```
if (requireNamespace("missForest", quietly = TRUE)) {
  set.seed(1)
  df <- data.frame(a = c(1, 2, NA, 4, 5), b = c(2, NA, 6, 8, 10))
  res <- missforest_oob_by_mtry(df, mtry = 1)
  res$oob_error
}
```

missforest_sweep_mtry *Impute a Data Frame by Sweeping ‘missForest’ Over ‘mtry’*

Description

Runs [missForest::missForest()] once for every candidate ‘mtry’, then keeps, ****for each column independently****, the imputation from whichever ‘mtry’ minimised that column’s out-of-bag error. Columns in the same data frame may therefore be imputed at different ‘mtry’ values.

Usage

```
missforest_sweep_mtry(
  data,
  exclude = NULL,
  mtry_values = NULL,
  ntree = 100,
  maxiter = 10,
  seed = TRUE,
  max_pct_missing = NULL,
  parallel = TRUE
)
```

Arguments

<code>data</code>	A data frame or tibble to impute. After ‘exclude’ is applied, any remaining column that is entirely ‘NA’ is refused with an error; see <code>[missforest_oob_by_mtry()]</code> ’s Details.
<code>exclude</code>	Character vector of columns to hold out of the imputation entirely – typically subject identifiers and time columns. They are removed before fitting and re-attached, unchanged, to the result. Default ‘NULL’.
<code>mtry_values</code>	Integer vector of ‘mtry’ values to sweep. Defaults to ‘1:(p - 1)’, where ‘p’ is the number of columns remaining after ‘exclude’.
<code>ntree, maxiter</code>	Passed to <code>[missforest_oob_by_mtry()]</code> .
<code>seed</code>	Controls parallel-safe random number generation; passed to <code>[furr::furr_options()]</code> . ‘TRUE’ (the default) draws reproducible streams from the current RNG state, so ‘set.seed()’ beforehand makes the whole sweep reproducible. A single integer seeds the sweep self-containedly, independent of ambient RNG state. See Reproducibility.
<code>max_pct_missing</code>	Optional proportion in ‘(0, 1]’. Columns missing a greater share than this are held out of the imputation and carried through unimputed, and are reported in ‘excluded_high_missing’. ‘NULL’ (the default) imputes every column regardless of how sparse it is. Imputing a column observed in a handful of rows manufactures values rather than recovering them, and nothing downstream can tell the difference; a threshold states that judgement as a rule that carries to the next data set, instead of naming the offending column inline. Note a threshold below ‘1’ subsumes the all-‘NA’ refusal, since an all-‘NA’ column exceeds every threshold.
<code>parallel</code>	Whether to evaluate the sweep with <code>[furr::future_map()]</code> . Default ‘TRUE’. When ‘FALSE’, runs sequentially via <code>[purrr::map()]</code> .

Value

A list of four elements:

- ‘**imp_data**’ A tibble of the imputed data, with ‘exclude’ columns re-attached and the original column order restored.
- ‘**oob_error**’ A tibble of every column’s OOB error at every swept ‘mtry’.
- ‘**best**’ The winning row per column – the subset of ‘oob_error’ minimising ‘error’ within each ‘column’, ties broken by first occurrence.
- ‘**excluded_high_missing**’ Columns held out by ‘max_pct_missing’, carried through unimputed. ‘character(0)’ when none were.

Reproducibility

‘missForest’ is stochastic, so an unseeded sweep is not reproducible. In parallel mode the seed ****must**** reach `[furr::future_map()]`’s ‘.options’; attaching ‘furr_options()’ to `[future::plan()]` instead is silently ignored (‘future’ warns about an unknown argument and proceeds unseeded). This function always passes ‘seed’ through to the ‘future_map()’ call.

Seeded this way, results are independent of the number of workers and of the ‘future’ plan, so a Windows ‘multisession’ run and a Linux ‘multicore’ run draw identical random streams. Note that identical streams do not by themselves guarantee bit-identical output across **architectures**: tree splits compare floating-point impurity sums, and differences in FMA contraction or BLAS between

x86_64 and arm64 can resolve a near-tie differently. Verify cross-architecture agreement rather than assuming it.

The sequential path (`parallel = FALSE`) seeds with `[set.seed()]` when `'seed'` is a number. It is internally reproducible but will not match the parallel path, which uses L'Ecuyer-CMRG streams.

Parallel plan

This function never calls `[future::plan()]` – the appropriate backend differs by machine, and choosing one here would override the caller's. Set a plan before calling, for example `'future::plan(future::multisession, workers = future::availableCores() - 1)'`.

See Also

`[missforest_oob_by_mtry()]`, `[missforest_impute_by_mtry()]`

Examples

```
if (requireNamespace("missForest", quietly = TRUE)) {
  df <- data.frame(
    id = 1:8,
    a = c(1, 2, NA, 4, 5, 6, NA, 8),
    b = c(2, NA, 6, 8, 10, 12, 14, NA),
    c = c(5, 4, 3, NA, 1, 2, 3, 4)
  )
  res <- missforest_sweep_mtry(df, exclude = "id", seed = 42, parallel = FALSE)
  res$best
}
```

missmap

Missingness Map

Description

Visualizes missingness across a data frame as a heatmap, either per row/column or aggregated by a grouping column.

Usage

```
missmap(
  df,
  by_column = NULL,
  na_list = NULL,
  row_order = FALSE,
  fill = c("auto", "binary", "count")
)
```

Arguments

<code>df</code>	A data frame or tibble.
<code>by_column</code>	Optional unquoted column name to aggregate missingness by (e.g. a site or subject id) instead of plotting every row individually. Missingness is summed per feature within each group. When <code>row_order = FALSE</code> , both the groups (x-axis) and the features (y-axis) are ordered by descending total missingness, exactly as in the per-row/column view.
<code>na_list</code>	Optional vector of additional values to treat as missing (beyond actual NA).
<code>row_order</code>	Logical; if FALSE (default) rows/features (or, in <code>by_column</code> mode, groups/features) are ordered by how much missingness they have.
<code>fill</code>	In <code>by_column</code> mode, how to color the tiles: "auto" (default) picks a discrete two-level "Missing"/"Present" fill (matching the per-row/column view) when every group has at most one contributing row – in which case the summed missingness count is always 0 or 1 and a count scale would be misleading – and otherwise falls back to a continuous "# missing" gradient. "binary" and "count" force one of those two behaviors regardless of group size. Ignored outside <code>by_column</code> mode.

Value

A ggplot object.

Examples

```
df <- data.frame(
  a = c(1, NA, 3, 4, NA),
  b = c(NA, NA, 3, 4, 5),
  site = c("A", "A", "B", "B", "B")
)
missmap(df)
missmap(df, by_column = site)
```

missranger_max_mtry *Largest 'mtry' 'missRanger' Will Accept For a Data Set*

Description

'missRanger' admits a **smaller** 'mtry' than 'missForest' on the same data, and the difference is easy to trip over: exceeding it surfaces as 'ranger's catch-all "User interrupt or internal error.", which names neither 'mtry' nor the real bound.

Usage

```
missranger_max_mtry(data)
```

Arguments

`data` A data frame or tibble, after any identifier columns have been excluded.

Details

‘missForest’ predicts every column from every other one, so its bound is simply `ncol(data) - 1`. ‘missRanger’ builds its predictor pool up over the first iteration instead. It starts with ‘completed’, the set of **fully-observed** columns, imputes the targets in increasing order of missingness, and adds each one to ‘completed’ as it goes. The first target therefore sees only the complete columns, and ‘mtry’ must fit *that* pool:

* With ‘k’ usable complete columns, the bound is ‘k’. * With none, the first target falls back to univariate imputation and the second sees a single predictor, so the bound is ‘1’ – no matter how wide the data is.

Constant columns are excluded from the count, matching ‘missRanger’'s own handling of them as features.

The practical consequence: a sweep is only informative when some columns are complete. That is the usual shape of the intended use – imputing a domain’s incomplete outputs against its fully-observed inputs – but a frame in which *every* column has missing values admits ‘mtry = 1’ alone, and there is nothing to sweep.

Value

A single integer: the largest admissible ‘mtry’, at least ‘1’.

See Also

`[missranger_sweep_mtry()]`

Examples

```
# Two complete columns -> mtry may be 1 or 2.
missranger_max_mtry(data.frame(a = c(1, NA, 3), b = c(1, 2, 3), c = c(4, 5, 6)))

# Every column has a gap -> only mtry = 1 is admissible.
missranger_max_mtry(data.frame(a = c(1, NA, 3), b = c(NA, 2, 3), c = c(4, 5, NA)))
```

```
missranger_oob_by_mtry
```

Run ‘missRanger’ at a Single ‘mtry’ and Report Per-Column OOB Error

Description

Fits `[missRanger::missRanger()]` once at a given ‘mtry’, returning both the completed data and a tidy, per-column out-of-bag error table. This is the unit of work swept over by `[missranger_sweep_mtry()]`, and the ‘ranger’-based counterpart to `[missforest_oob_by_mtry()]`.

Usage

```
missranger_oob_by_mtry(
  data,
  mtry,
  num.trees = 500,
  pmm.k = 0,
  maxiter = 10,
```

```

    seed = NULL,
    num.threads = 1,
    keep_forests = FALSE
  )

```

Arguments

<code>data</code>	A data frame or tibble containing the columns to impute. Any identifier or time columns should already be removed; see <code>[missranger_sweep_mtry()]</code> 's <code>'exclude'</code> argument. Columns that are entirely <code>'NA'</code> are refused – see Details.
<code>mtry</code>	Number of variables randomly sampled as candidates at each split. A single positive integer.
<code>num.trees</code>	Number of trees per forest. Passed to <code>[missRanger::missRanger()]</code> . Default <code>'500'</code> , that function's own default.
<code>pmm.k</code>	Number of predictive-mean-matching donors. <code>'0'</code> (the default) disables PMM and imputes with the forest prediction directly.
<code>maxiter</code>	Maximum number of imputation iterations. Default <code>'10'</code> .
<code>seed</code>	Integer seed passed to <code>[missRanger::missRanger()]</code> , or <code>'NULL'</code> (default) to leave it unseeded. See Reproducibility.
<code>num.threads</code>	Threads used by <code>'ranger'</code> <i>*within*</i> one fit. Defaults to <code>'1'</code> ; see the Threading section, which explains why the <code>'ranger'</code> default is the wrong one here.
<code>keep_forests</code>	Whether to retain the fitted forests on the returned object. Default <code>'FALSE'</code> . See Details.

Details

****Error rows cover imputed columns only.**** `'missForest'` run variablewise reports an OOB error for every column; `'missRanger'` fits forests only for columns that actually have missing values, so complete columns get no row. `[missranger_sweep_mtry()]` carries those columns through unchanged.

****Errors are not comparable across engines.**** `'missForest'` reports raw MSE and PFC. `'missRanger'` reports `'ranger'`'s scaled OOB prediction error – `'1 - R^2'` for numeric targets and classification error for categorical ones – where roughly `'1'` means no better than predicting the mean. Compare `'mtry'` values within one engine; do not compare an `'error'` column from one against the other.

****All-`'NA'` columns are refused.**** `'missRanger'` does not drop them the way `'missForest'` does; it silently returns them still entirely `'NA'`, producing a frame that looks imputed but is not. Such columns arise easily – a mapping-driven reader that fills absent source columns with a bare `'NA'` yields full-length logical columns that pass name-based schema checks.

****`'keep_forests'` is the seam for applying to new data.**** `'missForest'` discards its forests, which is why it cannot be applied to a later batch. `'missRanger'` can retain them, making `'predict()'` on new data possible. This function only exposes the option; a fit/apply pair built on it is separate work.

Value

A list of two elements, plus a third when `'keep_forests = TRUE'`:

`'ximp'` A tibble of the completed (imputed) data.

`'oob_error'` A tibble with columns `'column'`, `'error_type'`, `'error'`, and `'mtry'` – one row per **imputed** column. See Details.

`'fit'` The `'missRanger'` object, when `'keep_forests = TRUE'`.

Threading

‘ranger’ defaults to using every available core. Inside a parallel ‘mtry’ sweep that nests one thread pool inside another and oversubscribes the machine badly – workers times cores threads competing for the same cores. ‘num.threads’ therefore defaults to ‘1’ here, leaving parallelism to the sweep. Raise it only when running a single fit on an otherwise idle box.

Reproducibility

Unlike ‘missForest’, ‘missRanger’ takes its own ‘seed’ and uses it to seed the fit directly, so a seeded fit is fully determined by ‘seed’ regardless of ambient RNG state, worker count, or ‘future’ plan. That makes seeded ‘missranger_*’ sweeps reproducible by construction rather than by careful stream management. The architecture caveat still applies: identical seeds do not guarantee bit-identical results across x86_64 and arm64, because split selection compares floating-point impurity sums.

See Also

[missranger_sweep_mtry()], [missforest_oob_by_mtry()]

Examples

```
if (requireNamespace("missRanger", quietly = TRUE)) {
  df <- data.frame(
    a = c(1, 2, NA, 4, 5, 6, 7, 8),
    b = c(2, NA, 6, 8, 10, 12, 14, 16),
    c = c(5, 4, 3, NA, 1, 2, 3, 4)
  )
  res <- missranger_oob_by_mtry(df, mtry = 1, num.trees = 50, seed = 1)
  res$oob_error
}
```

missranger_sweep_mtry *Impute a Data Frame by Sweeping ‘missRanger’ Over ‘mtry’*

Description

Runs [missRanger::missRanger()] once for every candidate ‘mtry’, then keeps, **for each column independently**, the imputation from whichever ‘mtry’ minimised that column’s out-of-bag prediction error. The ‘ranger’-based counterpart to [missforest_sweep_mtry()], with the same return shape so the two engines can be compared on one data set.

Usage

```
missranger_sweep_mtry(
  data,
  exclude = NULL,
  mtry_values = NULL,
  num.trees = 500,
  pmm.k = 0,
  maxiter = 10,
  seed = NULL,
```

```

    num.threads = 1,
    max_pct_missing = NULL,
    parallel = TRUE
  )

```

Arguments

<code>data</code>	A data frame or tibble to impute. After ‘exclude’ is applied, any remaining column that is entirely ‘NA’ is refused with an error.
<code>exclude</code>	Character vector of columns to hold out of the imputation entirely – typically subject identifiers and time columns. They are removed before fitting and re-attached, unchanged, to the result. Default ‘NULL’.
<code>mtry_values</code>	Integer vector of ‘mtry’ values to sweep. Defaults to ‘1:missranger_max_mtry(data)’. Note this is not ‘1:(p - 1)’, the ‘missForest’ bound – see [missranger_max_mtry()] for why, and expect a shorter sweep than [missforest_sweep_mtry()] performs on the same data.
<code>num.trees</code> , <code>pmm.k</code> , <code>maxiter</code> , <code>num.threads</code>	Passed to [missranger_oob_by_mtry()].
<code>seed</code>	Integer seed passed to every fit in the sweep, or ‘NULL’ (default) for an unseeded sweep. See Reproducibility.
<code>max_pct_missing</code>	Optional proportion in ‘(0, 1]’. Columns missing a greater share than this are held out of the imputation and carried through unimputed, and are reported in ‘excluded_high_missing’. ‘NULL’ (the default) imputes every column regardless of how sparse it is. Imputing a column observed in a handful of rows manufactures values rather than recovering them, and nothing downstream can tell the difference; a threshold states that judgement as a rule that carries to the next data set, instead of naming the offending column inline. Note a threshold below ‘1’ subsumes the all-‘NA’ refusal, since an all-‘NA’ column exceeds every threshold.
<code>parallel</code>	Whether to evaluate the sweep with [furrr::future_map()]. Default ‘TRUE’. When ‘FALSE’, runs sequentially via [purrr::map()].

Value

A list of four elements, matching [missforest_sweep_mtry()]:

- ‘**imp_data**’ A tibble of the imputed data, with ‘exclude’ columns re-attached and the original column order restored.
- ‘**oob_error**’ A tibble of every imputed column’s OOB error at every swept ‘mtry’.
- ‘**best**’ The winning row per imputed column.
- ‘**excluded_high_missing**’ Columns held out by ‘max_pct_missing’, carried through unimputed. ‘character(0)’ when none were.

Reproducibility

The same ‘seed’ is passed to every fit in the sweep. That is deliberate: holding the seed fixed across ‘mtry’ values is a common-random-numbers comparison, so differences in OOB error reflect ‘mtry’ rather than sampling noise between runs. Because ‘missRanger’ seeds itself, a seeded sweep is reproducible independently of the ‘future’ plan and the number of workers – the property a mixed Windows/Linux/arm64 fleet needs – without depending on L’Ecuyer stream management. ‘furrr’'s own parallel-safe seeding is still enabled, which covers the unseeded case.

Parallel plan

This function never calls `[future::plan()]` – the appropriate backend differs by machine, and choosing one here would override the caller's. Set a plan before calling. Note that `'num.threads'` defaults to `'1'` so that `'ranger'`'s own threading does not nest inside the sweep; see `[missranger_oob_by_mtry()]`'s Threading section.

See Also

`[missforest_sweep_mtry()]` for the `'missForest'` engine.

Examples

```
if (requireNamespace("missRanger", quietly = TRUE)) {
  df <- data.frame(
    id = 1:12,
    a = c(1, 2, NA, 4, 5, 6, NA, 8, 9, 10, 11, NA),
    b = c(2, 4, 6, NA, 10, 12, 14, NA, 18, 20, NA, 24),
    c = c(5, 4, 3, NA, 1, 2, 3, 4, 5, NA, 3, 2)
  )
  res <- missranger_sweep_mtry(
    df, exclude = "id", num.trees = 50, seed = 42, parallel = FALSE
  )
  res$best
}
```

multiple_t_test

Multiple T-Tests Against One Classifier

Description

Runs [single_t_test](#) for every (or a chosen set of) numeric column against a single binary classifier.

Usage

```
multiple_t_test(
  .data,
  .var_list = names(dplyr::select(.data, dplyr::where(is.numeric))),
  .class,
  alternative = "two.sided",
  conf.level = 0.95,
  ...
)
```

Arguments

<code>.data</code>	A data frame or tibble.
<code>.var_list</code>	Character vector of column names to test (default: every numeric column in <code>.data</code>).
<code>.class</code>	Name (string) of a binary classification column.

alternative	One of "two.sided" (default), "greater", or "less".
conf.level	Confidence level for the interval.
...	Additional arguments passed to single_t_test (and on to t.test) — e.g. <code>paired</code> and <code>.id</code> .

Value

A tibble with one row per tested variable.

Examples

```
mtcars |>
  dplyr::mutate(am = factor(am)) |>
  multiple_t_test(class = "am")
```

my_summary_table	<i>Summary Table Function</i>
------------------	-------------------------------

Description

Summary Table Function

Usage

```
my_summary_table(
  data,
  var,
  na.rm = FALSE,
  data.output = TRUE,
  table.output = FALSE
)
```

Arguments

data	a dataframe or tibble
var	variable to summarize
na.rm	should NA values be removed, TRUE or FALSE, default is FALSE
data.output	should the data be output, TRUE or FALSE, default is TRUE
table.output	should the table be output, TRUE or FALSE, default is FALSE

Value

either a dataframe or readable table

Examples

```
mtcars |>
  dplyr::group_by(cyl) |>
  my_summary_table(mpg)

mtcars |>
  dplyr::group_by(cyl) |>
  my_summary_table(mpg, table.output = TRUE)
```

 nested_cv_coxnet

Nested Cross-Validation of a Penalized Cox Model

Description

Estimates how well the whole tuning procedure of `[cv_coxnet()]` generalizes. For each outer split of an `[rsample::nested_cv()]` object: ‘mixture’ and ‘penalty’ are chosen by `[cv_coxnet()]` on that split’s inner resamples; the model is refit on the outer analysis set with those settings; and the refit is scored on the outer assessment set, which played no part in tuning.

Usage

```
nested_cv_coxnet(
  object,
  preprocessor,
  subject_id = NULL,
  group = NULL,
  rule = c("min", "1se"),
  mixture = 1,
  penalty = NULL,
  metrics = NULL,
  eval_time = NULL,
  metric = NULL,
  covariates = c("path", "baseline"),
  trunc = 0.05,
  parallel = FALSE,
  importance = c("none", "loco_mp"),
  ...
)

## S3 method for class 'nested_cv_coxnet'
collect_metrics(x, ..., summarize = TRUE)
```

Arguments

<code>object</code>	An <code>[rsample::nested_cv()]</code> object.
<code>preprocessor</code>	A formula with a ‘Surv()’ outcome, or a <code>[recipes::recipe()]</code> whose outcome is a ‘Surv’ column.
<code>subject_id, group</code>	Column names, as in <code>[cv_coxnet()]</code> .

rule	Use <code>"min"</code> (<code>'lambda.min'</code>) or <code>"1se"</code> (<code>'lambda.1se'</code>) from each inner cross-validation.
mixture	Numeric vector of elastic-net mixing values to try.
penalty	Optional decreasing penalty path; by default <code>glmnet</code> 's path for each <code>'mixture'</code> .
metrics	A <code>[yardstick::metric_set()]</code> of survival metrics. Defaults to the integrated Brier score, concordance, time-specific Brier score, and time-dependent ROC AUC.
eval_time	Evaluation times, shared by all splits. Defaults to deciles of the event times in the full data.
metric	Name of the metric used to choose settings; the first metric in <code>'metrics'</code> by default.
covariates	<code>"path"</code> or <code>"baseline"</code> ; see Details.
trunc	Lower bound for the censoring probability in the Graf weights.
parallel	If <code>'TRUE'</code> , fit folds with <code>[furrr::future_map()]</code> ; set a <code>[future::plan()]</code> first.
importance	Method for evaluating feature importance on the outer analysis sets. Options are <code>"none"</code> (default) or <code>"loco_mp"</code> (Leave-One-Covariate-Out with MiniPatch ensembles).
...	Further arguments passed to <code>[glmnet::glmnet()]</code> , such as <code>'cox.ties'</code> or <code>'nlambda'</code> .
x	A <code>'nested_cv_coxnet'</code> object.
summarize	For <code>'collect_metrics()'</code> : <code>'TRUE'</code> for the mean and standard error over outer splits, <code>'FALSE'</code> for each split's metrics.

Details

Build `'object'` with grouped resampling at both levels, e.g. `'nested_cv(data, outside = group_vfold_cv(group = subject_id), inside = group_vfold_cv(group = subject_id))'`, so a subject's start/stop rows stay together.

The outer metrics describe the procedure, not one final model: the chosen settings can differ between outer splits. To fit a final model, run `[cv_coxnet()]` on all the data.

Value

A tibble, of class `'nested_cv_coxnet'`, with one row per outer split: `'id'`, the chosen `'mixture'` and `'penalty'`, `'metrics'` (outer assessment-set metrics), `'coefs'` (coefficients of the refit), and `'inner'` (the inner cross-validation's summarized metrics). If `'importance = "loco_mp"`, includes `'importance'` with tidy LOCO-MP statistical inference. `'tune::collect_metrics()'` averages `'metrics'` over outer splits.

See Also

`[cv_coxnet()]`, `[coxnet()]`, `[cbe_loco_mp_coxnet()]`

Examples

```
if (requireNamespace("glmnet", quietly = TRUE) &&
    requireNamespace("survival", quietly = TRUE) &&
    requireNamespace("rsample", quietly = TRUE) &&
    requireNamespace("yardstick", quietly = TRUE)) {
  set.seed(1)
  long <- do.call(rbind, lapply(1:90, function(i) {
    k <- sample(1:3, 1)
```

```

stops <- cumsum(stats::runif(k, 2, 6))
risk <- stats::rnorm(1)
data.frame(subject_id = i, tstart = c(0, utils::head(stops, -1)), tstop = stops,
           status = c(rep(0, k - 1), stats::rbinom(1, 1, stats::plogis(risk))),
           x1 = risk + stats::rnorm(k, sd = 0.2), x2 = stats::rnorm(k), x3 = stats::rnorm(k))
)))
folds <- rsample::nested_cv(
  long,
  outside = rsample::group_vfold_cv(group = subject_id, v = 3),
  inside = rsample::group_vfold_cv(group = subject_id, v = 3)
)
res <- nested_cv_coxnet(
  folds, survival::Surv(tstart, tstop, status) ~ x1 + x2 + x3,
  subject_id = "subject_id", mixture = c(0.5, 1),
  metrics = yardstick::metric_set(yardstick::brier_survival_integrated,
                                   yardstick::concordance_survival),
  nlambda = 20, cox.ties = "breslow"
)
res
}

```

nested_cv_joint_model *Nested Cross-Validation for Joint Models*

Description

Implements two-layer nested cross-validation on an [nested_cv](#) object to evaluate the joint model pipeline without tuning leakage.

Usage

```

nested_cv_joint_model(
  object,
  outcome = survival::Surv(time, status) ~ .,
  subject_id = NULL,
  engine = c("glmnet", "baguette", "stacks"),
  calibration = TRUE,
  parallel = FALSE,
  ...
)

```

Arguments

<code>object</code>	An <code>rsample::nested_cv</code> object.
<code>outcome</code>	Survival formula with a <code>Surv()</code> outcome.
<code>subject_id</code>	Optional subject identifier column.
<code>engine</code>	Modeling engine: "glmnet", "baguette", or "stacks".
<code>calibration</code>	Logical; whether to calibrate status predictions (default TRUE).
<code>parallel</code>	Logical; whether to run outer splits in parallel.
<code>...</code>	Additional arguments.

Value

An S3 object of class `c("nested_cv_joint_model", "tbl_df")`.

Note

Counting-process `Surv(start, stop, event)` outcomes are not yet supported here (only in [joint_model](#) itself): scoring needs one row per subject, but predictions are one row per interval. Use `Surv(time, status)` outcomes for cross-validation.

`normalize_safely`

Normalize File Paths Without Failing

Description

A vectorized `[normalizePath()]` that never errors or warns: blank and ‘NA’ inputs become ‘NA’, and a path that cannot be normalized comes back as it was. Paths need not exist. Separators are forward slashes on every platform.

Usage

```
normalize_safely(x)
```

Arguments

`x` Character vector of paths.

Value

A character vector of normalized paths, named by the inputs.

See Also

`[scan_data_io()]`

Examples

```
normalize_safely(c("data/../../data/file.csv", "", NA))
```

one_vs_rest_t_test	<i>One-vs-Rest T-Tests Across a Multi-Level Factor</i>
--------------------	--

Description

For a classifier with more than two levels, dichotomizes each level against all others in turn and runs [single_t_test](#).

Usage

```
one_vs_rest_t_test(  
  .data,  
  .var,  
  .class,  
  .id = NULL,  
  paired = FALSE,  
  alternative = "two.sided",  
  conf.level = 0.95  
)
```

Arguments

<code>.data</code>	A data frame or tibble.
<code>.var</code>	Name (string) of the continuous column to test.
<code>.class</code>	Name (string) of a classification column (2+ levels).
<code>.id</code>	Name (string, optional) of a subject/record identifier column; see single_t_test . Only meaningful when <code>paired = TRUE</code> .
<code>paired</code>	Logical; paired t-test.
<code>alternative</code>	One of "two.sided" (default), "greater", or "less".
<code>conf.level</code>	Confidence level for the interval.

Value

A tibble with one row per level of `.class`.

Examples

```
one_vs_rest_t_test(iris, "Sepal.Length", "Species")
```

package_deliverables	<i>Package Pipeline Deliverables into a Structured Zip Archive</i>
----------------------	--

Description

Traverses computational pipeline objects (or explicit file lists), gathers all rendered reports and data deliverables, verifies existence on disk, organizes them into stage-prefixed folders, and builds a distribution ZIP file.

Usage

```
package_deliverables(
  pipeline_objects,
  output_formats = c("pdf", "docx"),
  data_deliverables = list(),
  zip_path = NULL,
  include_pipeline_graph = TRUE
)
```

Arguments

pipeline_objects	A list of FilePath / FileOutputs objects (e.g. from the compute graph).
output_formats	Character vector of report formats to include ("pdf", "docx", "html", or "all").
data_deliverables	Optional list or vector of FilePath objects or file paths for data spreadsheets/RDS files.
zip_path	Destination path for the generated zip archive. If NULL, defaults to a times-tamped archive.
include_pipeline_graph	Logical; whether to include static/interactive pipeline DAG graphs (default: TRUE).

Value

Absolute path to the generated zip file.

parse_here_call_vec	<i>Resolve 'here::here()' Calls Found in Code Text</i>
---------------------	--

Description

For each string, finds the first 'here::here(...)' call and evaluates it with its quoted string arguments, giving the path the call builds in the current project. Calls whose arguments are not all string literals resolve from the literals alone.

Usage

```
parse_here_call_vec(x)
```

Arguments

`x` Character vector of code lines.

Value

A character vector the length of `'x'`: the resolved path, or `'NA'` where a line has no `'here::here()'` call with quoted arguments.

See Also

`[scan_data_io()], [find_code()]`

Examples

```
parse_here_call_vec(c("x <- readRDS(here::here('data', 'x.rds'))", "no call"))
```

pca_biplot

PCA Loadings Biplot

Description

A classic PCA biplot: the observation scores (`pca_model$x`) overlaid with variable loading vectors (`pca_model$rotation`) drawn as labeled arrows from the origin. Observations can optionally be colored by a categorical grouping variable with 95% concentration ellipses.

Usage

```
pca_biplot(
  pca_model,
  x = 1,
  y = 2,
  group = NULL,
  ellipse = FALSE,
  title = "PCA Biplot",
  percent = FALSE
)
```

Arguments

<code>pca_model</code>	A prcomp object.
<code>x, y</code>	Which principal components to plot on the x/y axes (default 1, 2).
<code>group</code>	Optional categorical vector (e.g. <code>iris\$Species</code>) to color observations by.
<code>ellipse</code>	Logical (default FALSE); if TRUE and group is provided, draws 95% confidence/concentration ellipses around each group.
<code>title</code>	Optional plot title (default "PCA Biplot").
<code>percent</code>	Logical (default FALSE); if TRUE, appends percent variance explained to axis labels.

Details

Loading vectors are unit-scale by construction and would be invisible next to the score cloud if plotted as-is, so they are rescaled so that their maximum extent is 80% of the score cloud's maximum extent (a standard biplot convention) before being drawn.

Value

A ggplot object.

Examples

```
pca_model <- prcomp(iris[, 1:4], center = TRUE, scale. = TRUE)
pca_biplot(pca_model, group = iris$Species, ellipse = TRUE)
```

pca_eqns	<i>PCA Equations</i>
----------	----------------------

Description

Writes out each principal component as a linear equation in the original features (abbreviated f1, f2, ... — see the returned labels table for what each abbreviation means).

Usage

```
pca_eqns(PC_mod, precision = 3)
```

Arguments

PC_mod	A prcomp object.
precision	Digits to round loadings to (default 3).

Value

A list with two tibbles: eqns (one row per component, with its equation as text) and labels (feature-number-to-name key).

Examples

```
pca_model <- prcomp(mtcars, center = TRUE, scale. = TRUE)
pca_eqns(pca_model)$eqns
pca_eqns(pca_model)$labels
```

`pca_feature_loading_heatmap`*PCA Feature-Loading Heatmap*

Description

Heatmap of each original feature's loading onto each principal component.

Usage

```
pca_feature_loading_heatmap(pca_model)
```

Arguments

`pca_model` A [prcomp](#) object.

Value

A ggplot object.

Examples

```
pca_model <- prcomp(mtcars, center = TRUE, scale. = TRUE)
pca_feature_loading_heatmap(pca_model)
```

`pca_loading_diff`*Difference in PCA Loadings Between Two Fits*

Description

Compares the variable loadings of two independently-fit [prcomp](#) objects on the same set of variables (e.g. the same domain's data fit at baseline vs. at a later timepoint), matching components positionally (both fits' PC1, both fits' PC2, ...).

Usage

```
pca_loading_diff(pca_baseline, pca_comparison, n_components = NULL)
```

Arguments

`pca_baseline` A [prcomp](#) object treated as the reference.

`pca_comparison` A [prcomp](#) object to compare against `pca_baseline`, fit on the same (or overlapping) variables.

`n_components` Number of leading components to compare. Defaults to NULL, meaning all components shared by both fits.

Details

PCA loading vectors are only unique up to sign: a component can flip orientation between two otherwise-equivalent fits without changing the pattern it represents. Naively differencing loadings would then show a spuriously large change (up to roughly double the loading) for a component that hasn't meaningfully changed at all. To avoid this, for each shared component `pca_comparison`'s loading vector is sign-aligned to `pca_baseline`'s: it is flipped (multiplied by -1) if doing so reduces the total absolute difference across variables relative to leaving it as-is. The difference is then computed as (sign-aligned comparison) minus baseline.

Variables are matched by name (the rownames of `$rotation`). If the two fits were built on different variable sets, only the intersection is used; no error is raised.

Value

A tibble with one row per shared variable (feature column) and one column per compared component (PC1, PC2, ...) holding the sign-aligned difference (comparison minus baseline).

Examples

```
set.seed(1)
baseline <- prcomp(mtcars, center = TRUE, scale. = TRUE)
comparison <- prcomp(mtcars[sample(nrow(mtcars)), ], center = TRUE, scale. = TRUE)
pca_loading_diff(baseline, comparison)
```

pca_loading_diff_heatmap

Heatmap of PCA Loading Differences Between Two Fits

Description

Renders the output of `pca_loading_diff` as a heatmap (feature by component), using the same visual language as `pca_feature_loading_heatmap`: a diverging fill scale centered at zero, so components/variables with little sign-aligned change are white and larger changes in either direction stand out in blue or red.

Usage

```
pca_loading_diff_heatmap(pca_baseline, pca_comparison, n_components = NULL)
```

Arguments

<code>pca_baseline</code>	A <code>prcomp</code> object treated as the reference.
<code>pca_comparison</code>	A <code>prcomp</code> object to compare against <code>pca_baseline</code> , fit on the same (or overlapping) variables.
<code>n_components</code>	Number of leading components to compare. Defaults to <code>NULL</code> , meaning all components shared by both fits.

Value

A ggplot object.

Examples

```
set.seed(1)
baseline <- prcomp(mtcars, center = TRUE, scale. = TRUE)
comparison <- prcomp(mtcars[sample(nrow(mtcars)), ], center = TRUE, scale. = TRUE)
pca_loading_diff_heatmap(baseline, comparison)
```

pca_percent_var_explained

Percent Variance Explained by Each Principal Component

Description

Percent Variance Explained by Each Principal Component

Usage

```
pca_percent_var_explained(pca_model)
```

Arguments

pca_model A [prcomp](#) object.

Value

A ggplot object showing per-component and cumulative variance explained.

Examples

```
pca_model <- prcomp(mtcars, center = TRUE, scale. = TRUE)
pca_percent_var_explained(pca_model)
```

pca_plot

Plot a PCA Fit

Description

One entry point to this package's PCA plots.

Usage

```
pca_plot(
  pca_model,
  type = c("variance", "heatmap", "bi", "biplot", "scree", "circle"),
  ...
)
```

Arguments

pca_model	A prcomp object. (Not 'x', which would capture the 'x' component argument of 'type = "bi"' and "biplot".)
type	One of "variance" (pca_percent_var_explained), "heatmap" (pca_feature_loading_heatmap), "bi" (plot_pca_bi), "biplot" (pca_biplot), "scree" (pca_scree_plot), or "circle" (pca_variables_plot).
...	Passed on to the underlying plot function (needed for type = "bi", which requires newdata and column; and optionally used by type = "bi" or "biplot" to pick components with x/y).

Details

This is an ordinary function rather than a 'plot()' method: **stats** already registers 'plot()' for 'prcomp' objects (a scree plot), and a package that replaces another package's method for a class it doesn't own changes 'plot()' for everyone who loads it. Call 'pca_plot(x, type =)', or [stats::screeplot()] for the base scree plot.

Value

A ggplot object.

Examples

```
pca_model <- prcomp(mtcars, center = TRUE, scale. = TRUE)
pca_plot(pca_model, type = "variance")
pca_plot(pca_model, type = "biplot", x = 1, y = 3)
```

pca_scree_plot	<i>PCA Scree Plot</i>
----------------	-----------------------

Description

Scree plot displaying the eigenvalues or proportion of variance explained by each principal component, with an optional horizontal reference line at eigenvalue = 1 (Kaiser-Guttman criterion).

Usage

```
pca_scree_plot(
  pca_model,
  metric = c("eigenvalue", "variance"),
  kaiser = TRUE,
  title = "PCA Scree Plot"
)
```

Arguments

pca_model	A prcomp object or output from proc_pca .
metric	Which metric to plot: "eigenvalue" (default, with Kaiser line at 1) or "variance" (percent of variance explained).
kaiser	Logical (default TRUE); draw horizontal dashed line at eigenvalue = 1.
title	Optional plot title.

Value

A ggplot object.

Examples

```
pca_model <- prcomp(iris[, 1:4], center = TRUE, scale. = TRUE)
pca_scree_plot(pca_model)
```

pca_variables_plot	<i>PCA Variable Correlation Circle</i>
--------------------	--

Description

Plots the projection of variables onto the unit circle for two principal components, showing correlations between original variables and principal components.

Usage

```
pca_variables_plot(
  pca_model,
  x = 1,
  y = 2,
  title = "Variables Correlation Circle"
)
```

Arguments

pca_model	A prcomp object or output from proc_pca .
x, y	Which principal components to plot (default 1, 2).
title	Optional plot title (default "Variables Correlation Circle").

Value

A ggplot object.

Examples

```
pca_model <- prcomp(iris[, 1:4], center = TRUE, scale. = TRUE)
pca_variables_plot(pca_model)
```

`pdf_to_rtf`*Convert a PDF's Text to Rich Text Format (RTF)*

Description

Extracts the text layer of a PDF with `[pdftools::pdf_text()]` and writes it to an RTF file in a monospaced font, one RTF page per PDF page, so the column alignment of text tables (such as SAS listings) is kept.

Usage

```
pdf_to_rtf(  
  pdf,  
  rtf = sub("\\.pdf$", ".rtf", pdf, ignore.case = TRUE),  
  font_size = 10,  
  overwrite = TRUE  
)
```

Arguments

<code>pdf</code>	Path to the input ‘.pdf’ file.
<code>rtf</code>	Path of the ‘.rtf’ file to write. Defaults to ‘pdf’ with its extension replaced by ‘.rtf’.
<code>font_size</code>	Font size in points.
<code>overwrite</code>	If ‘FALSE’, error when ‘rtf’ already exists.

Details

Only text is converted: images, fonts, and vector graphics are not. A scanned PDF with no text layer gives empty pages. Non-ASCII characters are written as RTF Unicode escapes, so the file is plain ASCII and opens the same in any locale.

Value

The path to ‘rtf’, invisibly.

See Also

`[create_toc_from_sas_pdf()]`

Examples

```
pdf <- system.file("templates", "example.pdf", package = "TempleCBE")  
rtf <- pdf_to_rtf(pdf, tempfile(fileext = ".rtf"))  
readLines(rtf, n = 3)
```

pformat

*Format p-values for Biostatistical and Clinical Reporting***Description**

Formats numeric p-values into clean, publication-ready strings (e.g. "p = 0.024", "p < 0.001", or "0.024").

Usage

```
pformat(p, accuracy = 0.001, add_p = TRUE, digits = NULL)
```

```
cbe_pformat(p, accuracy = 0.001, add_p = TRUE, digits = NULL)
```

Arguments

p	Numeric p-value or vector of p-values.
accuracy	Numeric precision threshold passed to label_pvalue (default 0.001).
add_p	Logical; if TRUE (the default), prepends "p = ", "p < ", or "p > ". If FALSE, returns only the formatted numeric string.
digits	Optional integer number of decimal places (overrides accuracy as $10^{(-digits)}$).

Value

A character vector of formatted p-values.

Examples

```
pformat(0.0241)
pformat(0.0001)
pformat(0.852, add_p = FALSE)
```

pipeline_config

*Get or set global pipeline configuration options***Description**

Decouples project-specific parameters (study name, stage display labels, stage color palettes) from core pipeline mechanics and classes.

Usage

```
pipeline_config(
  study_name = NULL,
  stage_labels = NULL,
  stage_colors = NULL,
  default_formats = NULL
)
```

Arguments

- study_name Character string for the main study title.
- stage_labels Named character vector mapping raw stage identifiers to human-readable labels.
- stage_colors Named list or character vector mapping stage identifiers to hex colors.
- default_formats Optional character vector of default output formats for pipeline renderers (e.g. `c("pdf", "docx")`). When set, this overrides the package default across `render` and `create_qmd_renderer`.

Value

A list containing the current configuration options.

pipeline_summary	<i>Produce an executive summary of the computational pipeline</i>
------------------	---

Description

Summarizes total nodes, edges, staleness status, artifact role counts, and a stage-by-role breakdown table.

Usage

```
pipeline_summary(object)

summarize_pipeline(object)
```

Arguments

- object A `tbl_graph`, `igraph`, or list of `FilePath` objects.

Value

Invisibly returns a tibble of the stage breakdown.

pi_anonymizer	<i>PI Name Generator and Anonymizer</i>
---------------	---

Description

Provides functions to procedurally generate synthetic last names, generate pseudonym tokens (sequential or HMAC-based), and anonymize investigator names (supporting both single scalar and vectorized inputs).

plot.features_percent_miss

Plot method for features_percent_miss objects

Description

Creates a bar chart visualizing feature missingness percentages for objects produced by [features_percent_miss](#).

Generates a horizontal bar chart visualizing the percentage of missing values per column.

Usage

```
## S3 method for class 'features_percent_miss'
plot(x, top_n = NULL, ...)

plot_features_percent_miss(data, top_n = NULL)
```

Arguments

x	An object of class <code>features_percent_miss</code> .
top_n	Optional integer to limit to the top N features with highest missingness.
...	Additional arguments (currently unused).
data	A data frame, tibble, or the output of <code>'features_percent_miss()'</code> .

Value

A ggplot object representing feature missingness.

A ggplot object.

Examples

```
res <- features_percent_miss(mtcars)
plot(res)

df <- data.frame(a = c(1, 2, NA, 4), b = c(NA, NA, 3, 4))
plot_features_percent_miss(df)
```

plot_categorical_associations

Plot Categorical Association Matrix Using corrplot

Description

Computes pairwise associations (Cramer's V or -log10 p-values) across all categorical/factor columns in a data frame and renders an association matrix plot via **corrplot**. This complements [correlation_plot](#) for numeric variables.

Usage

```
plot_categorical_associations(  
  data,  
  cols = NULL,  
  method = c("cramer_v", "p_value"),  
  title = "Categorical Association Matrix",  
  fill_colors = c("#FFFFFF", "#9D2235"),  
  tl.cex = 0.8,  
  number.cex = 0.75,  
  show_coef = TRUE,  
  ...  
)
```

Arguments

data	A data frame or tibble.
cols	Optional character vector of column names to include. If NULL, all factor and character columns are selected.
method	Association metric: "cramer_v" (Cramer's V correlation, default) or "p_value" ($-\log_{10}(p)$ from categorical tests).
title	Plot title (defaults to "Categorical Association Matrix").
fill_colors	Character vector for the color gradient (defaults to Temple palette).
tl.cex	Label character expansion (default 0.8).
number.cex	Numeric value character expansion inside circles (default 0.75).
show_coef	Logical (default TRUE); whether to print coefficients.
...	Additional arguments passed to corrplot .

Value

Invisibly, the association matrix.

Examples

```
df <- data.frame(  
  A = factor(sample(c("Yes", "No"), 50, replace = TRUE)),  
  B = factor(sample(c("High", "Low"), 50, replace = TRUE)),  
  C = factor(sample(c("Group1", "Group2", "Group3"), 50, replace = TRUE))  
)  
plot_categorical_associations(df)
```

plot_cox_forest	<i>Forest Plot of Hazard Ratios with Confidence Intervals</i>
-----------------	---

Description

Generates a publication/deck-ready horizontal forest plot of hazard ratios.

Usage

```
plot_cox_forest(
  data,
  label_col = "index_label",
  hr_col = "estimate",
  low_col = "conf.low",
  high_col = "conf.high",
  p_col = NULL,
  color = "#9D2235",
  scale = c("hr", "log_hr"),
  color_by = c("none", "significance"),
  order_by = c("none", "magnitude", "pvalue"),
  x_limits = NULL,
  x_breaks = NULL,
  title = NULL,
  caption = NULL,
  base_size = 13
)
```

Arguments

<code>data</code>	Data frame containing tidy Cox regression results (e.g. from <code>broom::tidy</code> or multi-model runs).
<code>label_col</code>	Column name for labels / terms (character string or unquoted).
<code>hr_col</code>	Column name for hazard ratios, always on the HR (not log) scale (default: "estimate").
<code>low_col</code>	Column name for lower CI bounds, on the HR scale (default: "conf.low").
<code>high_col</code>	Column name for upper CI bounds, on the HR scale (default: "conf.high").
<code>p_col</code>	Optional column name for p-values. Required for <code>color_by = "significance"</code> and <code>order_by = "pvalue"</code> .
<code>color</code>	Primary color for points and error bars (default: Temple Cherry "#9D2235").
<code>scale</code>	One of "hr" (default; plots hazard ratios with a reference line at 1) or "log_hr" (log-transforms <code>hr_col/low_col/high_col</code> for plotting, with a reference line at 0).
<code>color_by</code>	One of "none" (default) or "significance", which colors points/error bars by whether <code>p_col</code> is below 0.05 and adds a legend.
<code>order_by</code>	One of "none" (default; original row order), "magnitude" (sorts by <code>abs(log(estimate))</code> descending), or "pvalue" (sorts by <code>p_col</code> ascending).
<code>x_limits</code>	Optional 2-element numeric vector for x-axis limits.
<code>x_breaks</code>	Optional numeric vector of axis break points.
<code>title</code>	Optional plot title.
<code>caption</code>	Optional plot caption.
<code>base_size</code>	Base font size (default: 13).

Value

A `ggplot2` object.

See Also

[plot_cox_forest_multi()]

plot_cox_forest_multi *Forest Plot of Hazard Ratios for a Multivariable Cox Model*

Description

Generates a forest plot from a [cbe_cox_multi()] result, grouping terms into per-variable blocks (facets) so factor levels stay visually attached to their parent variable.

Usage

```
plot_cox_forest_multi(
  x,
  color = "#9D2235",
  scale = c("hr", "log_hr"),
  color_by = c("none", "significance"),
  order_by = c("none", "magnitude", "pvalue"),
  x_limits = NULL,
  x_breaks = NULL,
  title = NULL,
  caption = NULL,
  base_size = 13
)
```

Arguments

x	A cbe_cox_multi object from [cbe_cox_multi()].
color	Primary color for points and error bars (default: Temple Cherry "#9D2235").
scale	One of "hr" (default) or "log_hr". See [plot_cox_forest()].
color_by	One of "none" (default) or "significance". See [plot_cox_forest()].
order_by	One of "none" (default), "magnitude", or "pvalue", applied within each variable's block. See [plot_cox_forest()].
x_limits	Optional 2-element numeric vector for x-axis limits.
x_breaks	Optional numeric vector of axis break points.
title	Optional plot title.
caption	Optional plot caption.
base_size	Base font size (default: 13).

Value

A ggplot2 object.

See Also

[cbe_cox_multi()], [plot_cox_forest()]

plot_cox_marginal	<i>Marginal Event Probability / Relative Hazard Diagnostic Plot</i>
-------------------	---

Description

Plots a continuous predictor against a model-derived curve, with a 95 for scale = "prob", overlaid actual binary event observations.

Usage

```
plot_cox_marginal(
  fit,
  data,
  feature = NULL,
  status_col = "status",
  scale = c("prob", "hr", "log_hr"),
  color = "#9D2235",
  base_size = 12
)
```

Arguments

fit	A fitted <code>survival::coxph</code> model or a <code>cbe_cox</code> object.
data	Data frame containing survival inputs.
feature	Character name of continuous predictor column.
status_col	Character name of status/event column (0 = censored, 1 = event; default: "status"). Only used when scale = "prob".
scale	One of "prob" (default; predicted probability of the event), "hr" (relative hazard, i.e. <code>exp()</code> of the centered linear predictor), or "log_hr" (the centered linear predictor itself).
color	Primary color for the fitted curve and CI band (default: Temple Cherry "#9D2235").
base_size	Base font size (default: 12).

Value

A `ggplot2` object.

plot_cox_survival	<i>Model-Predicted Survival Curves Stratified by Predictor</i>
-------------------	--

Description

Generates stratified predicted survival curves from a Cox model. For continuous predictors, values are automatically binned into quantile strata (e.g. quartiles).

Usage

```
plot_cox_survival(
  fit,
  data,
  feature = NULL,
  id_col = NULL,
  n_tiles = 4,
  label_endpoints = TRUE,
  overlay_km = FALSE,
  base_size = 12
)
```

Arguments

fit	A fitted <code>survival::coxph</code> model or a <code>cbe_cox</code> object.
data	Data frame used for fitting the model.
feature	Character name of the predictor column.
id_col	Character name of subject ID column (default: "arl_number" or row number).
n_tiles	Number of quantile bins for continuous predictors (default: 4).
label_endpoints	Logical; if TRUE, repels labels for stratum values at the final time point.
overlay_km	Logical (default: FALSE). If TRUE, overlays observed Kaplan-Meier step curves (dashed) for each stratum on top of the Cox-predicted curves (solid), sharing color by stratum, with a linetype legend distinguishing "Cox-predicted" from "KM observed". The event time/status used are those the model was fit with (<code>fit\$y</code>).
base_size	Base font size (default: 12).

Value

A `ggplot2` object.

plot_dynamic_trajectory

Longitudinal Biomarker Trajectories by Cohort / Outcome

Description

Plots mean trajectory over time with standard error bars, stratified by outcome or treatment cohort (e.g. Survived vs. Died).

Usage

```
plot_dynamic_trajectory(
  data,
  value_col,
  time_col,
  group_col,
  facet_col = NULL,
```

```

    colors = c(Survived = "#6F6F6F", Died = "#9D2235"),
    base_size = 12
  )

```

Arguments

data	Long-format data frame containing longitudinal measurements.
value_col	Character name of measurement column.
time_col	Character name of time point column.
group_col	Character name of cohort/grouping column.
facet_col	Optional character name of domain or biomarker column to facet wrap.
colors	Named character vector of colors matching group levels. Defaults to c(Survived = "#6F6F6F", Died = "#9D2235").
base_size	Base font size (default: 12).

Value

A ggplot2 object.

plot_group_comparison *Group Comparison Bar Chart with Standard Errors*

Description

Faceted bar chart comparing biomarker means and standard errors across groups.

Usage

```

plot_group_comparison(
  data,
  value_col,
  group_col,
  facet_col = NULL,
  se_col = NULL,
  colors = c(Survived = "#6F6F6F", Died = "#9D2235"),
  base_size = 13
)

```

Arguments

data	Data frame containing summary means and standard errors (or raw data).
value_col	Column containing index values or means.
group_col	Cohort/outcome grouping column (e.g. "outcome_label").
facet_col	Column to facet wrap across (e.g. "index_label").
se_col	Optional standard error column. If NULL, calculated automatically.
colors	Named color vector (default: c(Survived = "#6F6F6F", Died = "#9D2235")).
base_size	Base font size (default: 13).

Value

A ggplot2 object.

plot_missingness	<i>Missing Data Audit Plot</i>
------------------	--------------------------------

Description

Generates dodged percentage bar charts comparing data completeness across modalities.

Usage

```
plot_missingness(data, base_size = 13)
```

Arguments

- data Data frame with columns: parameter, pct (fraction missing), method.
- base_size Base font size (default: 13).

Value

A ggplot2 object.

plot_pca_bi	<i>PCA Biplot</i>
-------------	-------------------

Description

Biplot of feature loading vectors against two principal components.

Usage

```
plot_pca_bi(pca_model, newdata, column, x = 1, y = 2)
```

Arguments

- pca_model A [prcomp](#) object.
- newdata Data to project onto pca_model.
- column Column in newdata to use as point labels.
- x, y Which principal components to plot on the x/y axes (default 1, 2).

Value

A ggplot object.

Examples

```
pca_model <- prcomp(mtcars, center = TRUE, scale. = TRUE)
mtcars2 <- tibble::rownames_to_column(mtcars, "model")
plot_pca_bi(pca_model, mtcars2, column = "model")
```

plot_survival_km	<i>Standardized Kaplan-Meier Survival Curve</i>
------------------	---

Description

Renders a clean Kaplan-Meier curve using the Temple Cherry palette, with censor marks, confidence ribbons, and percentage-formatted survival axis.

Usage

```
plot_survival_km(
  data,
  time_col,
  status_col,
  group_col = NULL,
  color = "#9D2235",
  title = "Kaplan-Meier Survival Estimate",
  caption = NULL,
  base_size = 13
)
```

Arguments

data	Data frame containing survival inputs.
time_col	Character string naming the follow-up time column.
status_col	Character string naming the event status column (1 = event, 0 = censored).
group_col	Optional character string naming a stratification grouping column.
color	Primary line/ribbon color (default: Temple Cherry "#9D2235").
title	Optional plot title.
caption	Optional plot caption.
base_size	Base font size (default: 13).

Value

A ggplot2 object.

predict.coxnet_model	<i>Predict From a 'coxnet' Model</i>
----------------------	--------------------------------------

Description

Predict From a 'coxnet' Model

Usage

```
## S3 method for class 'coxnet_model'
predict(
  object,
  new_data,
  type = c("linear_pred", "survival"),
  penalty = NULL,
  eval_time = NULL,
  increasing = TRUE,
  ...
)
```

Arguments

object	A [coxnet()] model.
new_data	A data frame (or matrix) of new predictors.
type	"linear_pred" for the linear predictor, or "survival" for survival probabilities at 'eval_time'.
penalty	The penalty to predict at; defaults to the model's 'penalty'.
eval_time	For 'type = "survival"', the times to predict survival at.
increasing	For 'type = "linear_pred"': if 'TRUE' (the default, as in tidymodels' censored package), the sign is flipped so larger values mean longer survival. Use 'FALSE' for glmnet's own sign, where larger values mean higher risk.
...	Not used.

Value

A tibble with one row per row of 'new_data': '.pred_linear_pred', or '.pred', a list-column of tibbles with 'eval_time' and '.pred_survival'.

Survival is $S(t | x) = \exp(-H_0(t)e^{x^\top \beta})$, with H_0 the Breslow estimate of the cumulative baseline hazard from the training data at 'penalty'. Each row's covariates are taken as constant from time 0; to predict along a subject's start/stop covariate path, and score it, use [cv_coxnet()].

predict.joint_model *Predict Method for Joint Models*

Description

Generates multi-paradigm predictions from a fitted [joint_model](#): dynamic survival probabilities from the Cox model, binary event probabilities from the status model (both raw and calibrated), and expected duration from the time model.

Usage

```
## S3 method for class 'joint_model'
predict(object, new_data = NULL, eval_time = NULL, ...)
```

Arguments

object	A joint_model object.
new_data	Optional new data frame to predict upon. If NULL, predicts on training data.
eval_time	Horizon times for survival probabilities. Defaults to the model's eval_time.
...	Additional arguments passed to underlying predict methods.

Value

A tibble with columns:

.pred_survival	Nested list of survival probability curves over eval_time.
.pred_status	Predicted event probability from the status model.
.pred_status_calibrated	Calibrated event probability (if calibration was enabled).
.pred_time	Predicted duration from the time model.
.pred_linear_pred	Linear predictor from the Cox model (higher = longer survival).
.pred_risk_score	Relative hazard risk score from the Cox model.

print.cbe_cox

Print Method for cbe_cox Object

Description

Print Method for cbe_cox Object

Usage

```
## S3 method for class 'cbe_cox'
print(x, ...)
```

Arguments

x	Object of class cbe_cox
...	Additional arguments (unused)

print.cbe_cox_check	<i>Print Method for cbe_cox_check Object</i>
---------------------	--

Description

Print Method for cbe_cox_check Object

Usage

```
## S3 method for class 'cbe_cox_check'
print(x, ...)
```

Arguments

x	Object of class cbe_cox_check
...	Additional arguments (unused)

print.cbe_cox_multi	<i>Print Method for cbe_cox_multi Object</i>
---------------------	--

Description

Print Method for cbe_cox_multi Object

Usage

```
## S3 method for class 'cbe_cox_multi'
print(x, ...)
```

Arguments

x	Object of class cbe_cox_multi
...	Additional arguments (unused)

print.cbe_km	<i>Print Method for cbe_km Object</i>
--------------	---------------------------------------

Description

Print Method for cbe_km Object

Usage

```
## S3 method for class 'cbe_km'
print(x, ...)
```

Arguments

x	Object of class cbe_km
...	Additional arguments (unused)

print.joint_model	<i>Print Method for Joint Models</i>
-------------------	--------------------------------------

Description

Print Method for Joint Models

Usage

```
## S3 method for class 'joint_model'  
print(x, ...)
```

Arguments

x	A joint_model object.
...	Additional arguments.

print_pipeline	<i>Print the dependency graph as a PNG image</i>
----------------	--

Description

Renders via igraph/ggraph with a layered (Sugiyama) DAG layout and canvas sizing adapted to the graph's own node density, rather than the previous Graphviz/DiagrammeR renderer, which laid large graphs out extremely wide and flat (e.g. the full compute graph rendered at 9019x1140px, with text too small to read and the arrows between boxes impossible to follow).

Usage

```
print_pipeline(  
  objects,  
  path,  
  drop_stages = character(0),  
  drop_edges_to = character(0),  
  title = NULL,  
  base_font_size = 9,  
  label_wrap_width = 18,  
  min_width_in = 8,  
  min_height_in = 3,  
  max_width_in = 40,  
  width_per_node_in = 2.1,  
  dpi = 200  
)
```

Arguments

objects	The full list of FilePath, FileUses, and FileOutputs objects, or an already-built igraph object (see <code>as_igraph()</code>).
path	The path where the PNG image will be saved.
drop_stages	Character vector of ‘stage’ values to omit entirely (node and its edges). Use "output_report" on a crowded overview graph to hide the "<X> PDF"/"<X> DOCX" leaf nodes, which add little beyond what their parent QMD node already conveys.
drop_edges_to	Character vector of node names whose <i>incoming</i> edges should be omitted, while keeping the node itself visible. Use this for a node like an introduction/summary report that every other report is wired to depend on purely to force it to render last, not because of a real analytical dependency – dropping just the edges avoids a dense fan-in that conveys no real relationship.
title	Optional title drawn on the plot itself. Leave NULL when the image is embedded with its own Quarto fig-cap, to avoid a duplicate- looking caption.
base_font_size	Base font size (pt) for node labels; the title uses ‘base_font_size + 4’.
label_wrap_width	Character width at which node labels are wrapped.
min_width_in, max_width_in	Minimum and maximum plot width, in inches. Width scales with ‘width_per_node_in’ between these bounds.
min_height_in	Minimum plot height, in inches; height also scales with the number of pipeline layers.
width_per_node_in	Width, in inches, added per node along the graph’s horizontal extent, before clamping to ‘min_width_in’/‘max_width_in’.
dpi	Resolution (dots per inch) of the saved PNG.

proc_pca

*Process and Plot Principal Component Analysis (PCA)***Description**

Extracts a per-component variance summary from a PCA fit. Accepts either an existing `prcomp` object or raw data, in which case `prcomp` is run first.

Usage

```
proc_pca(data, center = TRUE, scale = TRUE, ...)
```

Arguments

data	A ‘prcomp’ object, or a numeric matrix/data frame to fit PCA on.
center, scale	Passed to <code>prcomp</code> (as ‘center’ and ‘scale.’) when ‘data’ is raw data; ignored when ‘data’ is already a ‘prcomp’ object.
...	Further arguments passed to <code>prcomp</code> when ‘data’ is raw data.

Value

A tibble with one row per component: ‘component’, ‘eigenvalue’, ‘variance_pct’, ‘cum_variance_pct’.

Examples

```
proc_pca(prcomp(mtcars[, 1:4], scale. = TRUE))

# Or fit the PCA in one step
proc_pca(mtcars[, 1:4], scale = TRUE)
```

profvis_summary	<i>Summarize a ‘profvis’ Profile</i>
-----------------	--------------------------------------

Description

Tabulates the samples of a [profvis::profvis()] profile by function: memory, memory increments, call frequency, stack depth, and memory over time.

Usage

```
profvis_summary(prof)
```

Arguments

prof A ‘profvis’ object.

Details

Memory columns sum the sampled ‘memalloc’/‘meminc’ values, so they rank functions by how much memory was allocated while they were on the stack; they are not exact allocation totals.

Value

A named list of tibbles:

‘memory_by_function’ Summed ‘memalloc’, share, calls, and memory per call, by function.

‘memory_increment_by_function’ Summed ‘meminc’, share, calls, and increment per call, by function.

‘calls_by_function’ Samples per function and share.

‘deepest_calls’ Each function’s deepest sample, deepest first.

‘memory_over_time’ Summed ‘memalloc’ per sample time.

Examples

```
## Not run:
prof <- profvis::profvis(for (i in 1:50) sort(runif(1e5)))
profvis_summary(prof)$memory_by_function

## End(Not run)
```

range_norm	<i>Range Normalization</i>
------------	----------------------------

Description

Like `min_max_norm`, but treats every numeric column as coming from a single combined distribution — one global min/max is used to rescale all of them together, rather than normalizing each column independently. Unlike `min_max_norm`, non-numeric columns are kept (only the numeric ones are transformed) — matching the original implementation this was migrated from.

Usage

```
range_norm(obj)
```

Arguments

`obj` A numeric vector, matrix, data frame, or tibble.

Value

A normalized object of the same dimensions and class as `obj`.

Examples

```
range_norm(data.frame(x = c(1, 2, 3), y = c(-10, 0, 10)))
```

`read_database_from_excel`

Read an R Database from an Excel Workbook

Description

Reads all worksheets from an Excel workbook into a named list of tibbles, optionally extracting and attaching the METADATA sheet as an attribute.

Usage

```
read_database_from_excel(
  path,
  include_metadata = TRUE,
  metadata_sheet = "METADATA",
  ...
)

cbe_read_database(
  path,
  include_metadata = TRUE,
  metadata_sheet = "METADATA",
  ...
)
```

Arguments

`path` Path to the .xls/.xlsx file.

`include_metadata` Logical; if TRUE (default) and a sheet named `metadata_sheet` is found, it is separated and attached as attribute "metadata" on the returned database list.

`metadata_sheet` Name of the metadata sheet to recognize (default: "METADATA").

`...` Additional arguments passed to [read_excel](#).

Value

A named list of tibbles. If `include_metadata = TRUE` and the metadata sheet was found, it is accessible via `attr(result, "metadata")`.

Examples

```
## Not run:
db <- read_database_from_excel("clinical_database.xlsx")
meta <- attr(db, "metadata")

## End(Not run)
```

<code>read_data_manifest</code>	<i>Read Data Manifest File</i>
---------------------------------	--------------------------------

Description

Reads a CSV manifest file listing frozen dataset copies and their repo-relative sources. Expected columns: `file` (filename in destination folder) and `source` (repo-relative path).

Usage

```
read_data_manifest(dir = ".", manifest_file = "manifest.csv")
```

Arguments

`dir` Directory containing the manifest file (default: current directory).

`manifest_file` Name of the manifest CSV file (default: "manifest.csv").

Value

A data frame containing the manifest entries.

read_excel_multiple_headers

Read Excel Data With Multi-Row Column Headers

Description

Variant of [read_excel](#) for sheets where a column's name is split across multiple header rows — the rows are concatenated in order, joined with " | ".

Usage

```
read_excel_multiple_headers(path, n_header_rows, ...)
```

Arguments

path	Path to the .xls/.xlsx file.
n_header_rows	Number of header rows in the sheet.
...	Additional arguments passed to read_excel .

Value

A tibble.

Examples

```
## Not run:
read_excel_multiple_headers("workbook.xlsx", n_header_rows = 2)

## End(Not run)
```

read_mapped_section_data

Read One Section's Data Using a Column Mapping

Description

Generic schema-enforced reader: reads whichever file corresponds to an INDEX value in a mapping table, renames raw headers to standardized names, resolves duplicate columns per duplicate_action, and tolerates missing or unmapped columns.

Usage

```
read_mapped_section_data(
  mapping,
  index,
  file = NULL,
  data_dir = NULL,
  sheet = NULL,
  header_rows = 1
)
```

Arguments

mapping	Validated column mapping table.
index	Character string identifying the section/domain in mapping\$INDEX.
file	Optional explicit path to the data file.
data_dir	Directory to search for the section's file if file is NULL.
sheet	Sheet name or number for Excel files.
header_rows	Number of header rows to concatenate (default: 1).

Value

A tibble with standardized column names.

read_raw_table	<i>Read Raw Table with Multi-Row Header Support</i>
----------------	---

Description

Read Raw Table with Multi-Row Header Support

Usage

```
read_raw_table(file, sheet = NULL, header_rows = 1)
```

Arguments

file	Path to .csv or .xlsx file.
sheet	Sheet name or number (for Excel).
header_rows	Number of rows to combine into column headers (default: 1).

Value

A tibble with raw data.

read_search	<i>Search a Directory Tree for File-Read Calls</i>
-------------	--

Description

Convenience wrapper around [find_code](#) that searches for common data-import function calls (readRDS, read.csv, read_excel, [read_workbook](#), [read_excel_multiple_headers](#), etc.) across a directory tree.

Usage

```
read_search(path, include_comments = TRUE, workers = NULL)
```

Arguments

path	Character. Root directory to search.
include_comments	Logical. Passed through to find_code . Defaults to TRUE.
workers	Integer number of parallel workers to use when future/furrr are installed. Defaults to NULL, which caps at available cores minus one. Falls back to sequential search when future/furrr are not installed.

Value

A tibble of matches (see [find_code](#)) with an additional pattern column recording which read function matched.

Examples

```
## Not run:
read_search(here::here("analysis"))

## End(Not run)
```

read_workbook	<i>Read Every Sheet of an Excel Workbook</i>
---------------	--

Description

Variant of [read_excel](#) that reads every sheet in a workbook into a named list of tibbles, named after each sheet.

Usage

```
read_workbook(path, ...)
```

Arguments

path	Path to the .xls/.xlsx file.
...	Additional arguments passed to read_excel .

Value

A named list of tibbles, one per sheet.

Examples

```
## Not run:
read_workbook("workbook.xlsx")

## End(Not run)
```

render

*Render Quarto and R Markdown Documents to Multiple Formats With Timing***Description**

Renders one or more Quarto (.qmd) or R Markdown (.Rmd) documents to multiple output formats (by default both PDF and DOCX, or formats defined in YAML or compute graphs) sequentially or in parallel, recording execution duration for each document/format combination. Useful for workflows where reviewers need a Word (.docx) file for track changes/edits while publication and archival copies need PDF.

Usage

```
render(
  path,
  formats = getOption("pipeline.default_formats", c("pdf", "docx")),
  pattern = "\\.(qmd|Rmd|rmd)$",
  engine = c("auto", "quarto", "rmarkdown"),
  workers = NULL,
  ...
)

render_me(
  path,
  formats = getOption("pipeline.default_formats", c("pdf", "docx")),
  pattern = "\\.(qmd|Rmd|rmd)$",
  engine = c("auto", "quarto", "rmarkdown"),
  workers = NULL,
  ...
)
```

Arguments

- | | |
|---------|---|
| path | Target document(s) to render. Can be: <ul style="list-style-type: none"> • A character string specifying a file path or directory. • A character vector of file paths. • A list of file paths (e.g. <code>list("report1.qmd", "report2.Rmd")</code>). • An S4 compute graph object (<code>FileOutputs</code> or <code>FilePath</code>). • A list of compute graph objects, such as a topological render plan returned by get_render_plan or a pipeline list from pipeline_config. • A named list or list of lists specifying per-document paths and formats (e.g. <code>list("report1.qmd" = c("pdf", "docx"), "report2.Rmd" = "html")</code>). |
| formats | Output format(s) to generate. Defaults to <code>getOption("pipeline.default_formats", c("pdf", "docx"))</code> . Can be: <ul style="list-style-type: none"> • A character vector of formats (e.g. <code>c("pdf", "docx"), "html", "gfm"</code>). • <code>"yaml"</code> or <code>"auto"</code> (or <code>NULL</code>) to read and render the output formats declared in each document's own YAML frontmatter. |

- Standard shorthand formats ("pdf", "docx", "html", "gfm") are automatically translated to R Markdown equivalents ("pdf_document", "word_document", "html_document", "github_document") when rendering with **rmarkdown**.

When a document in a compute graph defines its own deliverable output formats (via `create_qmd_renderer`), those document-level formats override the default formats.

pattern	Optional regular expression to filter files when path is a directory. Defaults to <code>"\\.(qmd Rmd rmd)\$"</code> .
engine	Rendering engine: "auto" (default), "quarto", or "rmarkdown". When "auto", .qmd files are rendered via quarto , and .Rmd files are rendered via quarto if available or rmarkdown otherwise.
workers	Integer specifying the number of multisession parallel workers when rendering multiple documents. Defaults to NULL, which caps at available cores minus one. Ignored (rendering falls back to sequential) when the future/furrr packages are not installed.
...	Additional options passed directly to the rendering backend (<code>quarto::quarto_render</code> or <code>rmarkdown::render</code>), such as <code>params</code> , <code>execute_params</code> , <code>output_dir</code> , <code>clean</code> , or <code>quiet</code> .

Value

A data frame containing:

file	Base filename of the rendered document
output	Output format (e.g., "pdf", "docx")
duration	Execution time in seconds
status	"success" or error condition message

Examples

```
## Not run:
# Render a single Quarto report to both PDF and DOCX
render(here::here("analysis", "report.qmd"))

# Render a single R Markdown report to both PDF and Word
render(here::here("analysis", "report.Rmd"))

# Render using formats declared in document's own YAML
render("analysis/report.qmd", formats = "yaml")

# Render a compute graph render plan
# plan <- get_render_plan(pipeline)
# render(plan)

# Render all QMD and Rmd reports in a directory in parallel
render("analysis/")

## End(Not run)
```

rossi_data	<i>Rossi Recidivism Dataset</i>
------------	---------------------------------

Description

Loads the classic Rossi recidivism study dataset used in biostatistical benchmarks for survival models with time-varying coefficients and covariates (e.g., JSS Vol 61, Code 01). The study tracks 432 convicts released from Maryland state prisons over a 52-week follow-up period.

Usage

```
rossi_data(format = c("csv", "sas"))
```

Arguments

format	Either "csv" (default, reading 'inst/extdata/rossi.csv') or "sas" (reading the native 'inst/extdata/recid.sas7bdat' via [haven::read_sas()]).
--------	---

Value

A tibble with 432 rows containing recidivism follow-up data:

week Follow-up time until arrest or censoring (1–52 weeks).
arrest Event indicator (1 = arrested, 0 = censored/did not re-offend).
fin Financial aid treatment (1 = received financial aid, 0 = control).
age Age at release from prison in years.
race Race indicator (1 = Black, 0 = other).
wexp Prior full-time work experience (1 = yes, 0 = no).
mar Marital status (1 = married, 0 = unmarried).
paro Release on parole (1 = yes, 0 = no).
prio Number of prior convictions.
educ Education level coded as 2 (grades 2-5) through 6 (some college).
emp1–emp52 Weekly employment status indicator across the 52 weeks.

Source

Rossi, P. H., Berk, R. A., & Lenihan, K. J. (1980). *Money, work, and crime: Some experimental results*. Academic Press.

See Also

[cbe_sas_macro_path()], [run_sas_script()]

Examples

```
df <- rossi_data()
head(df[, 1:10])
```

rotation_matrix	<i>PCA Rotation Matrix (Loadings)</i>
-----------------	---------------------------------------

Description

PCA Rotation Matrix (Loadings)

Usage

```
rotation_matrix(PC_mod)

pca_loadings(PC_mod)
```

Arguments

PC_mod A [prcomp](#) object.

Value

A tibble of feature loadings onto each principal component, with a feature_num column ("f1", "f2", ...) for compact labeling in [pca_eqns](#).

Examples

```
pca_model <- prcomp(mtcars, center = TRUE, scale. = TRUE)
rotation_matrix(pca_model)
```

run_sas_script	<i>Run a SAS Program in Batch Mode</i>
----------------	--

Description

Runs a ‘.sas’ file with the SAS executable, writing its log and listing to separate folders next to the program (created if needed).

Usage

```
run_sas_script(
  path,
  sas_path = find_sas(),
  log_dir = file.path(dirname(path), "logs"),
  list_dir = file.path(dirname(path), "list")
)
```

Arguments

path Path to the ‘.sas’ program.
sas_path Path to the SAS executable. Defaults to [find_sas()].
log_dir Folder for the ‘.log’ file. Defaults to ‘logs/’ beside ‘path’.
list_dir Folder for the ‘.lst’ listing. Defaults to ‘list/’ beside ‘path’.

Value

The exit status of the SAS process (0 on success, 1 for warnings, 2 for errors).

See Also

[find_sas()], [cbe_sas_macro_path()]

Examples

```
## Not run:
programs <- list.files("validation", pattern = "\\\\.sas$", full.names = TRUE)
vapply(programs, run_sas_script, integer(1))

## End(Not run)
```

R_names

Clean Column Names, Preserving Originals as Labels

Description

Runs [clean_names](#) on df and stores each column's original name as a **labelled** variable label, so the human-readable original isn't lost when the column name itself becomes machine-friendly.

Usage

```
R_names(df)
```

Arguments

df A data frame or tibble.

Value

A tibble with cleaned column names and the original names stored as variable labels (see [var_label](#)).

Examples

```
df <- tibble::tibble(
  `name with spaces` = 1:10,
  `special * characters` = LETTERS[1:10]
)
R_names(df)
```

scale_color_cbe	<i>Discrete Color Scale for Temple CBE</i>
-----------------	--

Description

Discrete Color Scale for Temple CBE

Usage

```
scale_color_cbe(...)
```

Arguments

... Arguments passed to `ggplot2::scale_color_manual`

scale_colour_temple	<i>Temple Color and Fill Scales for ggplot2</i>
---------------------	---

Description

Temple Color and Fill Scales for ggplot2

Usage

```
scale_colour_temple(
  palette = c("main", "diverging", "sequential"),
  discrete = NULL,
  reverse = FALSE,
  midpoint = 0,
  ...
)

scale_color_temple(
  palette = c("main", "diverging", "sequential"),
  discrete = NULL,
  reverse = FALSE,
  midpoint = 0,
  ...
)

scale_fill_temple(
  palette = c("main", "diverging", "sequential"),
  discrete = NULL,
  reverse = FALSE,
  midpoint = 0,
  ...
)
```

Arguments

palette	One of "main", "diverging", or "sequential".
discrete	Logical; a discrete scale (default for "main") or a continuous gradient.
reverse	Logical; reverse the color order.
midpoint	For continuous "diverging" scales, the data value mapped to white (default 0).
...	Passed to discrete_scale , scale_colour_gradient2 , or scale_colour_gradientn .

Value

A ggplot2 scale.

Examples

```
library(ggplot2)
ggplot(mtcars, aes(wt, mpg, colour = factor(cyl))) +
  geom_point() +
  scale_colour_temple()

ggplot(faithfuld, aes(waiting, eruptions, fill = density)) +
  geom_raster() +
  scale_fill_temple("sequential", discrete = FALSE)
```

scale_fill_cbe	<i>Discrete Fill Scale for Temple CBE</i>
----------------	---

Description

Discrete Fill Scale for Temple CBE

Usage

```
scale_fill_cbe(...)
```

Arguments

...	Arguments passed to <code>ggplot2::scale_fill_manual</code>
-----	---

scan_data_io	<i>Audit Data File Read/Write Calls Against a Project's Files on Disk</i>
--------------	---

Description

Scans `code_path` for read and write calls (via [read_search/write_search](#)), resolves the file paths those calls reference — including `here::here()` calls, and a heuristic fallback for indirect references like `excel_file_paths[["ABG"]]` matched against filenames actually present on disk — and cross-references them against every matching file under `project_root`. The result classifies each such file as a write output the code produces, a read input the code consumes, both, or an orphan with no code reference found ("unknown") — useful for finding stale, missing, or undocumented data deliverables in a project.

Usage

```
scan_data_io(
  code_path,
  project_root = code_path,
  ext = "xlsx",
  strong_read_patterns = c("read_workbook", "read_excel_multiple_headers"),
  include_comments_write = FALSE,
  max_depth = Inf
)
```

Arguments

<code>code_path</code>	Directory of scripts (.R/.Rmd/.qmd) to search for read/write calls.
<code>project_root</code>	Root directory to inventory files under, and to resolve relative/here::here() paths against (useful when auditing a checkout at a different location than the one the scripts were written against). Defaults to <code>code_path</code> .
<code>ext</code>	File extension to audit, without the leading dot (e.g. "xlsx", "csv", "rds"). Defaults to "xlsx".
<code>strong_read_patterns</code>	Character vector of read_search pattern names that should always count as a read of a file with this extension, even on a line that doesn't otherwise mention <code>ext</code> (e.g. a call to a project-specific reader function). Defaults to <code>c("read_workbook", "read_excel_multiple_headers")</code> .
<code>include_comments_write</code>	Logical, passed through to write_search . Defaults to FALSE.
<code>max_depth</code>	Maximum folder depth below <code>project_root</code> to inventory files in (default Inf, no limit).

Value

A list with:

<code>writes</code>	Resolved file-write call sites.
<code>inputs</code>	Resolved file-read call sites.
<code>missing_write_dirs</code>	Directories containing a write target that doesn't exist on disk yet, and any matching-extension files present there now — a likely spot for a not-yet-(re)generated deliverable.
<code>files</code>	Every file with extension <code>ext</code> under <code>project_root</code> , tagged <code>write_output/workflow_input/</code> unknown and by which script(s) reference it.

Examples

```
## Not run:
scan_data_io(here::here("analysis"), ext = "xlsx")

## End(Not run)
```

sd.error	<i>Compute standard error</i>
----------	-------------------------------

Description

Compute standard error

Usage

```
sd.error(x, na.rm = FALSE)
```

Arguments

x	vector of data
na.rm	should NA values be removed, TRUE or FALSE, default is FALSE

Value

numeric value representing standard error

Examples

```
sd.error(mtcars$mpg)
```

significance_stars	<i>P-value Significance Stars</i>
--------------------	-----------------------------------

Description

P-value Significance Stars

Usage

```
significance_stars(p_value)
```

Arguments

p_value	A numeric vector of p-values.
---------	-------------------------------

Value

A character vector of significance stars.

Examples

```
significance_stars(c(0.0001, 0.02, 0.2, 0.8))
```

simulate_section_data	<i>Simulate Synthetic Section Data From a Schema Mapping</i>
-----------------------	--

Description

Generates synthetic data honoring the column roles declared in a validated mapping table. If the `pslongSim` package is available, it is leveraged for baseline subject scaffolding; otherwise, an internal robust generator produces demographic and longitudinal grids.

Usage

```
simulate_section_data(
  mapping,
  index,
  n_subjects = 20,
  seed = 1,
  missing_rate = 0,
  time_levels = NULL
)
```

Arguments

<code>mapping</code>	Validated column mapping table; see validate_column_mapping .
<code>index</code>	Character string identifying the section/domain in <code>mapping\$INDEX</code> to simulate.
<code>n_subjects</code>	Number of synthetic subjects to generate (default: 20).
<code>seed</code>	Random seed for reproducibility (default: 1).
<code>missing_rate</code>	Probability in $[0, 1]$ that any non-ID, non-time cell is set to NA.
<code>time_levels</code>	Optional character vector of levels for the section's <code>Time_var</code> column (default: <code>c("T0", "T1", "T2", "T3")</code>).

Value

A tibble shaped according to the standardized new columns for the specified section.

single_t_test	<i>Single T-Test, Tidied</i>
---------------	------------------------------

Description

Runs a t-test comparing `.var` between the two levels of `.class`, tidied into one row with group sizes/SDs and fold-change added.

Usage

```
single_t_test(
  .data,
  .var,
  .class,
  .id = NULL,
  alternative = "two.sided",
  conf.level = 0.95,
  paired = FALSE,
  ...
)
```

Arguments

<code>.data</code>	A data frame or tibble.
<code>.var</code>	Name (string) of the continuous column to test.
<code>.class</code>	Name (string) of a binary (2-level) classification column.
<code>.id</code>	Name (string, optional) of a subject/record identifier column. When <code>paired = TRUE</code> , pass this to pair observations by matching <code>.id</code> across the two groups rather than by row position — every id must have exactly one observation in each group. If <code>paired = TRUE</code> and <code>.id</code> is omitted, observations are paired by row order within each group, which silently produces meaningless results if the two groups aren't already sorted into corresponding order.
<code>alternative</code>	One of "two.sided" (default), "greater", or "less".
<code>conf.level</code>	Confidence level for the interval.
<code>paired</code>	Logical; paired t-test.
<code>...</code>	Additional arguments passed to t.test .

Value

A one-row tibble.

Examples

```
mtcars |>
  dplyr::mutate(am = factor(am)) |>
  single_t_test("mpg", "am")
```

step_famd

Factor Analysis of Mixed Data (FAMD) Recipe Step

Description

`'step_famd()'` creates a *specification* of a recipe step that converts a mix of numeric and categorical variables into principal components with Factor Analysis of Mixed Data (`[FactoMineR::FAMD()]`). It is the mixed-data counterpart of `[recipes::step_pca()]`: numeric variables are standardized and categorical variables are weighted, so each variable contributes on a comparable scale.

Usage

```
step_famd(
  recipe,
  ...,
  role = "predictor",
  trained = FALSE,
  num_comp = 2,
  threshold = NA,
  options = list(),
  res = NULL,
  columns = NULL,
  levels = NULL,
  prefix = "FAMD",
  keep_original_cols = FALSE,
  skip = FALSE,
  id = recipes::rand_id("famd")
)

## S3 method for class 'step_famd'
tidy(x, type = "coef", ...)
```

Arguments

recipe	A recipe object. The step will be added to the sequence of operations for this recipe.
...	One or more selector functions to choose variables for this step. See <code>[recipes::selections()]</code> .
role	Role of the new component columns (default <code>"predictor"</code>).
trained	A logical indicating whether the step has been trained.
num_comp	Number of components to keep (default 2); ignored when <code>'threshold'</code> is set. <code>'0'</code> leaves the data unchanged.
threshold	Fraction of total variance, in (0, 1], that the kept components should cover. Overrides <code>'num_comp'</code> when not <code>'NA'</code> .
options	A named list of further arguments to <code>[FactoMineR::FAMD()]</code> .
res	The fitted FAMD object, once trained.
columns	The selected column names, once trained.
levels	The factor levels of the categorical columns, once trained.
prefix	Prefix of the new column names (default <code>"FAMD"</code>), numbered by <code>[recipes::names0()]</code> . It differs from <code>[recipes::step_pca()]</code> 's <code>"PC"</code> , so both steps can be used in one recipe.
keep_original_cols	A logical to keep the original variables in the output. Defaults to <code>'FALSE'</code> .
skip	A logical. Should the step be skipped when the recipe is baked by <code>[recipes::bake()]</code> ?
id	A character string that is unique to this step.
x	A <code>'step_famd'</code> object.
type	For <code>'tidy()'</code> : <code>"coef"</code> (variable contributions) or <code>"variance"</code> .

Details

The selected variables must include at least one numeric and one categorical (factor, character, or logical) variable; use `[recipes::step_pca()]` for all-numeric data. Character and logical variables are treated as factors, with levels learned by `'prep()'`. Missing values, and categories that `'bake()'` meets but training did not, are errors: impute or collapse levels first, e.g. with `[recipes::step_impute_mode()]` or `[recipes::step_other()]`.

FAMD has at most $\min(n - 1, p + \sum_j (L_j - 1))$ dimensions, for n rows, p numeric variables, and categorical variables with L_j levels, which can be more than the number of variables. `'num_comp'` is capped there. With `'threshold'`, the step keeps the fewest components whose cumulative share of variance reaches `'threshold'`, choosing among all dimensions.

Frequency weights (`[hardhat::frequency_weights()]`) are passed to FAMD as row weights. Like other unsupervised recipe steps, such as `[recipes::step_pca()]`, it ignores importance weights.

Tidying

`'tidy()'` with `'type = "coef"'` returns each variable's percentage contribution to each component (`'value'`): FactoMineR's per-variable summary, which is comparable between numeric and categorical variables. With `'type = "variance"'`, it returns each component's variance, cumulative variance, percent variance, and cumulative percent variance, as for `[recipes::step_pca()]`.

Value

An updated recipe object.

References

Pagès J (2004). Analyse factorielle de données mixtes. **Revue de Statistique Appliquée**, 52(4), 93-111.

Examples

```
if (requireNamespace("FactoMineR", quietly = TRUE)) {
  library(recipes)
  rec <- recipe(~ ., data = iris) |>
    step_famd(all_predictors(), num_comp = 3)
  prepped <- prep(rec)
  head(bake(prepped, new_data = NULL))
  tidy(prepped, number = 1, type = "variance")
}
```

step_lencode_coxnet	<i>Supervised Linear Encoding of Factors via Penalized Cox Models</i>
---------------------	---

Description

`step_lencode_coxnet` creates a *specification* of a recipe step that encodes categorical predictors into a single numeric variable representing the estimated log hazard ratio from a penalized Cox proportional hazards model on a survival outcome (`Surv(time, status)` or `Surv(tstart, tstop, status)`). Inspired by `embed::step_lencode_glm`.

Usage

```

step_lencode_coxnet(
  recipe,
  ...,
  outcome = c("time", "status"),
  role = "predictor",
  trained = FALSE,
  penalty = 0.05,
  mixture = 1,
  mapping = NULL,
  skip = FALSE,
  id = recipes::rand_id("lencode_coxnet")
)

tunable.step_lencode_coxnet(x, ...)

## S3 method for class 'step_lencode_coxnet'
required_pkgs(x, ...)

```

Arguments

<code>recipe</code>	A recipe object.
<code>...</code>	One or more selector functions to choose variables to be encoded. Must select factor or character columns.
<code>outcome</code>	A call to <code>recipes::vars()</code> selecting the survival outcome columns, e.g. <code>recipes::vars(time, status)</code> or <code>recipes::vars(tstart, tstop, status)</code> .
<code>role</code>	For model terms created by this step, what analysis role should they be assigned? By default, the function assumes that the new variables will be used as predictors in a model.
<code>trained</code>	A logical indicating whether the recipe has been trained.
<code>penalty</code>	A non-negative numeric value specifying the L1/L2 penalty for coxnet. Default is 0.05.
<code>mixture</code>	Elastic net mixing parameter between 0 and 1. Default is 1 (Lasso).
<code>mapping</code>	A named list of numeric vectors containing the level-to-score mappings, generated during <code>prep()</code> .
<code>skip</code>	A logical indicating whether the step should be skipped when the recipe is baked.
<code>id</code>	A unique identifier for the step.

Value

An updated version of recipe with the new step added.

Examples

```

## Not run:
if (requireNamespace("recipes", quietly = TRUE) && requireNamespace("survival", quietly = TRUE)) {
  lung <- survival::lung
  lung_df <- na.omit(lung[, c("time", "status", "sex", "ph.ecog")])
  lung_df$ph.ecog <- factor(lung_df$ph.ecog)
}

```

```

lung_df$sex <- factor(lung_df$sex)

rec <- recipes::recipe(time + status ~ ., data = lung_df) |>
  step_lencode_coxnet(ph.ecog, sex, outcome = c("time", "status"))
prepped <- recipes::prep(rec)
baked <- recipes::bake(prepped, new_data = NULL)
}

## End(Not run)

```

step_lencode_joint_model

Supervised Linear Encoding of Factors via Joint Survival-Status-Time Model

Description

step_lencode_joint_model creates a *specification* of a recipe step that encodes categorical predictors into a composite numeric risk score learned from a joint_model blending survival, binary status classification, and follow-up duration.

Usage

```

step_lencode_joint_model(
  recipe,
  ...,
  outcome = c("time", "status"),
  role = "predictor",
  trained = FALSE,
  engine = "glmnet",
  penalty = 0.05,
  mapping = NULL,
  skip = FALSE,
  id = recipes::rand_id("lencode_joint_model")
)

## S3 method for class 'step_lencode_joint_model'
required_pkgs(x, ...)

```

Arguments

recipe	A recipe object.
...	One or more selector functions to choose variables to be encoded.
outcome	A call to recipes::vars() selecting the survival outcome columns.
role	Role for the encoded variables. Default is "predictor".
trained	A logical indicating whether the step has been trained.
engine	Engine passed to joint_model(): "glmnet", "baguette", or "stacks".
penalty	Penalty parameter for the model. Default is 0.05.
mapping	A named list of level-to-score mappings generated during prep().
skip	Logical; skip step when baking? Default is FALSE.
id	A unique identifier for the step.

Value

An updated version of recipe.

stop_if_invalid_manifest

Stop Execution if Data Copies are Invalid or Stale

Description

Halts report execution with an actionable error message if any snapshot copy does not match its upstream source.

Usage

```
stop_if_invalid_manifest(validation, dir = ".")
```

Arguments

validation	Validation tibble returned by validate_data_manifest .
dir	Destination directory name for reporting.

Value

Invisibly returns validation if all files are valid.

summarize_section_by_time

Summarize One Section By Its Time Variable

Description

Generates summary tables (using **gtsummary** by default, or optionally **arsenal**) grouped by whichever column is flagged with `Time_var == TRUE` in the mapping dictionary.

Usage

```
summarize_section_by_time(
  df,
  mapping,
  index,
  id_cols = NULL,
  engine = c("gtsummary", "arsenal")
)
```

Arguments

df	Data frame already renamed via read_mapped_section_data .
mapping	Validated column mapping table.
index	Character string identifying the section/domain.
id_cols	Column names to exclude from summary (defaults to columns with ID_var == TRUE).
engine	Character string specifying the summary engine: "gtsummary" (the default) or "arsenal".

Value

A [tbl_summary](#) object (when engine = "gtsummary") or an `arsenal::tableby` summary object (when engine = "arsenal").

```
summarize_tune_results
```

Tune Over 'alpha' for Every Split of a Resample

Description

Calls `[tune_over_alpha()]` on each split of 'object' and row-binds the successful fits. Fits that errored are dropped. A split draws its own random 'alpha' values, reproducibly: the map runs with `'furrr_options(seed = TRUE)'`.

Usage

```
summarize_tune_results(
  object,
  ...,
  num_alpha_values = 10,
  num_fixed = 6,
  alphas = NULL,
  formulas = NULL,
  progress = FALSE
)
```

Arguments

object	A resampling object with a 'splits' column (e.g. one element of the 'inner_resamples' column of <code>[rsample::nested_cv()]</code>).
...	Arguments passed to <code>[glmnet_IBS()]</code> ('recipe', 'feature_names', 'time_data', 'id_col', 'cox.ties', ...).
num_alpha_values	Total number of 'alpha' values in the grid.
num_fixed	Number of evenly spaced values from 0 to 1 in the grid.
alphas	Optional numeric vector of 'alpha' values to use instead of the generated grid; in formula mode, one per formula.
formulas	Optional character vector of '+'-separated feature sets, or a list of one-sided formulas, each fit as a separate model (see <code>[glmnet_IBS()]</code> 's 'formula').
progress	Show a progress bar.

Value

A tibble of [glmnet_IBS()] results with an ‘inner_resamples_splits’ column giving the split’s position in ‘object\$splits’.

See Also

[glmnet_IBS()], [tune_over_alpha()], [nested_cv_coxnet()]

summary.tbl_graph	<i>S3 summary method for tbl_graph objects</i>
-------------------	--

Description

S3 summary method for tbl_graph objects

Usage

```
## S3 method for class 'tbl_graph'
summary(object, ...)
```

Arguments

object A tbl_graph object.
... Additional arguments (unused; for S3 method consistency).

Value

A tibble of the stage breakdown if ‘object’ is a pipeline graph (see [pipeline_summary](#)), otherwise the default summary() for the next applicable class.

SumNa	<i>Count Total Missing (NA) Values</i>
-------	--

Description

Calculates the total number of missing values across a vector, matrix, or data frame, optionally treating additional values (e.g. "NA" as a literal string, or "") as missing too.

Usage

```
SumNa(x, na_list = NULL)
```

Arguments

x A vector, matrix, or data frame.
na_list Optional vector of additional values to treat as missing, beyond actual NA.

Value

An integer representing the total count of missing values.

Examples

```
SumNa(c(1, 2, NA, 4, NA))
SumNa(data.frame(a = c(1, NA), b = c(NA, 2)))
SumNa(c(1, NA, "NA", 4), na_list = "NA")
```

surv_subject_truth	<i>Collapse Survival Outcomes to One Row per Subject</i>
--------------------	--

Description

Turns right-censored or counting-process (start/stop) survival outcomes into the one-row-per-subject, right-censored truth that yardstick's survival metrics require.

Usage

```
surv_subject_truth(truth, subject_id = NULL)
```

Arguments

truth	A [survival::Surv()] object: 'Surv(time, event)' or 'Surv(start, stop, event)'.
subject_id	Subject identifiers, one per element of 'truth'. Required for counting-process data; for right-censored data each row is its own subject when 'subject_id' is 'NULL'.

Details

For counting-process data, a subject's time is the stop time of their last interval and their event status is the status of that interval. Rows are checked first: within a subject, intervals must not overlap, and an event may only occur in the last interval (recurrent events are not supported).

Value

A tibble with one row per subject, sorted by 'subject_id': '.subject_id', '.entry' (start of the first interval; 0 for right-censored data), and '.truth', a right-censored 'Surv' object.

See Also

[censoring_km()], [graf_weights()], [add_graf_weights()]

Examples

```
if (requireNamespace("survival", quietly = TRUE)) {
  long <- data.frame(
    subject_id = c(1, 1, 2, 3, 3),
    tstart = c(0, 5, 0, 0, 4),
    tstop = c(5, 9, 6, 4, 10),
    status = c(0, 1, 0, 0, 0)
  )
  surv_subject_truth(
    survival::Surv(long$tstart, long$tstop, long$status),
    subject_id = long$subject_id
  )
}
```

table_two_by_two

*Formatted 2x2 Contingency Table with Exact Test***Description**

Builds an institutional 2x2 contingency table with row counts, row percentages, marginal totals, odds ratio estimates, confidence intervals, and exact test p-values. If any cell count is zero, it automatically defaults to the mid-p version of Central Fisher's exact test (midp = TRUE).

Usage

```
table_two_by_two(
  data,
  row_var,
  col_var,
  row_label = row_var,
  col_label = col_var,
  midp = NULL,
  conf.level = 0.95
)
```

Arguments

data	Data frame containing the categorical variables.
row_var	Character name of the row variable.
col_var	Character name of the column variable.
row_label	Optional display label for the row variable.
col_label	Optional display label for the column variable.
midp	Logical or NULL (default). If NULL, defaults to TRUE when any cell in the 2x2 table is zero, and FALSE otherwise.
conf.level	Confidence level for the odds ratio confidence interval (default 0.95).

Value

A list containing:

table	Tibble with formatted cells and marginal totals.
p_value	Exact test p-value.
estimate	Estimated odds ratio.
conf.int	Confidence interval for the odds ratio.
method	Test method name.
midp	Logical indicating whether mid-p adjustment was used.
has_zero	Logical indicating whether any table cell had zero count.
note	Formatted institutional table note.

temple_brand_path	<i>Path to the Bundled Temple brand.yml</i>
-------------------	---

Description

A copy of the quarto_temple_brand extension's brand.yml, for tools that read brand files directly, such as quarto::theme_brand_ggplot2() or bslib::bs_theme(brand =).

Usage

```
temple_brand_path()
```

Value

Path to brand.yml inside the installed package.

Examples

```
temple_brand_path()
```

temple_colors	<i>Temple University Brand Colors</i>
---------------	---------------------------------------

Description

Hex codes for the Temple University palette used by the `quarto_temple_brand` Quarto extension, so R graphics match branded reports. The palette follows Temple's current brand (<https://liberalarts.temple.edu/marcom/logos-and-brand>): primary cherry, white, and black; secondary Clear Skies and Book Nook; formal accents Academic Gold, Diamond Acres, Founder's Garden, and Night Owl; and casual accents Owl's Eye, Conwell Blue, Upward Momentum, and Cherry Blossom.

Usage

```
temple_colors(...)
```

Arguments

... Optional color names (e.g. "cherry", "night-owl"). With none, every color is returned.

Value

A named character vector of hex codes.

See Also

[temple_pal](#), [scale_colour_temple](#), [theme_temple](#)

Examples

```
temple_colors()
temple_colors("cherry", "night-owl")
```

temple_pal	<i>Temple Color Palettes</i>
------------	------------------------------

Description

Returns a palette function that generates n Temple brand colors.

Usage

```
temple_pal(palette = c("main", "diverging", "sequential"), reverse = FALSE)
```

Arguments

palette	One of "main", "diverging", or "sequential".
reverse	Logical; reverse the color order.

Details

- "main" Qualitative: cherry, Night Owl, Owl’s Eye, Founder’s Garden, Upward Momentum, Diamond Acres, black. At most 7 colors.
- "diverging" Night Owl through white to cherry, for values centered at zero (correlations, loadings, differences).
- "sequential" Book Nook to cherry.

Value

A function taking n and returning n hex codes.

Examples

```
temple_pal()(3)
temple_pal("diverging")(5)
```

theme_cbe	<i>Temple University CBE ggplot2 Themes, Palettes, and Formatters</i>
-----------	---

Description

Consistent institutional branding and clean publication/presentation styling for Temple University Center for Biostatistics & Epidemiology.
A minimal, publication-ready ggplot2 theme with clean gridlines and bold strip text.

Usage

```
theme_cbe(base_size = 11, base_family = "")
```

Arguments

base_size	Base font size (default: 11)
base_family	Base font family (default: "")

Value

A ggplot2 theme object

theme_cbe_deck	<i>CBE ggplot2 Theme for Presentation Decks</i>
----------------	---

Description

Optimized for slides (PowerPoint/Keynote) with larger typography and bottom legend.

Usage

```
theme_cbe_deck(base_size = 14)
```

Arguments

base_size Base font size (default: 14)

Value

A ggplot2 theme object

theme_temple	<i>Temple ggplot2 Theme</i>
--------------	-----------------------------

Description

A minimal theme with Temple cherry titles and facet strips, matching reports rendered with the quarto_temple_brand extension.

Usage

```
theme_temple(base_size = 11, base_family = "")
```

Arguments

base_size Base font size in points.

base_family Base font family. The brand's body typeface is "Faustina" (headings use "Roboto"); the default "" uses the device font, because naming a font that isn't installed makes devices warn.

Value

A ggplot2 theme.

Examples

```
library(ggplot2)
ggplot(mtcars, aes(wt, mpg)) +
  geom_point(colour = temple_colors("cherry")) +
  facet_wrap(~cyl) +
  labs(title = "Weight vs. mileage") +
  theme_temple()
```

tidy.cbe_loco_mp_coxnet

Tidy a LOCO-MP Cox Model Object

Description

Extracts feature importance and inference metrics into a clean tibble.

Usage

```
## S3 method for class 'cbe_loco_mp_coxnet'
tidy(x, ...)
```

Arguments

x	A cbe_loco_mp_coxnet object.
...	Not used.

Value

A tibble with term, importance, std_error, statistic, p_value, p_adjusted, conf_low, conf_high.

tidy.coxnet_model

Tidy the Coefficients of a 'coxnet' Model

Description

Tidy the Coefficients of a 'coxnet' Model

Usage

```
## S3 method for class 'coxnet_model'
tidy(x, penalty = NULL, ...)
```

Arguments

x	A [coxnet()] model.
penalty	The penalty to report coefficients at; defaults to the model's 'penalty'.
...	Not used.

Value

A tibble with 'term', 'estimate' (log hazard ratio per unit of the predictor; 0 for predictors the penalty removed), and 'penalty'.

tidy.joint_model	<i>Tidy Method for Joint Models</i>
------------------	-------------------------------------

Description

Extracts and aligns nonzero coefficients across the survival (coxnet), binary classification (status), and duration regression (time) models.

Usage

```
## S3 method for class 'joint_model'
tidy(x, ...)
```

Arguments

x	A joint_model object.
...	Additional arguments.

Value

A tibble with comparative coefficients per feature.

tidy_tmerge_cox	<i>Tidy Construction of Counting-Process (Start-Stop) Survival Data</i>
-----------------	---

Description

Merges repeated longitudinal measurement data with event/censoring data (and optional baseline covariates) to build tidy start-stop survival intervals (tstart, tstop, event, event_label). This provides a tidy, pipe-friendly alternative to SAS DATA-step macros and survival::tmerge for creating counting-process datasets for time-dependent Cox proportional hazards models.

Usage

```
tidy_tmerge_cox(
  measure_df,
  event_df,
  id = "subject_id",
  measure_time = "time",
  event_time = "event_time",
  event_type = "event_type",
  baseline_df = NULL,
  post_event = c("exclude", "include")
)
```

Arguments

measure_df	A data frame containing repeated longitudinal measurements.
event_df	A data frame containing subject event or censoring times and event types.
id	Column name identifying subjects across data frames (default: "subject_id").
measure_time	Column name in measure_df indicating measurement observation times (default: "time").
event_time	Column name in event_df indicating event or censoring time (default: "event_time").
event_type	Column name in event_df indicating event label or description (default: "event_type").
baseline_df	Optional data frame containing time-fixed baseline covariates keyed by id.
post_event	Character string: "exclude" (default) drops measurements occurring after the subject's event time; "include" extends event_time to the maximum observed measurement time.

Value

A tibble with columns tstart, tstop, event (1 if event occurred at tstop, 0 otherwise), event_label, and all merged measurement and baseline covariates.

See Also

[cbe_cox_multi()]

tumor_long	<i>Convert Wide Tumor Data to Counting Process (Start/Stop) Format</i>
------------	--

Description

Replicates the exact SAS DATA step transformation from SAS/STAT User's Guide (Example 85.7: "Tumor1" dataset). For each subject, intervals (T1, T2] are constructed across observation times where the time-dependent covariate NPap changes, with Status equal to Dead only on the terminal interval and 0 elsewhere.

Usage

tumor_long(data = NULL)

Arguments

data	A data frame structured like tumor_wide(). If NULL, tumor_wide() is used by default.
------	--

Value

A tibble with columns:

- ID** Subject ID.
- Time** Original death or censoring time.
- Dead** Original death status.
- Dose** Dose level.

T1 Start time of the risk interval.

T2 Stop time of the risk interval.

NPap Number of papillomas active during interval (T1, T2].

Status Event status at T2 (1 = event, 0 = censored/continuing).

Source

SAS/STAT User's Guide, Example 85.7: Time-Dependent Repeated Measurements of a Covariate.

https://support.sas.com/documentation/cdl/en/statug/68162/HTML/default/statug_phreg_examples07.htm

tumor_wide	<i>Internal SAS Benchmark Tumor Dataset (Wide Format)</i>
------------	---

Description

Transcribes the 45-animal tumor-promoting agent study from the SAS/STAT User's Guide (Example 85.7 / 91.7: "Time-Dependent Repeated Measurements of a Covariate"). In this study, 45 rodents were exposed to a carcinogen and randomized to three dose levels of a tumor-promoting agent (1.0, 2.5, 10.0). The number of papillomas was observed repeatedly across 15 observation times (weeks 27, 34, 37, 41, 43, 45, 46, 47, 49, 50, 51, 53, 65, 67, 71).

Usage

```
tumor_wide()
```

Value

A tibble with 45 rows and 19 columns:

ID Subject ID (1–45).

Time Event or censoring time in weeks.

Dead Censoring status (1 = dead, 0 = censored).

Dose Dose level of the promoting agent (1.0, 2.5, or 10.0).

P1–P15 Number of papillomas counted at observation weeks 27, 34, 37, 41, 43, 45, 46, 47, 49, 50, 51, 53, 65, 67, 71. Missing (NA) beyond death or study termination.

Source

SAS/STAT User's Guide, Example 85.7: Time-Dependent Repeated Measurements of a Covariate.

https://support.sas.com/documentation/cdl/en/statug/68162/HTML/default/statug_phreg_examples07.htm

tune_over_alpha

*Tune a Penalized Cox Model Over a Grid of ‘alpha’ Values***Description**

Runs [glmnet_IBS()] on one resampling split for each value in an ‘alpha’ grid, or for each of a set of candidate formulas, in parallel via [furrr::future_map()]. Set a [future::plan()] first to run in parallel; this function never sets one.

Usage

```
tune_over_alpha(
  object,
  ...,
  num_alpha_values = 10,
  num_fixed = 6,
  alphas = NULL,
  formulas = NULL,
  progress = FALSE
)
```

Arguments

object	An ‘rsplit’.
...	Arguments passed to [glmnet_IBS()] (‘recipe’, ‘feature_names’, ‘time_data’, ‘id_col’, ‘cox.ties’, ...).
num_alpha_values	Total number of ‘alpha’ values in the grid.
num_fixed	Number of evenly spaced values from 0 to 1 in the grid.
alphas	Optional numeric vector of ‘alpha’ values to use instead of the generated grid; in formula mode, one per formula.
formulas	Optional character vector of ‘+’-separated feature sets, or a list of one-sided formulas, each fit as a separate model (see [glmnet_IBS()]’s ‘formula’).
progress	Show a progress bar.

Details

****Grid mode**** (the default). The grid is ‘num_fixed’ evenly spaced values from 0 to 1, plus ‘num_alpha_values - num_fixed’ values drawn uniformly at random, one per gap between consecutive fixed values (cycling through the gaps if there are more random draws than gaps).

****Formula mode**** (‘formulas’ given). One fit per formula, each with its own ‘alpha’: the matching element of ‘alphas’, or a value drawn uniformly from 0 to 1 when ‘alphas’ is ‘NULL’. Each result gains a ‘formula’ column.

Random draws use the R session’s random number stream, so call [set.seed()] first for reproducible values. Model fits run with ‘furrr_options(seed = TRUE)’, so they are reproducible too.

To tune ‘alpha’ and the penalty together in one call, see [cv_coxnet()].

Value

A list named by ‘alpha’, one element per fit, each the output of `[purrr::safely()]`: a list with ‘result’ (the `[glmnet_IBS()]` tibble, or ‘NULL’) and ‘error’ (‘NULL’, or the condition).

See Also

`[glmnet_IBS()]`, `[summarize_tune_results()]`, `[cv_coxnet()]`

use_template_brand

Install the Temple Brand Quarto Extension Into a Project

Description

Installs **quarto_template_brand** with `quarto::quarto_add_extension()`, creating a minimal `_quarto.yml` first if there isn’t one (Quarto only applies a brand extension inside a project). Documents in path can then use format: `template-html`, `template-pdf` (LaTeX title page), `template-typst`, or `template-revealjs`, and every format picks up the brand colors, fonts, and logo.

Usage

```
use_template_brand(
  path = ".",
  extension = "jkylearmstrong-temple/quarto_template_brand",
  quiet = FALSE,
  check_root = TRUE
)
```

Arguments

path	Project directory (created if missing). Defaults to the working directory.
extension	Extension source passed to <code>quarto add</code> : a GitHub org/repo, a URL, or a local path.
quiet	Logical; suppress Quarto’s output.
check_root	Logical; warn when path isn’t the current project’s root as found by here (i.e. the nearest ancestor with a <code>.Rproj</code> , <code>.git</code> , or similar marker). Set to <code>FALSE</code> to install into a report’s own subfolder without the warning, e.g. for a one-off report that won’t grow siblings.

Details

Install once, at the project root. Quarto resolves a project’s `_extensions` from any subfolder of that project, so a single install at the root (where `_quarto.yml` lives) covers every report under it. Calling `use_template_brand()` again for each report’s own subfolder instead creates a separate `_extensions` copy per report, which can drift out of sync as the extension is updated. See [create_report](#)’s `install_brand` argument for the one-report shortcut, which installs into that report’s own folder and is fine when there’s only ever going to be one.

Value

Invisibly, a list with `path`, `quarto_yaml`, `created_quarto_yaml`, and `extension_dir`.

See Also

`create_report` with `template_name = "template"`.

Examples

```
## Not run:
# Install once at the project root; every report below it shares the
# extension automatically.
use_template_brand(here::here())
create_report(here::here("analysis"), template_name = "template")
create_report(here::here("reports"), template_name = "template", filename = "q3")

# A single, one-off report: install alongside it and skip the check.
use_template_brand("analysis", check_root = FALSE)
create_report("analysis", template_name = "template")

## End(Not run)
```

validate_column_mapping

Validate a Plug-and-Play Column Mapping

Description

Verifies that a mapping dictionary has the required schema and logical constraints: columns INDEX, old, new, X_var, Y_var, ID_var, Time_var, duplicate_of, and duplicate_action.

Usage

```
validate_column_mapping(mapping)
```

Arguments

mapping A data frame/tibble containing the column mapping rules.

Value

TRUE, invisibly, if valid; otherwise raises an error listing all issues.

validate_data_manifest

Validate Data Snapshot Copies Against Upstream Sources

Description

Compares destination files against upstream source files checking: 1. Existence of both source and copy. 2. Cryptographic MD5 hash equality. 3. Object-level equivalence via `identical(readRDS(), readRDS())` for `.rds` files.

Usage

```
validate_data_manifest(dir = ".", manifest = NULL, project_root = here::here())
```

Arguments

<code>dir</code>	Destination directory containing the snapshot copies.
<code>manifest</code>	Data frame with columns <code>file</code> and <code>source</code> .
<code>project_root</code>	Root directory for resolving source paths (default: <code>here::here()</code>).

Value

A tibble with validation status per file.

<code>visualize_pipeline</code>	<i>Visualize FilePath dependencies with stage-based coloring and staleness check</i>
---------------------------------	--

Description

Visualize FilePath dependencies with stage-based coloring and staleness check

Usage

```
visualize_pipeline(  
  all_objects,  
  extract_graph_code = FALSE,  
  stage_colors = NULL  
)
```

Arguments

<code>all_objects</code>	The full list of <code>FilePath</code> , <code>FileUses</code> , and <code>FileOutputs</code> objects.
<code>extract_graph_code</code>	Logical. If <code>TRUE</code> , the graph source code will be extracted and returned as a character string.
<code>stage_colors</code>	Optional named list of stage colors. If <code>NULL</code> , uses <code>pipeline_config()</code> options.

visualize_pipeline_interactive	<i>Generate an interactive visNetwork visualization of the computation graph</i>
--------------------------------	--

Description

Generate an interactive visNetwork visualization of the computation graph

Usage

```
visualize_pipeline_interactive(  
  all_objects,  
  direction = "LR",  
  hierarchical = TRUE,  
  title = getOption("pipeline.study_name", "Computational Pipeline"),  
  subtitle = "Computational Pipeline"  
)
```

Arguments

all_objects	A list of FilePath objects or an igraph object.
direction	Layout direction: "LR" (Left-to-Right, default) or "UD" (Top-to-Bottom).
hierarchical	Logical, whether to use a hierarchical DAG layout. Default is TRUE.
title	Character, main title – defaults to getOption("pipeline.study_name", "Computational Pipeline").
subtitle	Character, the document-specific name (e.g. "Computational Pipeline", or a subgraph's own name from export_subgraph()).

Value

A visNetwork htmlwidget.

volcano_plot	<i>Volcano Plot</i>
--------------	---------------------

Description

Volcano Plot

Usage

```
volcano_plot(  
  .data,  
  log2_fold_change,  
  log_p,  
  var,  
  alpha = 0.05,  
  highlight_significant = TRUE  
)
```

Arguments

.data	A data frame, e.g. the output of <code>multiple_t_test</code> .
log2_fold_change	Unquoted column of log2 fold-change.
log_p	Unquoted column of $-\log_{10}(\text{p.value})$.
var	Unquoted column of variable labels.
alpha	Significance threshold for the reference line/shading (default 0.05).
highlight_significant	Logical (default TRUE); shade the region above the $p < \alpha$ line.

Value

A ggplot object.

Examples

```
mtcars |>
  dplyr::mutate(am = factor(am)) |>
  multiple_t_test(.class = "am") |>
  volcano_plot(log2_fold_change = log2_fold_change, log_p = log_p, var = var)
```

words	<i>Combine Character Vector into Natural-Language Words</i>
-------	---

Description

Combine Character Vector into Natural-Language Words

Usage

```
words(x, none = "none")
```

Arguments

x	Character vector
none	Fallback text when vector is empty (default: "none")

Value

Character string (e.g., "A, B, and C")

write_database_to_excel

Write an R Database or Data Frames to an Excel Workbook

Description

Writes a named list of data frames (an R database) or a single data frame to a multi-worksheet Excel workbook, with an optional automated data dictionary worksheet.

Usage

```
write_database_to_excel(
  x,
  path,
  include_metadata = TRUE,
  metadata = NULL,
  metadata_sheet = "METADATA",
  ...
)
```

```
write_workbook(
  x,
  path,
  include_metadata = TRUE,
  metadata = NULL,
  metadata_sheet = "METADATA",
  ...
)
```

```
cbe_write_database(
  x,
  path,
  include_metadata = TRUE,
  metadata = NULL,
  metadata_sheet = "METADATA",
  ...
)
```

Arguments

x	A named list of data frames (an R database) or a single data frame/tibble.
path	File path to save the .xlsx workbook.
include_metadata	Logical; if TRUE (default), computes and appends a data dictionary sheet generated via get_dataset_info .
metadata	Optional pre-computed metadata data frame. If NULL and include_metadata = TRUE, metadata is computed automatically.
metadata_sheet	Character string naming the metadata worksheet (default: "METADATA").
...	Additional arguments passed to write_xlsx .

Value

The file path invisibly.

Examples

```
## Not run:
db <- list(patients = data.frame(id = 1:3), labs = data.frame(id = 1:3, val = 4:6))
write_database_to_excel(db, "clinical_database.xlsx")

## End(Not run)
```

write_search	<i>Search a Directory Tree for File-Write Calls</i>
--------------	---

Description

Convenience wrapper around [find_code](#) that searches for common data-export function calls (`write_xlsx`, `saveRDS`, `write.csv`, etc.) across a directory tree.

Usage

```
write_search(path, include_comments = FALSE, workers = NULL)
```

Arguments

path	Character. Root directory to search.
include_comments	Logical. Passed through to find_code . Defaults to FALSE.
workers	Integer number of parallel workers to use when future/furrr are installed. Defaults to NULL, which caps at available cores minus one. Falls back to sequential search when future/furrr are not installed.

Value

A tibble of matches (see [find_code](#)) with an additional pattern column recording which write function matched.

Examples

```
## Not run:
write_search(here::here("analysis"))

## End(Not run)
```

zip_render

*Render a Quarto Document and Zip It With Its Dependencies***Description**

Renders input in an isolated build directory, copies in any resources it references (explicitly, or heuristically detected from quoted paths / `here::here()` calls), and zips the outputs together with the source and sidecar files.

Usage

```
zip_render(
  input,
  formats = c("html", "pdf", "docx"),
  resources = NULL,
  detect = c("heuristic", "none"),
  build_dir = NULL,
  zip_name = NULL,
  copy_back_dir = NULL,
  include_sources = TRUE,
  overwrite = TRUE,
  verbose = TRUE
)
```

Arguments

input	Path to the input .qmd file.
formats	Character vector of output formats (e.g. <code>c("html", "pdf", "docx")</code> or <code>"all"</code>).
resources	Optional character vector of extra files to include, absolute or project-relative.
detect	<code>"heuristic"</code> (default; scans the .qmd for likely file paths) or <code>"none"</code> .
build_dir	Staging directory; defaults to a fresh temp directory.
zip_name	Name of the resulting zip; defaults to <code><input-stem>.zip</code> .
copy_back_dir	Where to copy the finished zip; defaults to <code>dirname(input)</code> .
include_sources	Logical (default <code>TRUE</code>); include the .qmd and sidecar bib/tex/css files.
overwrite	Logical (default <code>TRUE</code>); overwrite an existing zip at the destination.
verbose	Logical (default <code>TRUE</code>); print progress messages.

Details

Heuristic detection inspects quoted file paths and `here::here()` calls in the document. For YAML front-matter resources, pass them explicitly via the ‘resources’ argument or include them in the document’s YAML; this function will attempt to parse YAML when present to pick up top-level resource lists.

Value

Invisibly, a list with the build directory, detected/copied resources, render outputs, and final zip path.

Examples

```
## Not run:
zip_render("report.qmd", formats = c("html", "pdf"))

## End(Not run)
```

zip_reports

Package Multiple Already-Rendered Reports Into an Indexed Zip

Description

Given a data frame describing an ordered set of already-rendered Quarto reports grouped into stage folders, copies their existing PDF/DOCX/HTML outputs into a staged build directory, builds an Excel index with hyperlinks to each report, optionally bundles arbitrary data-deliverable files alongside them, and zips the result.

Usage

```
zip_reports(
  reports,
  output_formats = c("pdf", "docx"),
  data_deliverables = character(),
  zip_name = NULL,
  output_dir = getwd(),
  docx_from_pdf = NULL
)
```

Arguments

reports	A data frame/tibble with one row per report and columns name (display name), path (path to the .qmd source — the corresponding .pdf/.docx/.html output paths are derived by swapping its extension), and optionally stage (subfolder to group the report under; NA/empty becomes "99_Other") and description.
output_formats	Character vector of formats to look for and include: any of "pdf", "docx", "html", or "all" for all three. Defaults to c("pdf", "docx").
data_deliverables	Character vector of additional file paths to copy into a top-level data/ folder in the zip. Defaults to character().
zip_name	Name of the output zip file. Defaults to all_reports_<date>.zip.
output_dir	Directory the finished zip is written to. Defaults to getwd().
docx_from_pdf	Optional function function(src_pdf, dest_docx) used to create dest_docx from an existing src_pdf when a report requests "docx" output but no (or a stale) .docx file exists next to its .qmd source. If NULL (default), missing/stale DOCX files are silently skipped rather than generated.

Details

Unlike `zip_render` (which renders and zips a single document), `zip_reports` does not render anything itself — it packages outputs that have already been rendered, e.g. by `render` (or `render_me`). Report order and staging is entirely the caller’s decision (for example, the topological sort of a project’s own dependency graph) — pass reports pre-ordered. If multiple reports share the same source stem (for example `analysis1/analysis.qmd` and `analysis2/analysis.qmd`), staged output names are disambiguated to prevent overwrites within a stage/format folder.

Value

The path to the created zip file.

Examples

```
## Not run:
reports <- data.frame(
  name = c("Introduction", "Results"),
  path = c("analysis/intro.qmd", "analysis/results.qmd"),
  stage = c("00_intro", "01_results"),
  description = c("Project overview", "Primary analysis results")
)
zip_reports(reports, output_formats = c("pdf", "docx"))

## End(Not run)
```

<code>z_norm</code>	<i>Z-Score Standard Normalization</i>
---------------------	---------------------------------------

Description

Standardizes numeric features to have mean = 0 and standard deviation = 1.

Usage

```
z_norm(x, na.rm = TRUE)
```

Arguments

- `x` A numeric vector, matrix, or data frame.
- `na.rm` Logical; whether to ignore NA values (default TRUE).

Value

Z-score standardized numeric object.

Examples

```
z_norm(c(10, 20, 30, 40, 50))
```

Index

`%!in%(infix_helpers)`, 99
`%flike%(infix_helpers)`, 99
`%ilike%(infix_helpers)`, 99
`%like%(infix_helpers)`, 99
`%notin%(infix_helpers)`, 99
`%plike%(infix_helpers)`, 99

`add_graf_weights`, 7
`add_missing`, 8
`add_p`, 30, 46
`anonymize_pi`, 9
`apply_database_metadata`, 11
`as.package`, 16
`as_igraph`
 (`as_igraph.cbe_database_relationships`), 12
`as_igraph.cbe_database_relationships`, 12
`as_pipeline_graph`, 12, 13
`as_tbl_graph.cbe_database_relationships`
 (`cbe_database_relationships`), 26
`as_tbl_graph.list`, 13
`audit_report_deliverables`, 13
`autoplot.cbe_compare_df`, 14
`autoplot.cbe_database_relationships`, 14
`autoplot.cbe_key_integrity`, 15
`autoplot.cbe_loco_mp_coxnet`, 15
`autoplot.cv_coxnet` (`cv_coxnet-methods`), 69
`autoplot.cv_joint_model`, 16

`build_manual_versioned`, 16

`calculate_fences`, 17
`cbe_balloon_plot`, 17
`cbe_bar_plot`, 18
`cbe_check_key_integrity`, 15
`cbe_check_key_integrity`
 (`cbe_database_relationships`), 26
`cbe_compare_df`, 14, 19, 26
`cbe_contingency_plot`, 17, 18, 20, 35, 36, 41, 44
`cbe_cox_check`, 22
`cbe_cox_multi`, 23
`cbe_cox_single`, 24
`cbe_cox_table`, 25
`cbe_database_label` (`cbe_labels`), 38
`cbe_database_label<-` (`cbe_labels`), 38
`cbe_database_name` (`cbe_labels`), 38
`cbe_database_name<-` (`cbe_labels`), 38
`cbe_database_relationships`, 12, 14, 26
`cbe_database_venn`, 27
`cbe_dataset_label` (`cbe_labels`), 38
`cbe_dataset_label<-` (`cbe_labels`), 38
`cbe_docx_review_extract`, 28
`cbe_exact2x2`, 29, 31, 35, 44
`cbe_exact2x2_ci`, 31
`cbe_explain`, 32
`cbe_explain_status` (`cbe_explain`), 32
`cbe_explain_survival`, 33
`cbe_explain_time` (`cbe_explain`), 32
`cbe_factor_reference`, 34
`cbe_find_shared_keys`
 (`cbe_database_relationships`), 26
`cbe_four_quadrant_report`, 35
`cbe_get_dataset_labels` (`cbe_labels`), 38
`cbe_get_id_cols` (`cbe_variable_roles`), 48
`cbe_get_outcomes` (`cbe_variable_roles`), 48
`cbe_get_predictors`
 (`cbe_variable_roles`), 48
`cbe_get_roles` (`cbe_variable_roles`), 48
`cbe_get_time_cols` (`cbe_variable_roles`), 48
`cbe_heatmap_plot`, 36
`cbe_km_single`, 37
`cbe_labels`, 38
`cbe_loco_mp_coxnet`, 39
`cbe_mosaic_plot`, 41
`cbe_pairwise_combos`, 41
`cbe_palette`, 42
`cbe_pformat` (`pformat`), 135

- cbe_predict_parts_shap, 42
- cbe_read_database
 - (read_database_from_excel), 152
- cbe_sas_macro_dir, 43
- cbe_sas_macro_path, 43
- cbe_set_dataset_labels (cbe_labels), 38
- cbe_set_roles (cbe_variable_roles), 48
- cbe_square_plot, 44
- cbe_survex_loss_brier
 - (cbe_survex_loss_ibs), 45
- cbe_survex_loss_ibs, 45
- cbe_test_categorical, 46
- cbe_theme_survival, 48
- cbe_variable_roles, 48, 49
- cbe_write_database
 - (write_database_to_excel), 190
- censoring_km, 50
- check_docx_toolchain, 51
- chisq.test, 21, 35, 45, 47
- clean_names, 51, 161
- collapse_by_stage, 52
- collect_metrics.cv_coxnet
 - (cv_coxnet-methods), 69
- collect_metrics.nested_cv_coxnet
 - (nested_cv_coxnet), 121
- combn, 60
- compute_graph, 52
- controls_datasets, 53
- convert_pdf_to_docx, 54
- convert_pdfs_to_docx, 53
- copy_data_manifest, 55
- cor, 57, 84
- cor.test, 60
- corr_test_all, 60
- correlation_diff, 55, 56
- correlation_diff_heatmap, 56
- correlation_plot, 56, 57, 57–59, 137
- correlation_plot_split, 58, 58
- corrplot, 58, 138
- coxnet, 61, 103
- create_qmd_renderer, 62, 79, 81, 136, 158
- create_report, 63, 185, 186
- create_toc_from_sas_pdf, 64
- cutree, 59
- cv_coxnet, 65, 103, 104
- cv_coxnet-methods, 69
- cv_joint_model, 70, 104
- database_metadata_io, 71
- delete_nul_files, 71
- demo_cbe_mapping, 72
- detect_outliers, 72
- discrete_scale, 163
- distribution_plot, 73
- distribution_test, 73
- exact2x2, 30, 47
- export_interactive_pipeline, 52, 74
- export_subgraph, 52, 74
- extract_all_xlsx_tokens, 75
- extract_surv_components, 76
- extract_win_posix_paths, 76
- extract_yaml_formats, 77
- features_percent_miss, 9, 77, 137
- file_meta_fs, 82
- FileOutputs, 78, 78, 79, 81, 82
- FileOutputs-class, 79
- FilePath, 78, 79, 79–82
- FilePath-class, 80
- FileUses, 79, 81, 81, 82
- FileUses-class, 82
- find_code, 83, 155, 156, 191
- find_correlation, 84
- find_python, 85
- find_sas, 86
- find_section_file, 86
- find_soffice, 87
- findCorrelation, 84
- fisher.test, 21, 35, 45, 47
- fit_n_rank, 87
- flag_outliers, 72, 88
- flike (infix_helpers), 99
- fmt_hr, 89
- fmt_num, 89
- fmt_p, 90
- fmt_pct, 90
- fmt_sig, 91
- formula, 47
- generate_last_names, 91
- generate_pi_names, 92
- generate_pseudonym_token, 93
- get_dataset_info, 71, 94, 190
- get_model_parameters, 95
- get_render_plan, 96, 157
- ggplot, 14, 27
- ggVennDiagram, 27
- glmnet, 103
- glmnet_IBS, 96, 104
- graf_weights, 98
- grep, 83
- grepl, 100
- group_vfold_cv, 70
- here, 185

- `ilike (infix_helpers)`, 99
- `infix_helpers`, 99
- `install.packages`, 100
- `install.packages.no_lock`, 100
- `is.int`, 101
- `is_normal`, 73, 101, 102
- `is_poisson`, 73, 102
-
- `join_pipelines`, 104
- `joint_model`, 70, 94, 103, 124, 146
-
- `keep_only`, 105
- `km_summary_to_prism`, 106
-
- `label_pvalue`, 135
- `like (infix_helpers)`, 99
- `list_stage_data_artifacts`, 107
-
- `make_excel_names`, 107
- `manhattan_plot`, 108
- `min_max_norm`, 73, 109, 152
- `missforest_impute_by_mtry`, 109
- `missforest_oob_by_mtry`, 110
- `missforest_sweep_mtry`, 111
- `missmap`, 113
- `missranger_max_mtry`, 114
- `missranger_oob_by_mtry`, 115
- `missranger_sweep_mtry`, 117
- `multiple_t_test`, 108, 119, 189
- `my_summary_table`, 120
-
- `nested_cv`, 123
- `nested_cv_coxnet`, 121
- `nested_cv_joint_model`, 104, 123
- `normalize_safely`, 124
- `notin (infix_helpers)`, 99
-
- `one_vs_rest_t_test`, 125
-
- `package_deliverables`, 126
- `parse_here_call_vec`, 126
- `pca_biplot`, 127, 132
- `pca_eqns`, 128, 160
- `pca_feature_loading_heatmap`, 129, 130, 132
- `pca_loading_diff`, 56, 129, 130
- `pca_loading_diff_heatmap`, 56, 130
- `pca_loadings (rotation_matrix)`, 160
- `pca_percent_var_explained`, 131, 132
- `pca_plot`, 131
- `pca_scree_plot`, 132, 132
- `pca_variables_plot`, 132, 133
- `pdf_to_rtf`, 134
- `pformat`, 21, 35, 45, 135
-
- `pi_anonymizer`, 136
- `pipeline_config`, 135, 157
- `pipeline_summary`, 136, 174
- `plan`, 47
- `plike (infix_helpers)`, 99
- `plot.features_percent_miss`, 137
- `plot_categorical_associations`, 137
- `plot_cox_forest`, 138
- `plot_cox_forest_multi`, 140
- `plot_cox_marginal`, 141
- `plot_cox_survival`, 141
- `plot_dynamic_trajectory`, 142
- `plot_features_percent_miss`
 (`plot.features_percent_miss`), 137
- `plot_group_comparison`, 143
- `plot_missingness`, 144
- `plot_pca_bi`, 132, 144
- `plot_survival_km`, 145
- `prcomp`, 127–133, 144, 150, 160
- `predict.censoring_km (censoring_km)`, 50
- `predict.coxnet_model`, 145
- `predict.cv_coxnet (cv_coxnet-methods)`, 69
- `predict.joint_model`, 146
- `print.cbe_cox`, 147
- `print.cbe_cox_check`, 148
- `print.cbe_cox_multi`, 148
- `print.cbe_km`, 148
- `print.joint_model`, 149
- `print_pipeline`, 52, 149
- `proc_contents (get_dataset_info)`, 94
- `proc_pca`, 132, 133, 150
- `profvis_summary`, 151
-
- `R_names`, 161
- `range_norm`, 73, 152
- `read_data_manifest`, 153
- `read_database_from_excel`, 152
- `read_database_metadata`
 (`database_metadata_io`), 71
- `read_excel`, 153, 154, 156
- `read_excel_multiple_headers`, 154, 155
- `read_mapped_section_data`, 154, 173
- `read_raw_table`, 155
- `read_search`, 155, 163, 164
- `read_workbook`, 155, 156
- `render`, 136, 157, 194
- `render_me`, 194
- `render_me (render)`, 157
- `required_pkgs.step_lencode_coxnet`
 (`step_lencode_coxnet`), 169

- required_pkgs.step_lencode_joint_model
 (step_lencode_joint_model), 171
- rossi_data, 159
- rotation_matrix, 160
- run_sas_script, 160
- scale_color_cbe, 162
- scale_color_template
 (scale_colour_template), 162
- scale_colour_gradient2, 163
- scale_colour_gradientn, 163
- scale_colour_template, 162, 177
- scale_fill_cbe, 163
- scale_fill_template
 (scale_colour_template), 162
- scan_data_io, 163
- sd.error, 165
- show, FileOutputs-method
 (FileOutputs-class), 79
- show, FilePath-method (FilePath-class),
 80
- show, FileUses-method (FileUses-class),
 82
- significance_stars, 165
- simulate_section_data, 166
- single_t_test, 119, 120, 125, 166
- step_famd, 167
- step_lencode_coxnet, 169
- step_lencode_joint_model, 171
- stop_if_invalid_manifest, 172
- summarize_pipeline (pipeline_summary),
 136
- summarize_section_by_time, 172
- summarize_tune_results, 173
- summary.tbl_graph, 174
- SumNa, 77, 174
- Surv, 103
- surv_subject_truth, 175
- t.test, 120, 167
- table_two_by_two, 176
- tbl_summary, 173
- temple_brand_path, 177
- temple_colors, 177
- temple_pal, 177, 178
- tests, 47
- theme_cbe, 178
- theme_cbe_deck, 179
- theme_template, 177, 179
- tidy.cbe_loco_mp_coxnet, 180
- tidy.coxnet_model, 180
- tidy.cv_coxnet (cv_coxnet-methods), 69
- tidy.joint_model, 181
- tidy.step_famd (step_famd), 167
- tidy_tmerge_cox, 181
- tumor_long, 182
- tumor_wide, 183
- tunable.step_lencode_coxnet
 (step_lencode_coxnet), 169
- tune_over_alpha, 184
- use_template_brand, 63, 64, 185
- validate_column_mapping, 48, 166, 186
- validate_data_manifest, 55, 172, 186
- var_label, 161
- vfold_cv, 70
- visualize_pipeline, 52, 187
- visualize_pipeline_interactive, 52, 188
- volcano_plot, 188
- words, 189
- write_database_metadata
 (database_metadata_io), 71
- write_database_to_excel, 190
- write_search, 163, 164, 191
- write_workbook
 (write_database_to_excel), 190
- write_xlsx, 190
- z_norm, 73, 194
- zip_render, 192, 194
- zip_reports, 193