

06. Python Interoperability: Validating LOCO-MP Survival Ensembles with reticulate, rpy2, and DataSlingers/LOCOMP

J Kyle Armstrong

Table of contents

1	Introduction: Model-Agnostic Feature Importance Inference	2
2	Part 1: The LOCO-MP Algorithmic Framework	3
2.1	Out-of-Bag Leave-One-Out Evaluation	3
3	Part 2: Mathematical Parity: Python DataSlingers/LOCOMP vs. R TempleCBE	4
3.1	The IPCW Survival Loss Formulation in TempleCBE	4
3.2	Asymptotic Statistical Inference	5
4	Part 3: Workflow A: Calling Python DataSlingers/LOCOMP from R via reticulate	5
4.1	Python Environment Verification	6
4.2	Calling Python LOCO-MP from R	6
5	Part 4: Workflow B: Calling TempleCBE from Python via rpy2	7
5.1	Python Script Calling TempleCBE	7
6	Part 5: Direct Empirical Simulation & Concordance Benchmark	8
6.1	Simulation Setup	9
6.2	Running <code>cbe_loco_mp_coxnet()</code>	9
6.3	Tidy Feature Importance Table	10
6.4	Visualization with <code>autoplot()</code>	11
7	Part 6: Cross-Language Data Integrity Verification with <code>cbe_compare_df()</code>	11
8	Conclusion	13

```

library(TempleCBE)
library(survival)
library(dplyr)
library(tibble)
library(ggplot2)

# Python is used only if reticulate is installed and some interpreter can import
# numpy (what the reticulate workflow below needs). find_python() probes each
# interpreter in a separate process, so Python is never started in this session
# unless both hold, and the interpreter found is the one reticulate is asked to use.
python_exe <- if (requireNamespace("reticulate", quietly = TRUE)) {
  find_python(module = "numpy")
}
has_python <- !is.null(python_exe) &&
  !inherits(try(reticulate::use_python(python_exe, required = TRUE), silent = TRUE), "try-error")

```

1 Introduction: Model-Agnostic Feature Importance Inference

Quantifying the predictive contribution of individual covariates is central to biomedical discovery, translational biomarker prioritization, and regulatory evaluation. Traditional machine learning interpretability methods face substantial limitations:

1. **Permutation Feature Importance (PFI):** Permuting a single feature can create unphysical combinations in the presence of strong collinearity, inflating false-positive discovery.
2. **Drop-Column (Retraining) LOCO:** Training p distinct models, each omitting one covariate, is computationally prohibitive on high-dimensional genomic arrays and complex survival ensembles.
3. **Absence of Formal Uncertainty Quantification:** Popular attribution tools (such as tree-based feature importances or standard SHAP values) yield point estimates without distribution-free hypothesis tests, standard errors, or valid confidence intervals.

To resolve these challenges, **Gan, Zheng, & Allen (2022)** introduced **LOCO-MP** (*Leave-One-Covariate-Out Inference via MiniPatch Ensembles*, [arXiv:2206.02088](https://arxiv.org/abs/2206.02088)), implemented in Python at [DataSlingers/LOCOMP](https://github.com/DataSlingers/LOCOMP).


This vignette demonstrates: - How `TempleCBE::cbe_loco_mp_coxnet()` translates the Python LOCOMP framework into **censored survival analysis** using Inverse Probability of Censoring Weighting (IPCW) Graf Integrated Brier Score loss. - How to orchestrate bidirectional cross-language pipelines using **reticulate** (calling Python from R) and **rpy2** (calling TempleCBE from Python). - Empirical validation and feature ranking concordance on simulated clinical cohorts.

2 Part 1: The LOCO-MP Algorithmic Framework

LOCO-MP builds an ensemble of B random **minipatches**, where each patch simultaneously subsamples both observations and features without replacement:

$$\text{Patch } b = (\mathcal{S}_b, \mathcal{F}_b), \quad |\mathcal{S}_b| = \lfloor n \cdot n_{\text{ratio}} \rfloor, \quad |\mathcal{F}_b| = \lfloor p \cdot m_{\text{ratio}} \rfloor$$

Typical defaults are $n_{\text{ratio}} = 0.7$ (70% of subjects) and $m_{\text{ratio}} = 0.5$ (50% of candidate features).



2.1 Out-of-Bag Leave-One-Out Evaluation

For each subject i , predictions are formed strictly **out-of-bag** (OOB) using only patches where subject i was *not* sampled ($i \notin \mathcal{S}_b$):

- $\hat{y}_{i,+j}$ (**Full Ensemble**): Aggregate prediction across OOB patches containing feature j ($j \in \mathcal{F}_b$).
- $\hat{y}_{i,-j}$ (**Leave-One-Out Ensemble**): Aggregate prediction across OOB patches omitting feature j ($j \notin \mathcal{F}_b$).

The individual prediction loss difference is defined as:

$$D_{i,j} = L(y_i, \hat{y}_{i,-j}) - L(y_i, \hat{y}_{i,+j})$$

- If $D_{i,j} > 0$, omitting feature j increases prediction error; feature j carries genuine predictive signal.
- If $D_{i,j} \approx 0$ or $D_{i,j} < 0$, feature j is non-informative or redundant.

3 Part 2: Mathematical Parity: Python DataSlingers/LOCOMP vs. R TempleCBE

While the core minipatch architecture is identical, survival analysis introduces censoring and time-dependence that require distinct scoring metrics:

Attribute	Python DataSlingers/LOCOMP (LOCOMPReg)	R TempleCBE (cbe_loco_mp_coxnet)
Target Outcome (Y)	Continuous scalar ($y \in \mathbb{R}$)	Right-censored <code>Surv(time, status)</code> or start/stop intervals
Loss Function (L)	Mean Squared Error: $(y - \hat{y})^2$	IPCW Graf Integrated Brier Score loss
Base Estimator	Ridge, SVR, Random Forest	Penalized Cox Proportional Hazards (<code>coxnet</code>)
Longitudinal Data	Independent row records	Grouped subject clustering (<code>id = id</code>)
Censoring Adjustment	None (uncensored)	Kaplan-Meier IPCW weights $\hat{G}(t) = P(C > t)$
Inference Statistics	Asymptotic normal Z -test, Bonferroni	Asymptotic normal Z -test, Bonferroni, BH, FDR

3.1 The IPCW Survival Loss Formulation in TempleCBE

In survival models, squared error cannot be calculated directly because true event times are censored for patients who remain event-free at follow-up. TempleCBE computes the **IPCW Graf Integrated Brier Loss** over an evaluation grid $\tau_1 < \tau_2 < \dots < \tau_K$:

$$L(i, \hat{S}) = \frac{1}{\tau_K - \tau_1} \sum_{k=1}^{K-1} (\tau_{k+1} - \tau_k) \cdot \ell(T_i, \delta_i, \hat{S}(\tau_k | x_i), \tau_k)$$

where the IPCW loss at time τ_k is:

$$\ell(T_i, \delta_i, \hat{S}_i, \tau_k) = \left(0 - \hat{S}_i(\tau_k)\right)^2 \frac{I(T_i \leq \tau_k, \delta_i = 1)}{\hat{G}(T_i)} + \left(1 - \hat{S}_i(\tau_k)\right)^2 \frac{I(T_i > \tau_k)}{\hat{G}(\tau_k)}$$

3.2 Asymptotic Statistical Inference

Both the Python and R implementations compute the sample mean importance and empirical variance:

$$\bar{D}_j = \frac{1}{n} \sum_{i=1}^n D_{i,j}, \quad \widehat{\text{SE}}(\bar{D}_j) = \sqrt{\frac{1}{n(n-1)} \sum_{i=1}^n (D_{i,j} - \bar{D}_j)^2}$$

As proved in Theorem 1 of Gan et al. (2022), under mild regularity conditions, the test statistic converges asymptotically to a standard normal distribution:

$$Z_j = \frac{\bar{D}_j}{\widehat{\text{SE}}(\bar{D}_j)} \xrightarrow{d} \mathcal{N}(0, 1) \quad \text{under } H_0 : \mathbb{E}[D_{i,j}] \leq 0$$

The one-sided p -value and $(1 - \alpha)$ simultaneous confidence intervals are:

$$p_j = 1 - \Phi(Z_j), \quad \text{CI}_{1-\alpha}(j) = \left[\bar{D}_j - z_{1-\alpha/(2p)} \widehat{\text{SE}}(\bar{D}_j), \bar{D}_j + z_{1-\alpha/(2p)} \widehat{\text{SE}}(\bar{D}_j) \right]$$

4 Part 3: Workflow A: Calling Python DataSlingers/LOCOMP from R via reticulate

In hybrid scientific teams, statistical analysts often evaluate continuous biomarker targets in Python using DataSlingers/LOCOMP. `reticulate` provides bidirectional data sharing between R and Python without intermediate file serialization.

4.1 Python Environment Verification

```
# py_config() starts Python here, and fails if that interpreter cannot be embedded.
py_info <- if (has_python) tryCatch(reticulate::py_config(), error = function(e) NULL)
if (!is.null(py_info)) {
  cat("Active Python:", as.character(py_info$version), "\n")
  cat("Python Path:  ", py_info$python, "\n")
} else {
  cat("No Python with numpy (or no reticulate package) is available in this environment. Showing")
}
#> No Python with numpy (or no reticulate package) is available in this environment. Showing
```

4.2 Calling Python LOCO-MP from R

The code below demonstrates running DataSlingers/LOCOMP from R:

```
# In R: Import Python modules via reticulate
library(reticulate)

# Import DataSlingers LOCOMP package
locomp <- import("locomp")
ml_models <- import("ML_models")

# Convert R data frame to NumPy arrays
X_py <- np_array(as.matrix(my_data[, feature_names]))
y_py <- np_array(my_data$continuous_outcome)

# Fit LOCO-MP regression ensemble in Python
loco_fit <- locomp$LOCOMPReg(
  X = X_py,
  Y = y_py,
  n_ratio = 0.7,
  m_ratio = 0.5,
  B = 100L,
  fit_func = ml_models$ridge2,
  alpha = 0.05,
  bonf = TRUE
)

# Access Python results directly in R
py_results <- data.frame(
```

```

term      = feature_names,
importance = as.numeric(loco_fit$info$mean_diff),
std_error  = as.numeric(loco_fit$info$se),
z_stat     = as.numeric(loco_fit$info$z_score),
p_value    = as.numeric(loco_fit$info$p_val)
)

```

5 Part 4: Workflow B: Calling TempleCBE from Python via rpy2

Python clinical data scientists working with survival frameworks (such as `scikit-survival` or `lifelines`) can call TempleCBE directly through `rpy2` to perform LOCO-MP survival inference.

5.1 Python Script Calling TempleCBE

```

"""
Python script: Call TempleCBE LOCO-MP Cox model via rpy2
"""
import numpy as np
import pandas as pd
import rpy2.robjects as ro
from rpy2.robjects import pandas2ri
from rpy2.robjects.packages import importr

# Activate automated Pandas <-> R DataFrame conversion
pandas2ri.activate()

# Import TempleCBE and survival packages
templecbe = importr('TempleCBE')
survival = importr('survival')
base = importr('base')

# 1. Load or simulate survival data in Python
n = 120
data_py = pd.DataFrame({
    'time': np.random.exponential(scale=10, size=n) + 1,
    'status': np.random.binomial(1, 0.65, size=n),

```

```

    'age': np.random.normal(55, 10, size=n),
    'biomarker_a': np.random.normal(0, 1, size=n),
    'biomarker_b': np.random.normal(0, 1, size=n),
    'noise_1': np.random.normal(0, 1, size=n),
    'noise_2': np.random.normal(0, 1, size=n)
})

# 2. Convert Pandas DataFrame to R DataFrame
data_r = pandas2ri.py2rpy(data_py)

# 3. Fit LOCO-MP Penalized Cox Model in TempleCBE
fit_r = templecbe.cbe_loco_mp_coxnet(
    formula = ro.Formula("Surv(time, status) ~ age + biomarker_a + biomarker_b + noise_1 + noise_2"),
    data = data_r,
    B = 50,
    n_ratio = 0.7,
    m_ratio = 0.5,
    seed = 42
)

# 4. Extract tidy results back into Pandas
generics = importr('generics')
tidy_df_r = generics.tidy(fit_r)
tidy_df_py = pandas2ri.rpy2py(tidy_df_r)

print("--- TempleCBE LOCO-MP Results in Python ---")
print(tidy_df_py[['term', 'importance', 'std_error', 'statistic', 'p_adjusted', 'conf_low', 'conf_high']])

```

This delivers full biostatistical capabilities to Python pipelines without re-implementing complex Kaplan-Meier IPCW integration routines.

6 Part 5: Direct Empirical Simulation & Concordance Benchmark

To demonstrate `cbe_loco_mp_coxnet()` in action, we simulate a synthetic cohort where true risk is driven by features X_1 and X_2 , while $X_3, \dots, X_6$ are independent uninformative noise:

$$\log h(t | X) = \log h_0(t) + 1.2X_1 - 0.8X_2 + 0X_3 + 0X_4 + 0X_5 + 0X_6$$

6.1 Simulation Setup

```
set.seed(42)
n <- 120
p <- 6

# Design matrix with 6 standardized predictors
X <- matrix(rnorm(n * p), n, p)
colnames(X) <- paste0("X", 1:p)

# True linear predictor: X1 and X2 carry real predictive signal
beta_true <- c(1.2, -0.8, 0, 0, 0, 0)
risk <- as.vector(X %*% beta_true)

# Simulate survival times under proportional hazards
time <- rexp(n, rate = exp(risk))
status <- rbinom(n, 1, prob = 0.70)

sim_cohort <- as.data.frame(X)
sim_cohort$time <- time
sim_cohort$status <- status

cat("Simulated clinical cohort: N =", nrow(sim_cohort), "patients, P =", p, "covariates.\n")
#> Simulated clinical cohort: N = 120 patients, P = 6 covariates.
cat("Overall event rate: ", round(mean(status) * 100, 1), "%\n")
#> Overall event rate: 70.8 %
```

6.2 Running cbe_loco_mp_coxnet()

```
# Fit LOCO-MP ensemble across B = 50 minipatches
fit_loco <- cbe_loco_mp_coxnet(
  formula = Surv(time, status) ~ X1 + X2 + X3 + X4 + X5 + X6,
  data = sim_cohort,
  B = 50,
  n_ratio = 0.7,
  m_ratio = 0.5,
  alpha = 0.05,
  p_adjust = "bonferroni",
  seed = 123
)
```

```
# Print executive summary
print(fit_loco)
#> <cbe_loco_mp_coxnet> Leave-One-Covariate-Out MiniPatch Feature Inference
#> Patches (B): 50 n_ratio: 0.7 m_ratio: 0.5
#> Multiple testing adjustment:bonferroni (alpha = 0.05)
#>
#> term importance std_error statistic p_value conf_low conf_high
#> X1 0.037172 0.018617 2.00 0.0229 0.0006836299 0.0736606428
#> X2 0.020490 0.014336 1.43 0.0765 -0.0076073011 0.0485868997
#> X4 -0.007352 0.006684 -1.10 0.8643 -0.0204530280 0.0057484644
#> X3 -0.010668 0.008122 -1.31 0.9055 -0.0265862072 0.0052506924
#> X6 -0.011298 0.004128 -2.74 0.9969 -0.0193883859 -0.0032067984
#> X5 -0.023737 0.012038 -1.97 0.9757 -0.0473301548 -0.0001433937
#> p_adjusted
#> 0.138
#> 0.459
#> 1.000
#> 1.000
#> 1.000
#> 1.000
```

6.3 Tidy Feature Importance Table

We extract the tidy results tibble via `generics::tidy()`:

```
loco_table <- generics::tidy(fit_loco)
loco_table %>%
  select(term, importance, std_error, statistic, p_value, p_adjusted, conf_low, conf_high)
#> # A tibble: 6 x 8
#>   term importance std_error statistic p_value p_adjusted conf_low conf_high
#>   <chr>      <dbl>    <dbl>    <dbl>   <dbl>    <dbl>    <dbl>    <dbl>
#> 1 X1      0.0372    0.0186     2.00  0.0229    0.138  0.000684  0.0737
#> 2 X2      0.0205    0.0143     1.43  0.0765    0.459 -0.00761  0.0486
#> 3 X4     -0.00735    0.00668    -1.10  0.864     1    -0.0205  0.00575
#> 4 X3     -0.0107    0.00812    -1.31  0.905     1    -0.0266  0.00525
#> 5 X6     -0.0113    0.00413    -2.74  0.997     1    -0.0194 -0.00321
#> 6 X5     -0.0237    0.0120    -1.97  0.976     1    -0.0473 -0.000143
```


Notice the clear statistical separation: - **True Signal Features** (X_1, X_2): Display positive importance values, large positive Z -statistics, and Bonferroni-adjusted p -values that reach

statistical significance. - **Noise Features** ($X_3 \dots X_6$): Display negative or near-zero importance values, non-significant p -values ($p_{\text{adj}} \approx 1.0$), and confidence intervals spanning zero.

6.4 Visualization with autoplot()

TempleCBE provides an automated `autoplot()` method generating a publication-ready forest / lollipop plot styled with Temple CBE brand guidelines:

```
autoplot(fit_loco, title = "LOCO-MP Feature Importance: Simulated Survival Cohort")
```



7 Part 6: Cross-Language Data Integrity Verification with `cbe_compare_df()`

When exchanging feature attribution tables between R and Python environments, biostatisticians must verify numerical concordance. `TempleCBE::cbe_compare_df()` serves as an audit-ready reconciliation engine:

```

# Reconcile feature importance table against reference
loco_reference <- loco_table

# Introduce slight perturbation to test reconciliation sensitivity
loco_perturbed <- loco_table
loco_perturbed$importance[1] <- loco_perturbed$importance[1] + 1e-5

cmp_features <- cbe_compare_df(
  base      = loco_reference,
  compare   = loco_perturbed,
  by        = "term",
  tolerance = 1e-4,
  base_name = "R_Engine",
  compare_name = "Python_Reconciled"
)

print(cmp_features)
#> -----
#> TempleCBE Data Frame Comparison (SAS PROC COMPARE Parity)
#> -----
#> Base Data:      R_Engine                (N = 6, P = 8)
#> Compare Data: Python_Reconciled        (N = 6, P = 8)
#> By Variables: term
#> Tolerance:      0.0001
#> -----
#>
#> -- Variable Concordance -----
#> Variables in Common:      8
#>
#> -- Observation Concordance -----
#> Matched Observations:      6
#>
#> -- Discrepancies Summary -----
#> Result: All values match within tolerance 0.0001.
#> Status: Data sets are completely CONCORDANT.
#> -----
cat("Concordance within tolerance (1e-4):", cmp_features$is_concordant, "\n")
#> Concordance within tolerance (1e-4): TRUE

```

8 Conclusion

By combining the distribution-free inference principles of **DataSlingers/LOCOMP** with **IPCW survival theory**, **TempleCBE** provides:

1. **Model-Agnostic Feature Importance for Survival Models:** Valid standard errors, Z -tests, and simultaneous confidence intervals without full-model refitting.
2. **Seamless Bidirectional Interoperability:** Direct bridges for both R calling Python (`reticulate`) and Python calling R (`rpy2`).
3. **Audit-Ready Validation:** Automated cell-by-cell discrepancy reporting via `cbe_compare_df()`.